DTAI
Declaratieve Talen en Artificiële Intelligentie

KATHOLIEKE UNIVERSITEIT LEUVEN
Faculteit Ingenieurswetenschappen

Exposé: An ontology for machine learning experimentation

Joaquin Vanschoren, K.U.Leuven (Belgium), U. Leiden (The Netherlands)
Larisa Soldatova, University of Aberystwyth (UK)

DM Ontology Jamboree 2010

DTAI
Declaratieve Talen en Artificiële Intelligentie

KATHOLIEKE UNIVERSITEIT LEUVEN
Faculteit Ingenieurswetenschappen

Exposé: An ontology for machine learning experimentation

Joaquin Vanschoren, K.U.Leuven (Belgium), U. Leiden (The Netherlands)
Larisa Soldatova, University of Aberystwyth (UK)

DM Ontology Jamboree 2010

Overview

Ontology lessons
Exposé ontology
Use cases

Ontology lessons

What did we learn from other ontologies

Ontology design

- Start from accepted classes & properties (top-level ontologies, e.g. OBI, RO)
- If possible, reuse prior ontologies to build on their knowledge/consensus
- Use ontology design patterns: reusable patterns for recurrent problems
  - http://ontologydesignpatterns.org
- Check clarity, consistency, extensibility, minimal commitment

Ontology recap:
OntoDM (Panov et al., ’09,’10)

- Aim: unified framework for DM research, builds on BFO

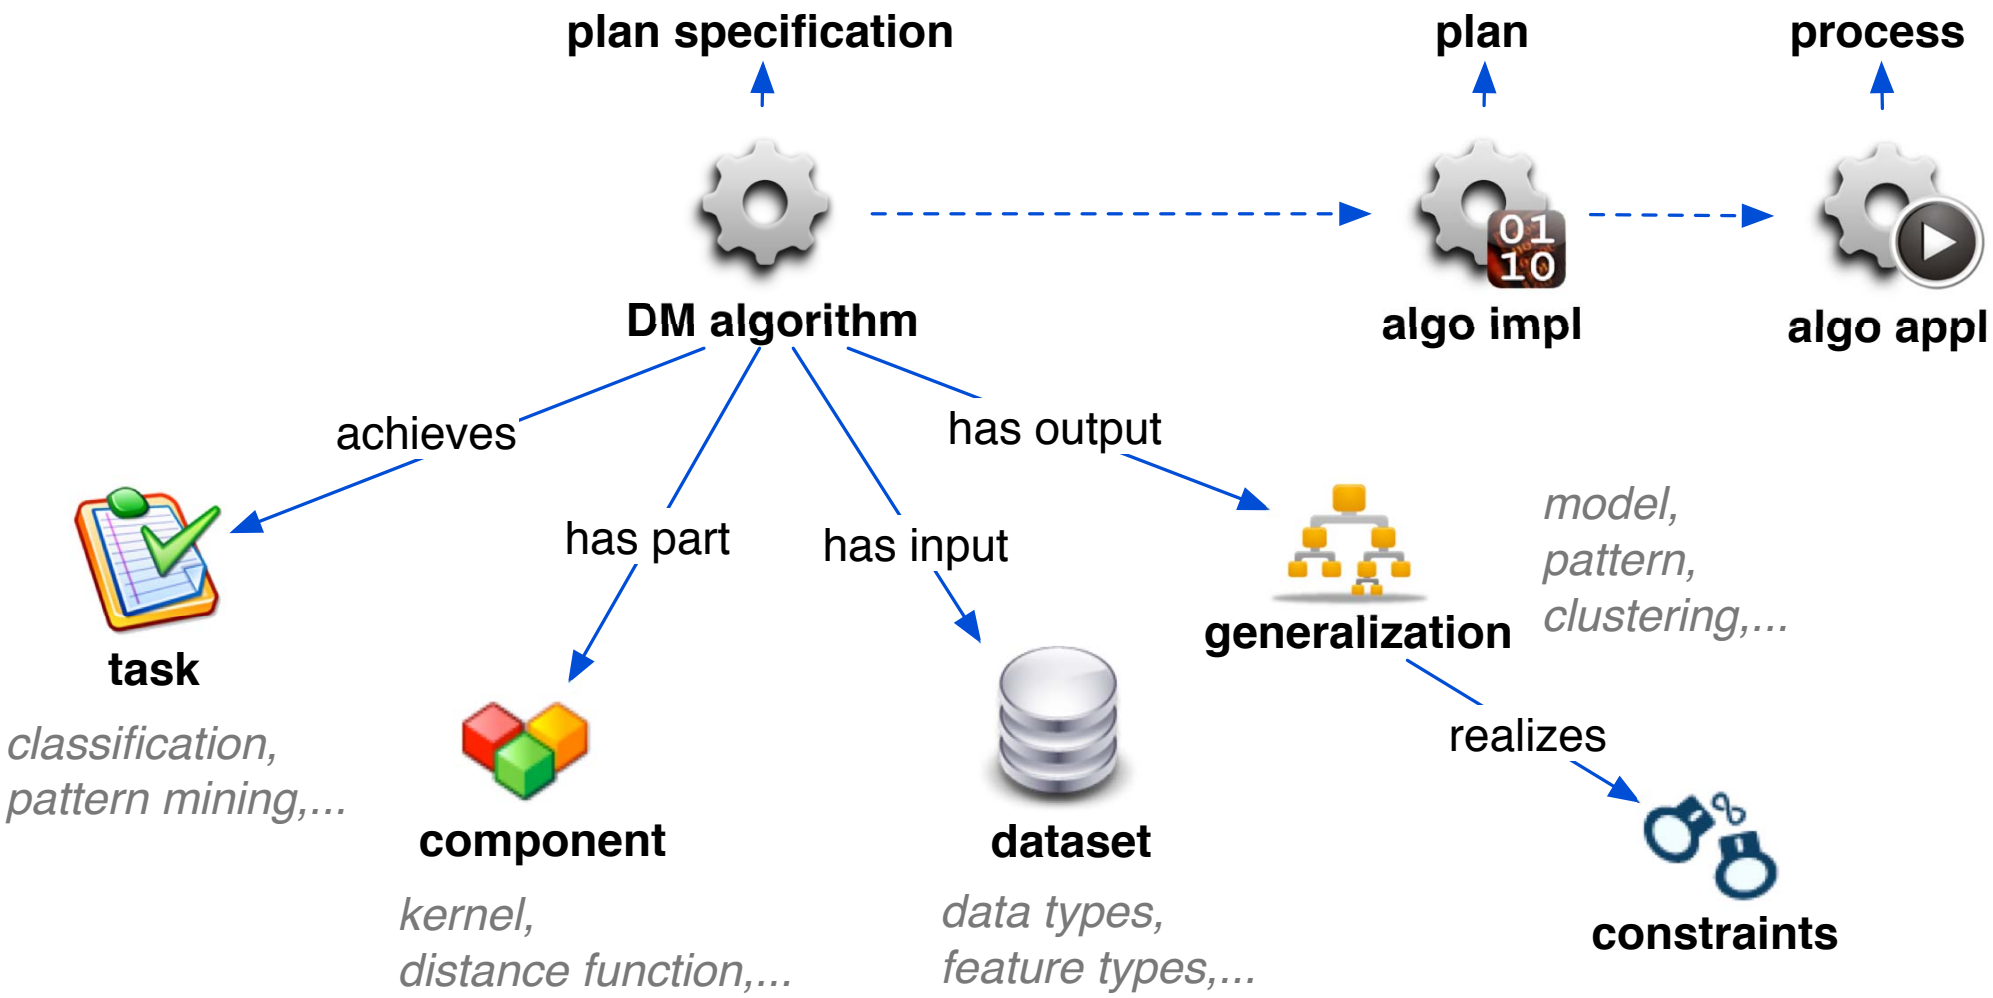

Ontology recap:
OntoDM (Panov et al., ’09,’10)

- Aim: unified framework for DM research, builds on BFO

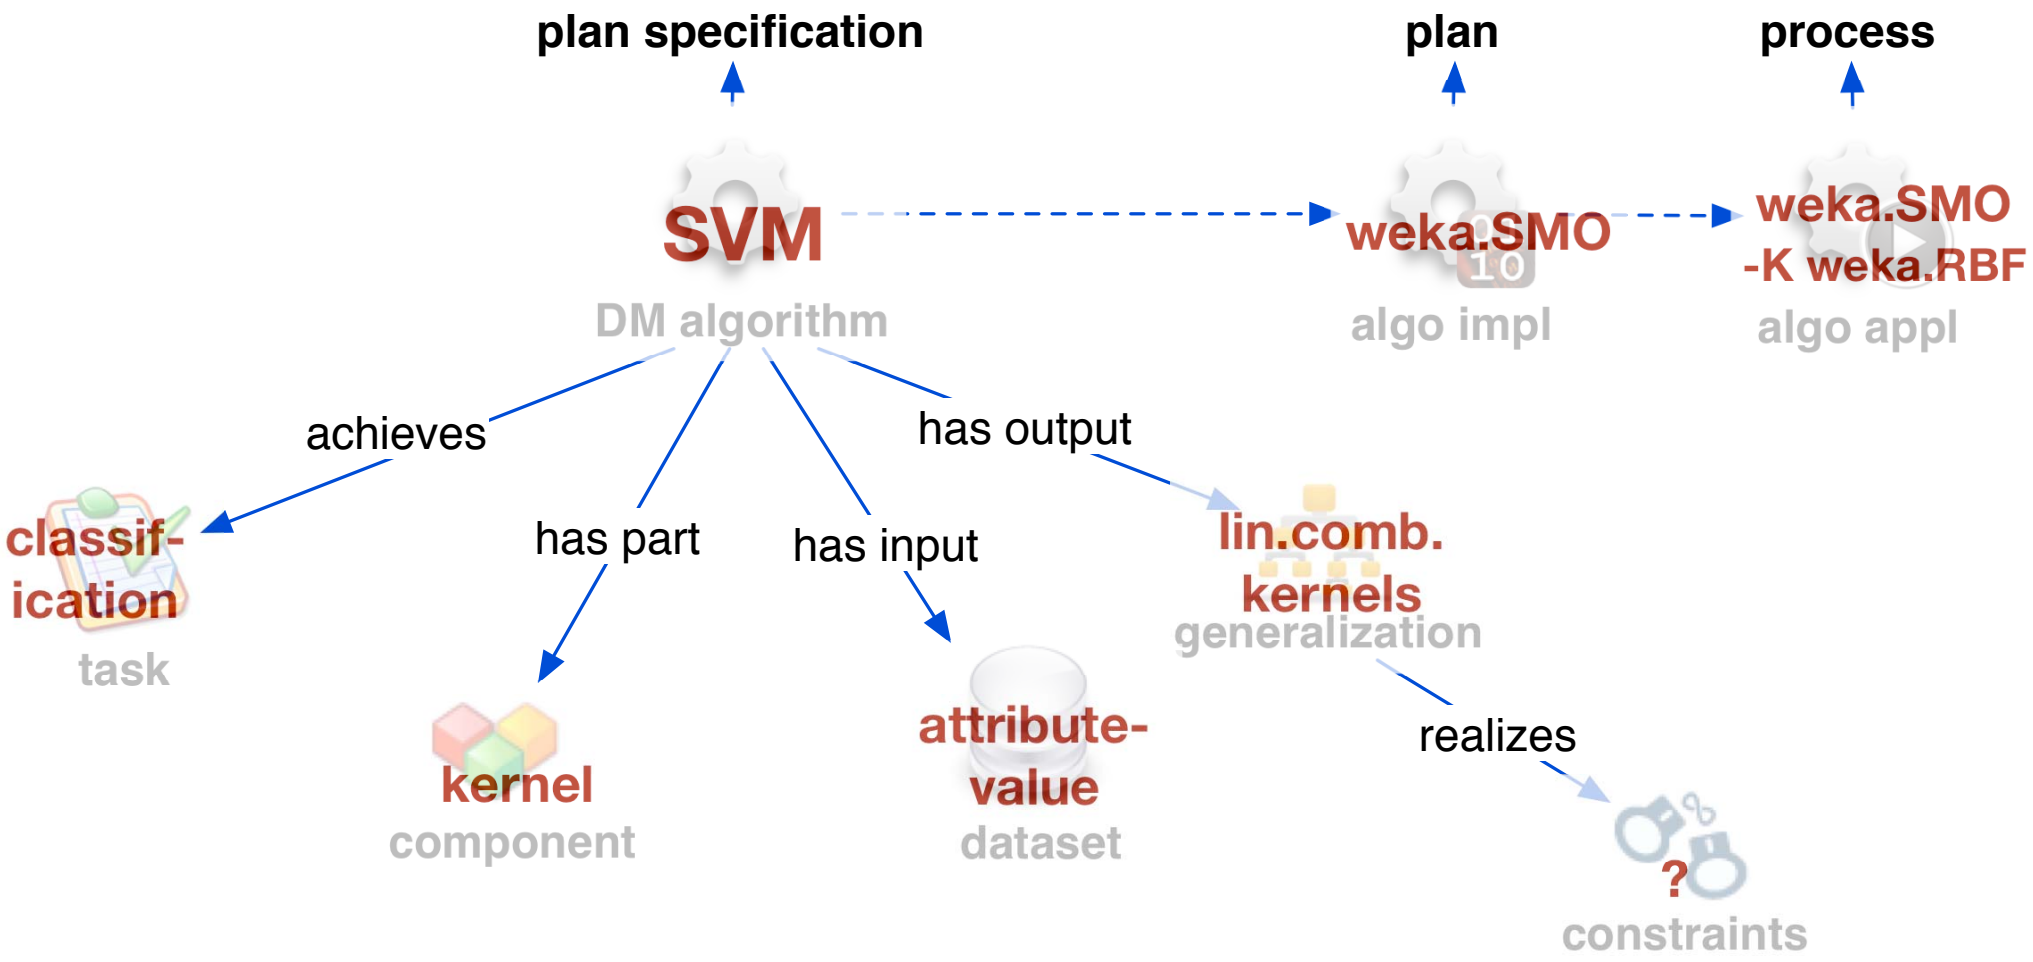

Ontology recap: DMOP (Hilario et al., '09)

- Model internal structure of learning algorithms

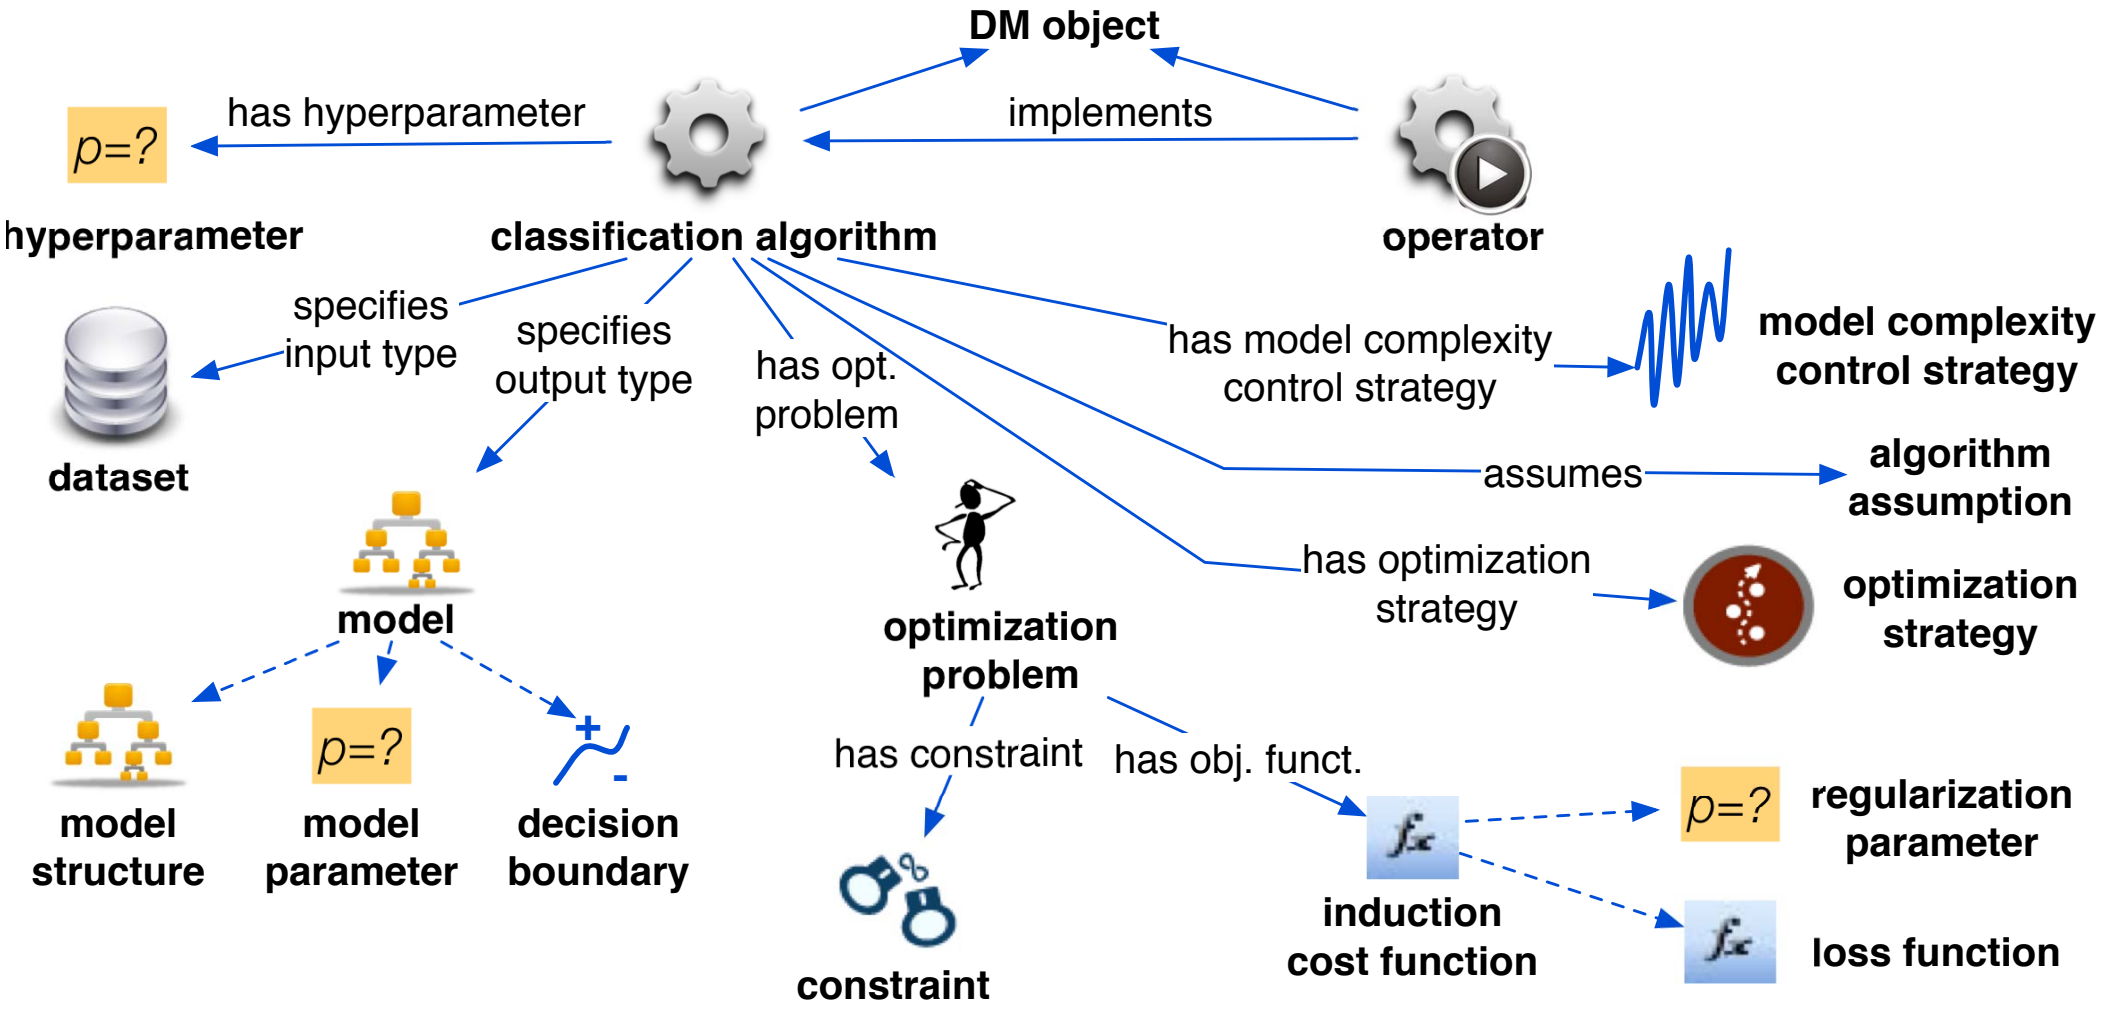

Ontology recap: DMOP (Hilario et al., '09)

- Model internal structure of learning algorithms

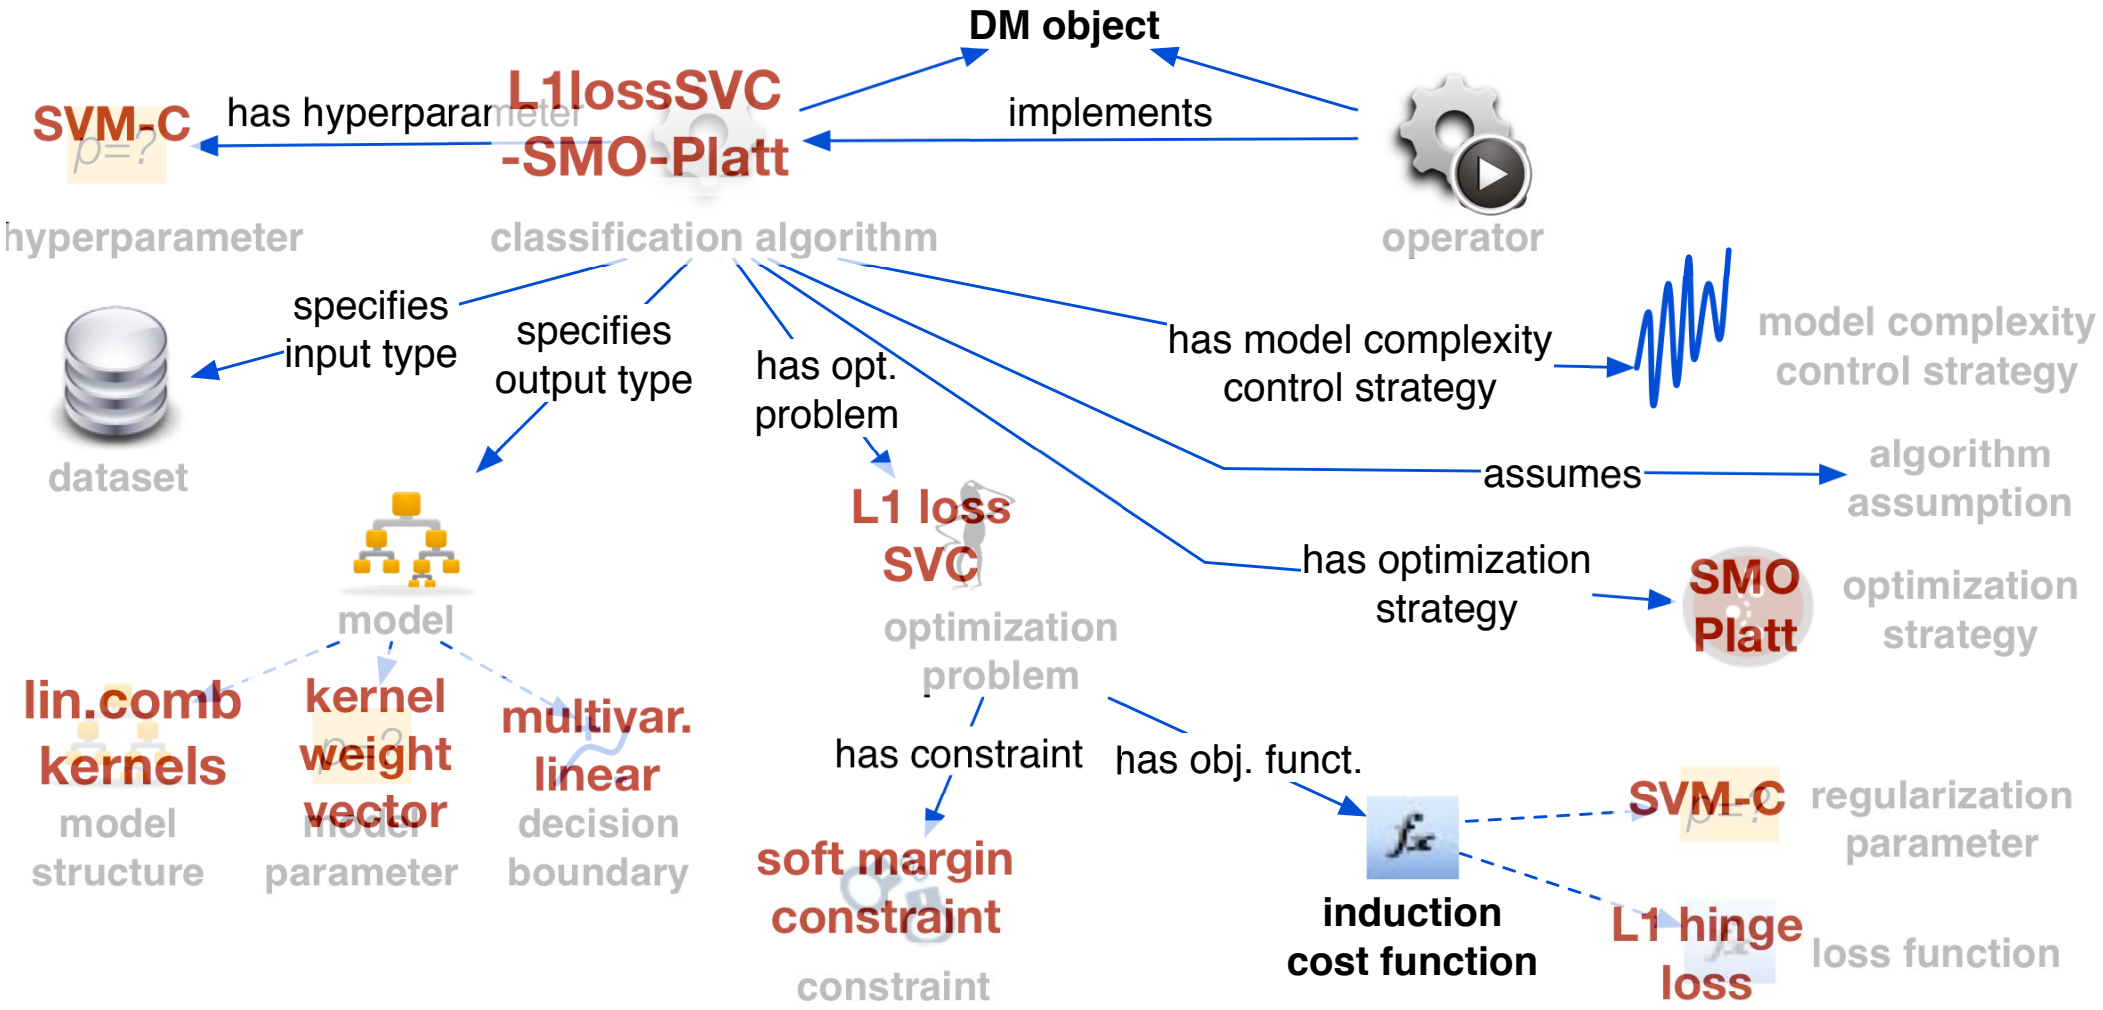

Ontology recap:
DMWF (Kietz et al., ’09)

- Reason about KD operators: in/outputs, conditions/effects (SWRL rules)

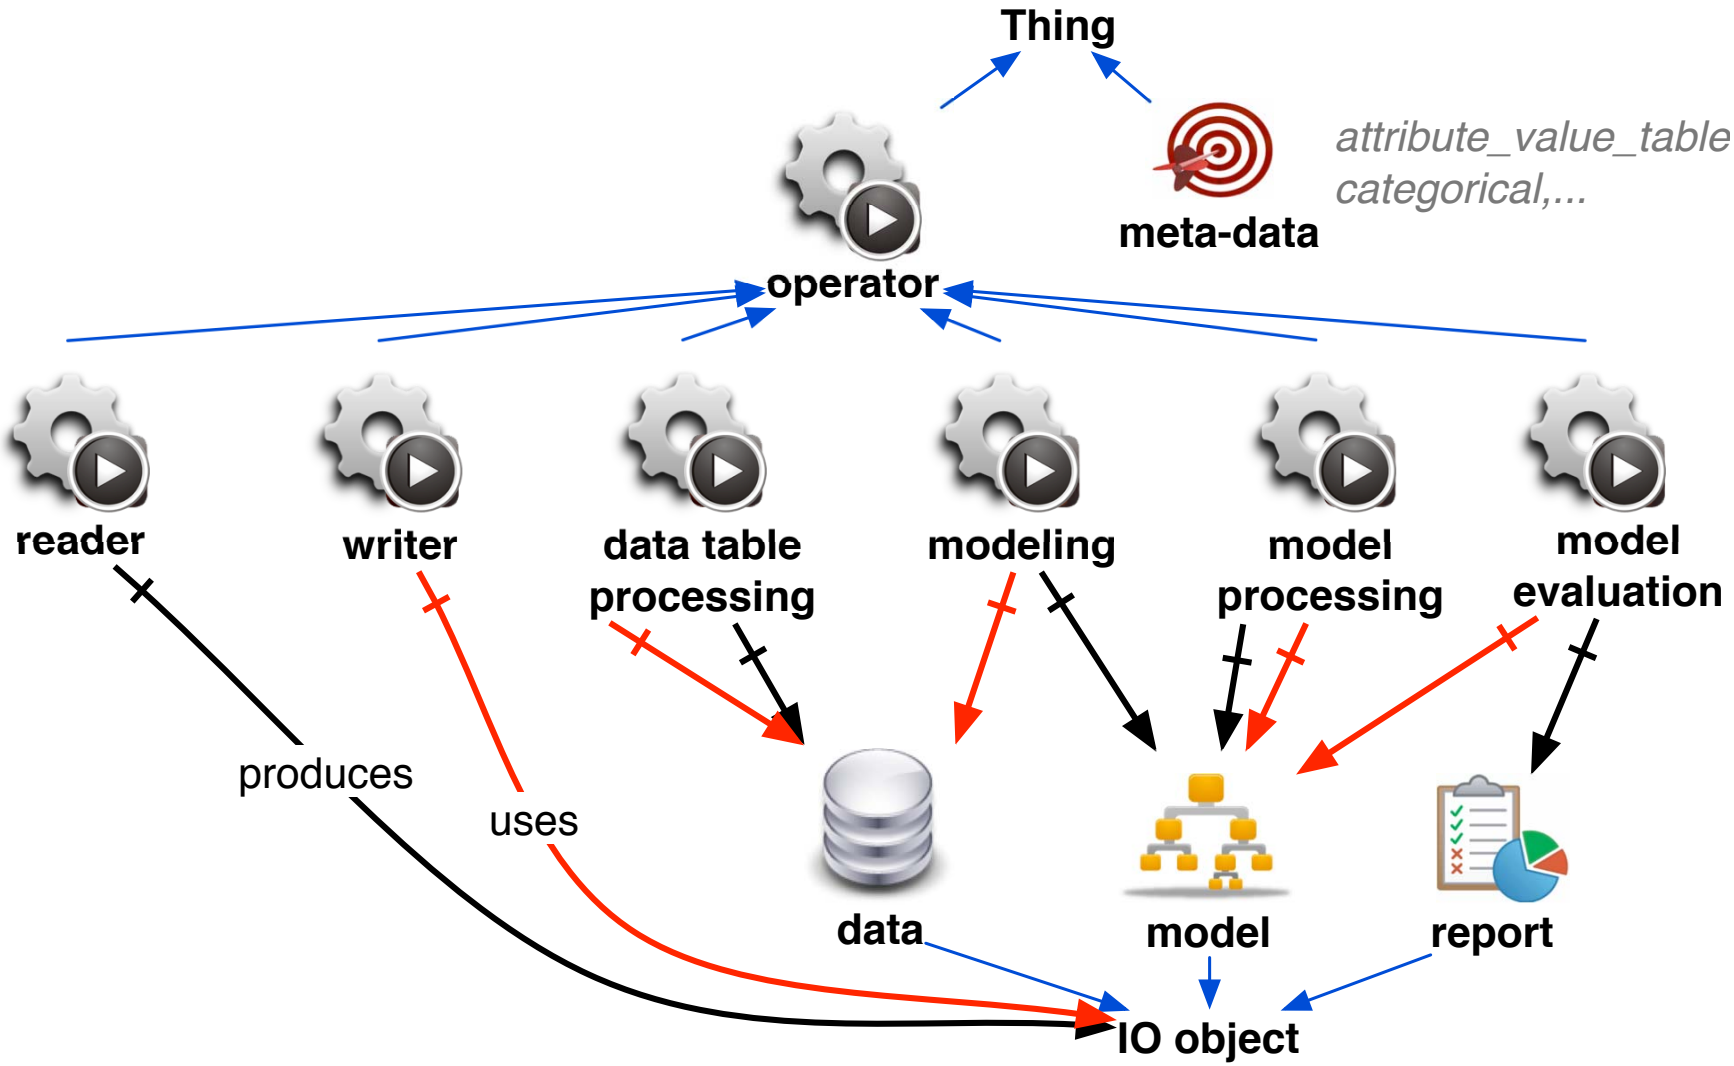

Ontology recap:
DMWF (Kietz et al., ’09)

- Reason about KD operators: in/outputs, conditions/effects (SWRL rules)

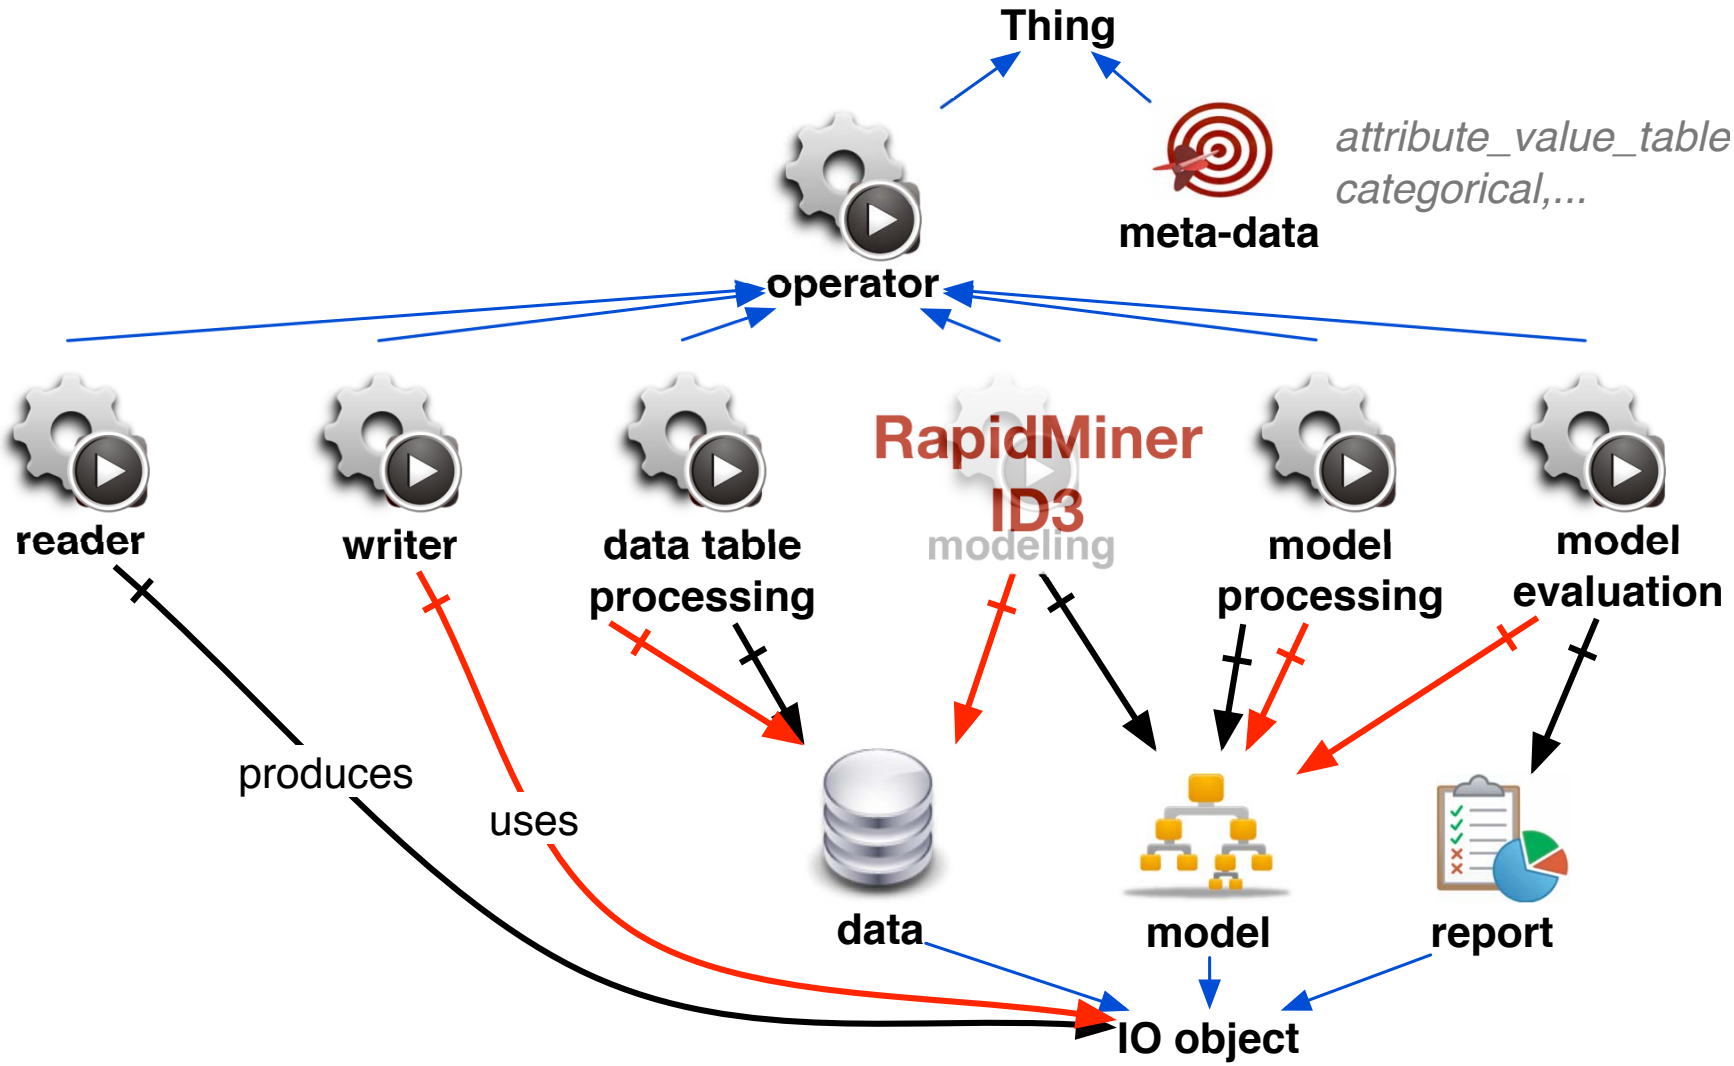

Ontology recap:
DMWF (Kietz et al., ’09)

- Reason about KD operators: in/outputs, conditions/effects (SWRL rules)

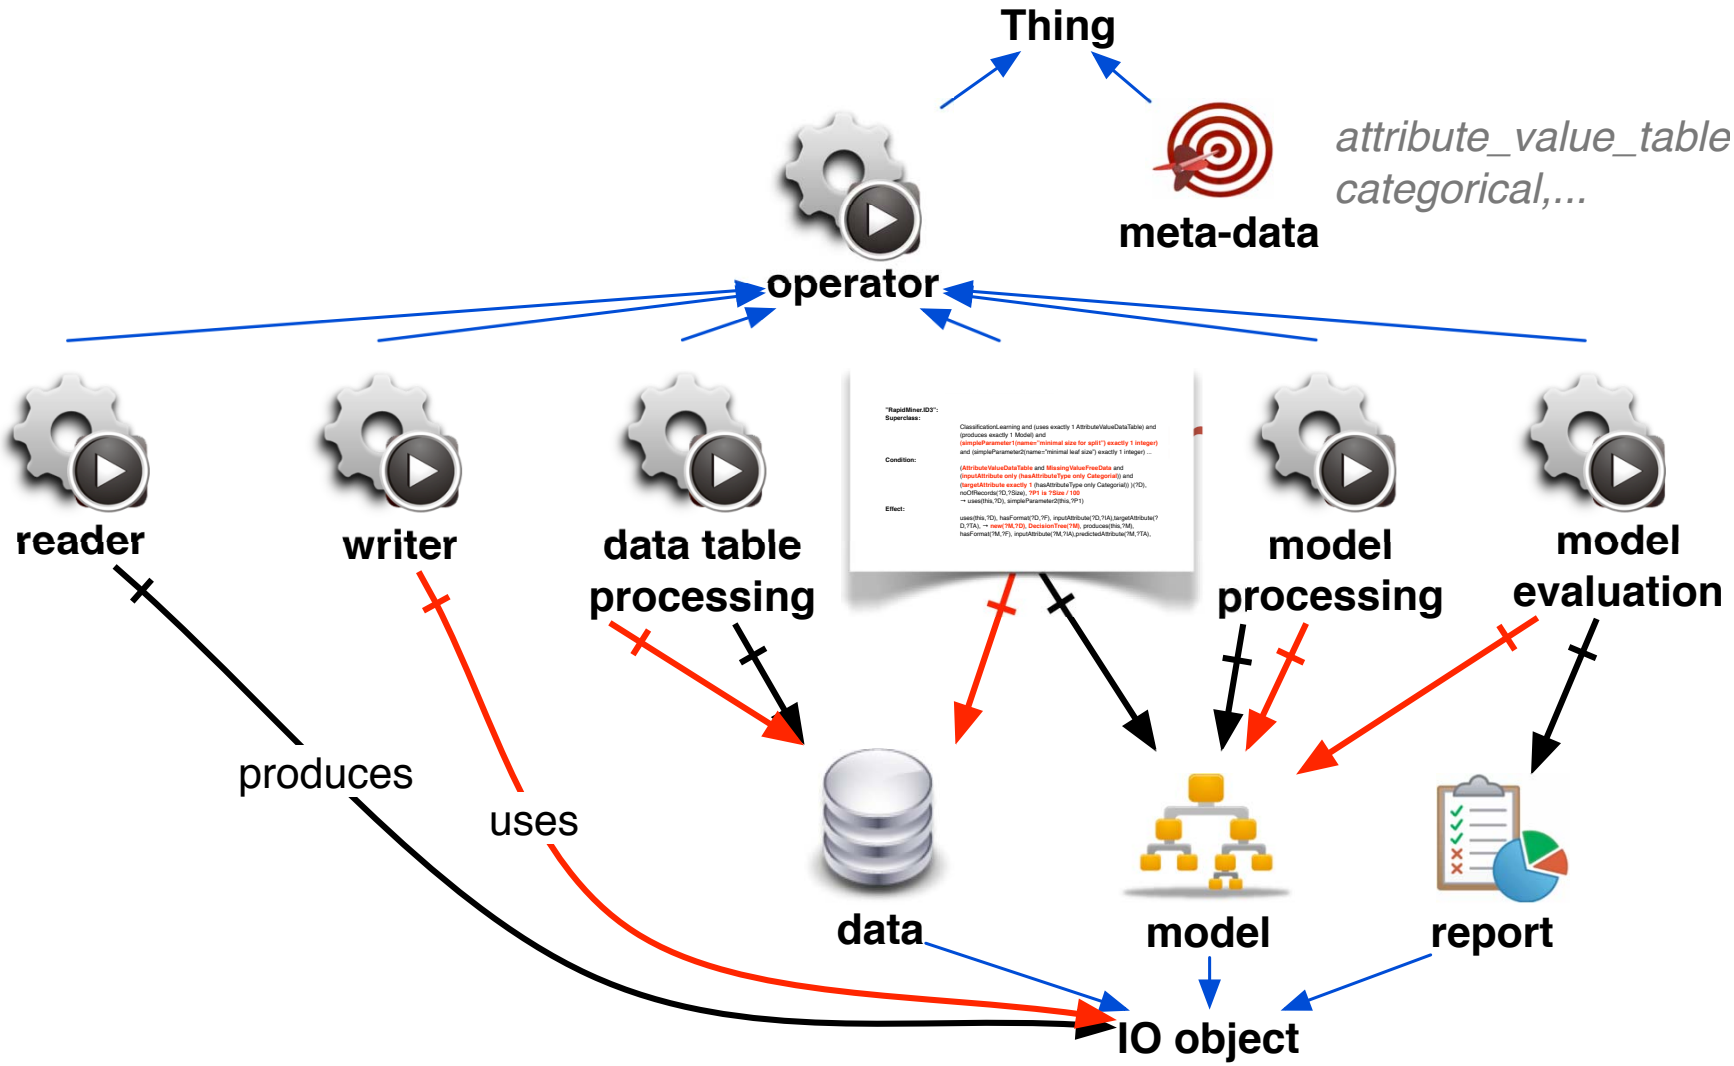

Ontology recap:
DMWF (Kietz et al., ’09)

- Reason about KD operators: in/outputs, conditions/effects (SWRL rules)

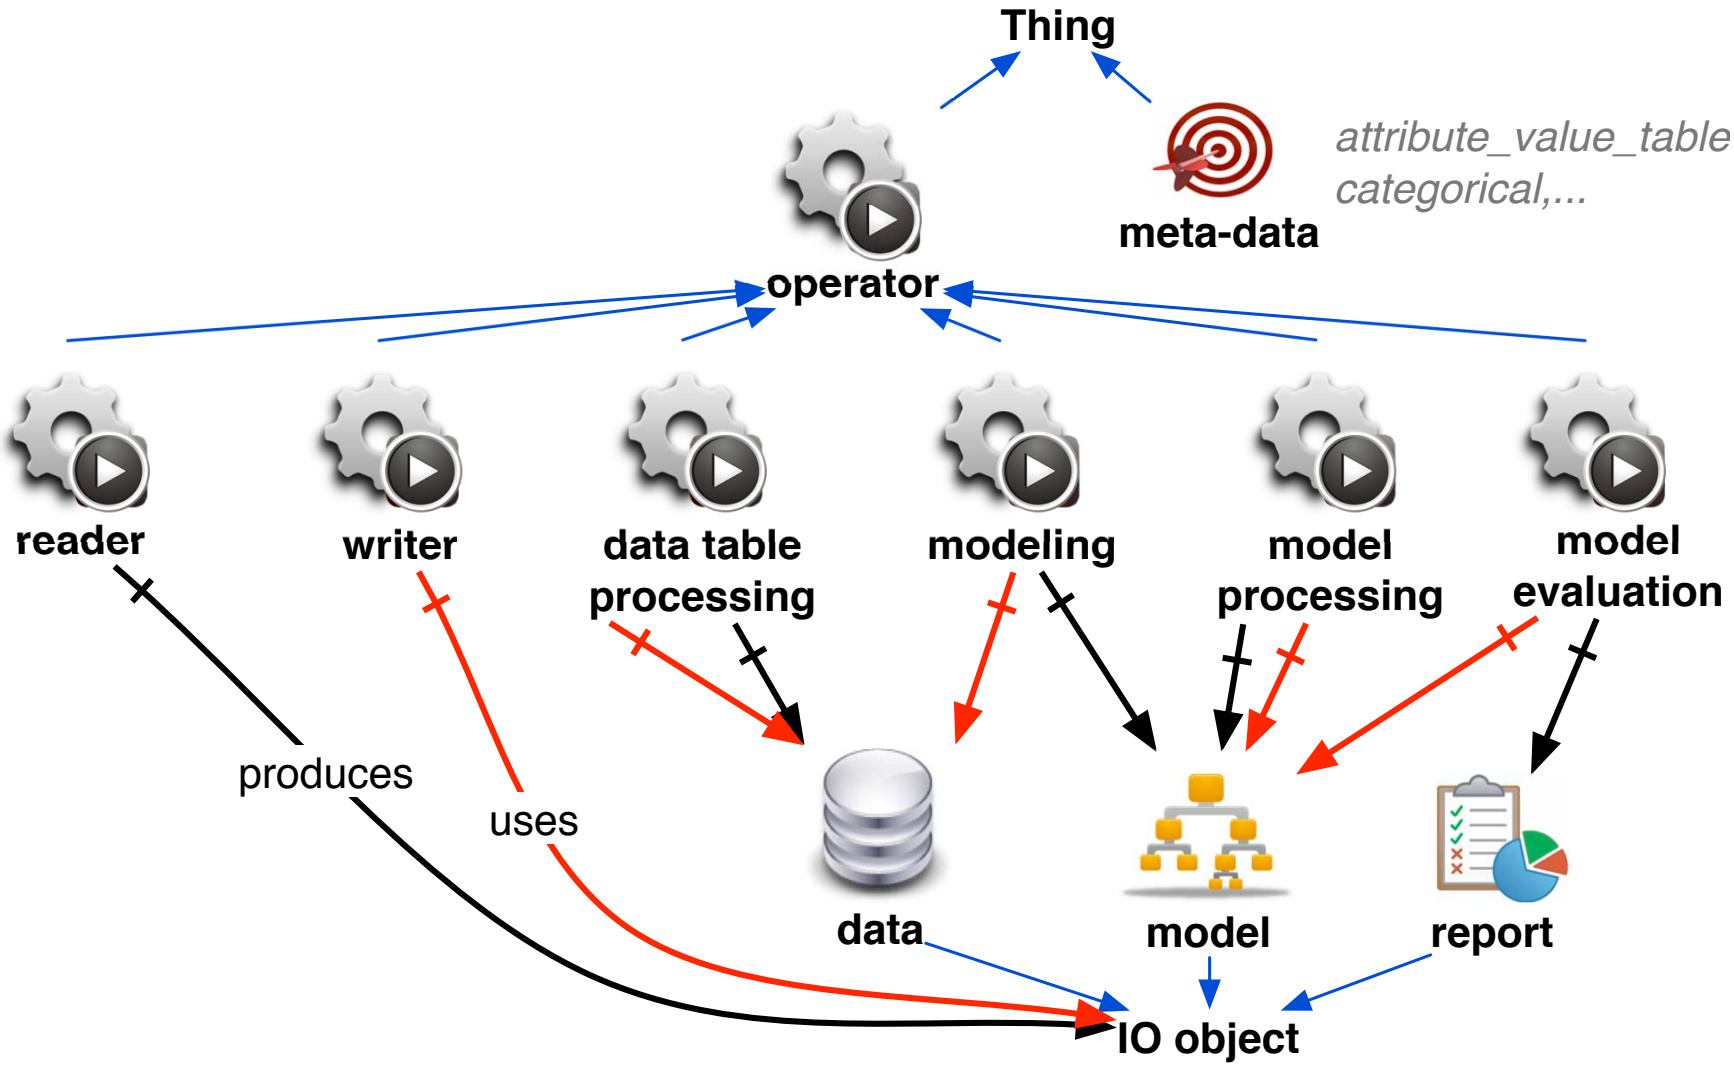

Ontology recap:
DMWF (Kietz et al., ’09)

- Reason about KD operators: in/outputs, conditions/effects (SWRL rules)

”RapidMiner.ID3”:
Superclass:
ClassificationLearning and (uses exactly 1 AttributeValueDataTable) and
(produces exactly 1 Model) and
(simpleParameter1(name=”minimal size for split”) exactly 1 integer) and
(simpleParameter2(name=”minimal leaf size”) exactly 1 integer) ...
Condition:
(**AttributeValueDataTable** and **MissingValueFreeData** and
(**inputAttribute only (hasAttributeType only Categorial**)) and
(**targetAttribute exactly 1** (hasAttributeType only Categorial)))(?D),
noOfRecords(?D,?Size), **?P1 is ?Size / 100**
→ uses(this,?D), simpleParameter2(this,?P1)
Effect:
uses(this,?D), hasFormat(?D,?F), inputAttribute(?D,?IA),targetAttribute(?D,?TA),
→ **new(?M,?D)**, **DecisionTree(?M)**, produces(this,?M), hasFormat(?M,?F),
inputAttribute(?M,?IA),predictedAttribute(?M,?TA),

Ontology recap:
EXPO (Soldatova and King, ’06)

- Make goal and structure of scientific experiments more explicit

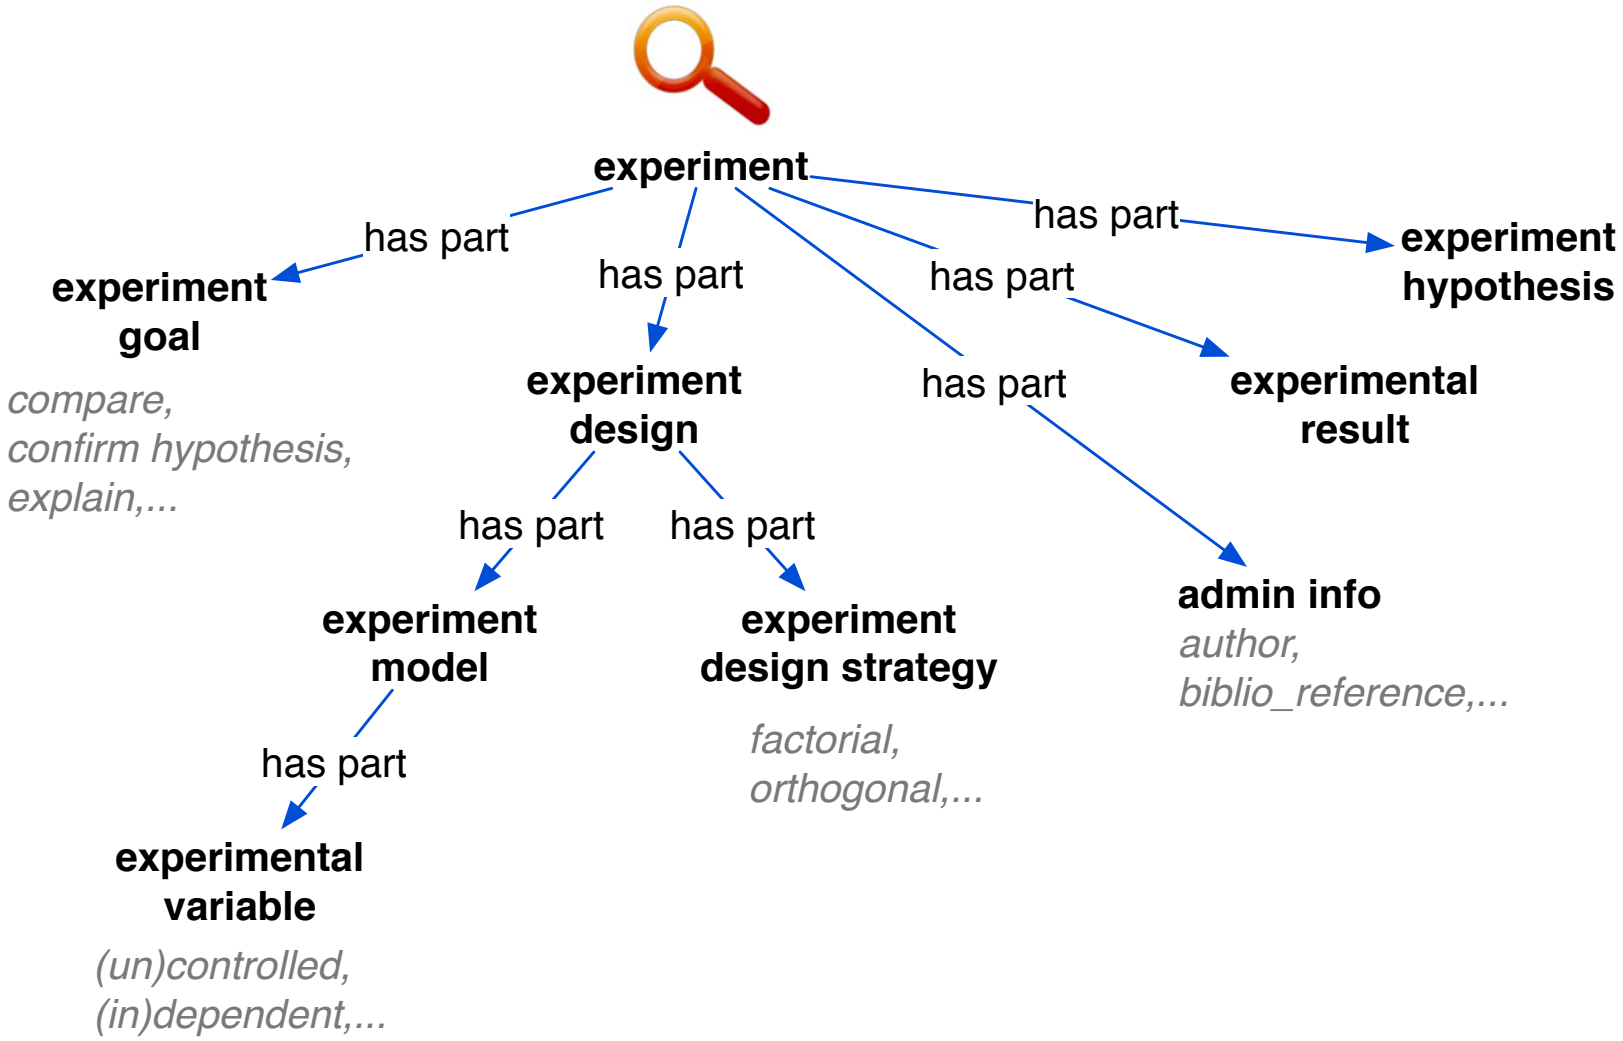

Ontology recap:
EXPO (Soldatova and King, ’06)

- Make goal and structure of scientific experiments more explicit

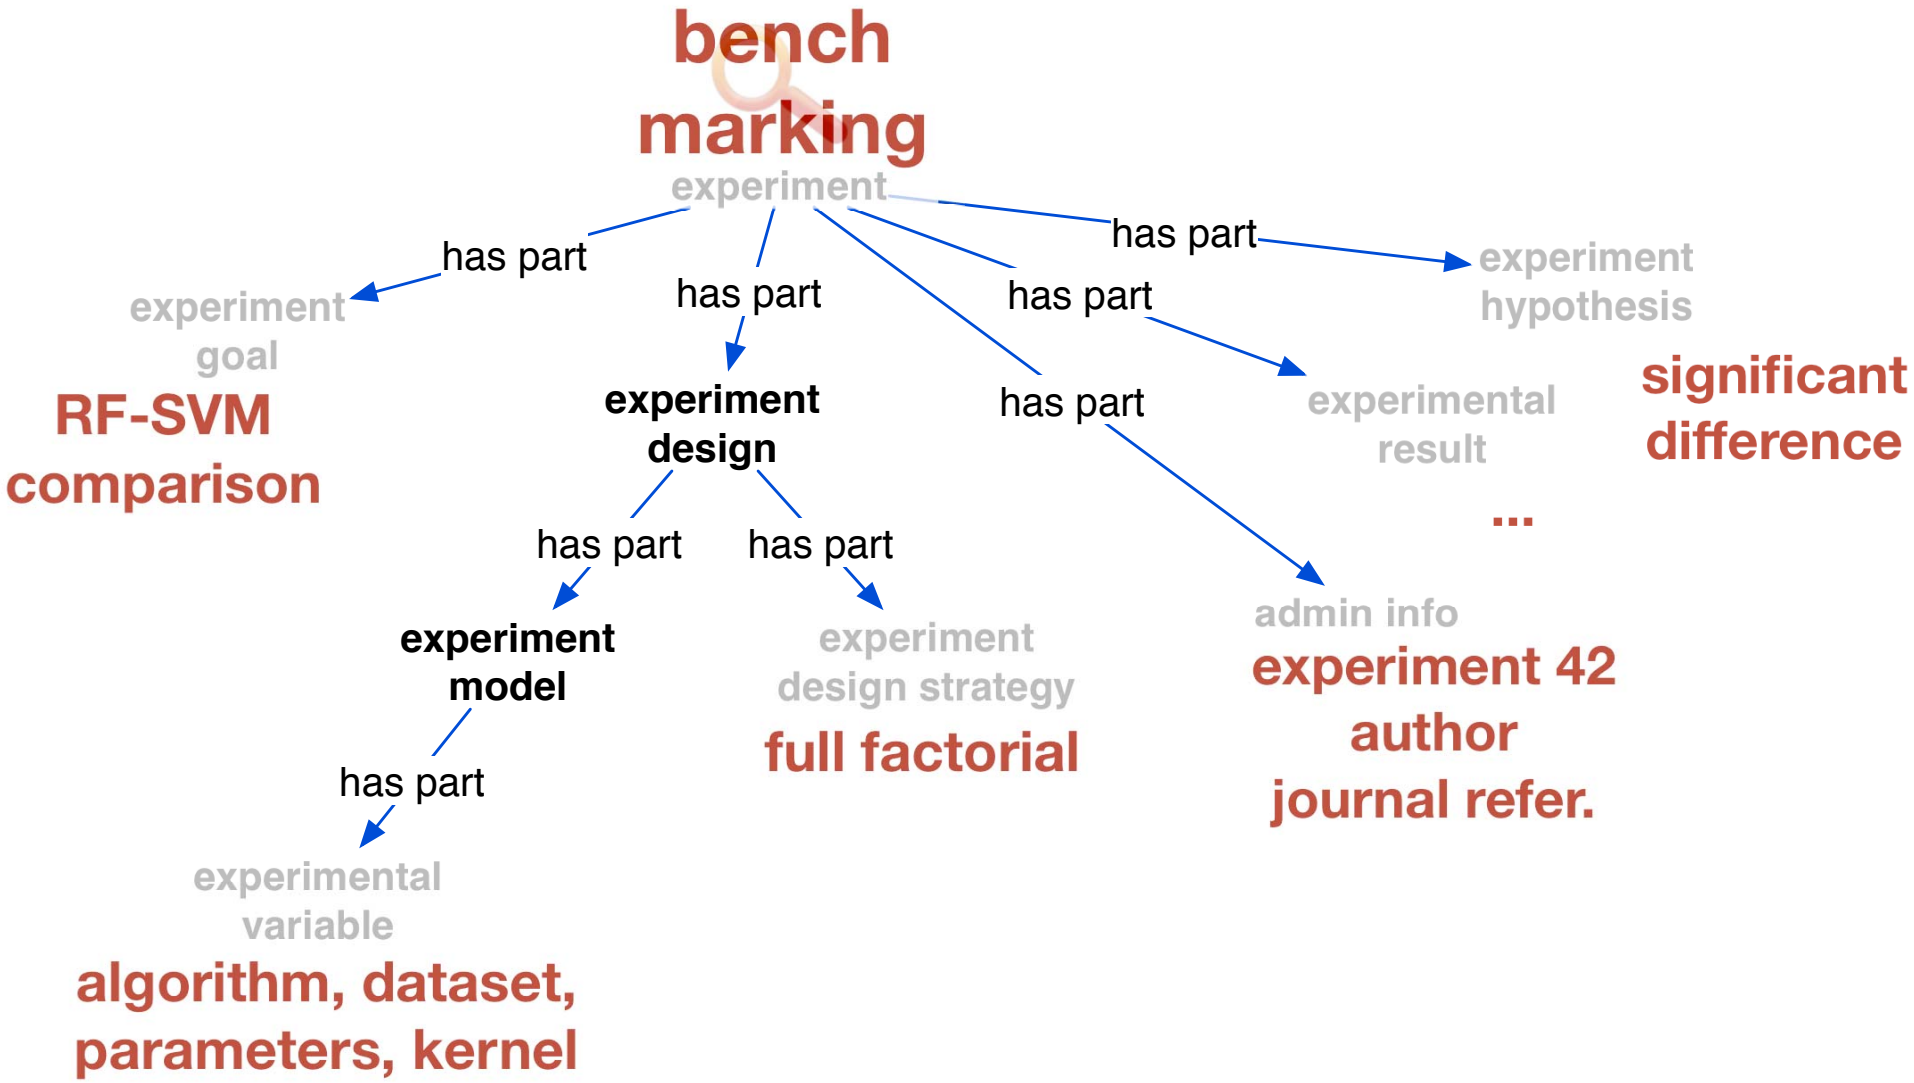

Exposé

an ontology for data mining experimentation

Context

- Giant, public database(s) of data mining experiments
- We need:
  - Common language to share experiments (through DM tools)
  - Intuitive ways to store and query experimental results
- We want:
  - Interoperable ontology: OntoDM for top-level, DMOP for detailed properties of learning algorithms
  - Driven by actual experiments submitted to database
    - New algorithms -> ideally, described by author
    - Instances automatically extracted from database

Problem 1: Experiments

What is a machine learning experiment?
What do we need to know about it?

Exposé: Experiments

hp: has participant
hd: has description

experiment workflow

KD workflow

Workflow:
has inputs, outputs, operators (participants)

Exposé: Experiments

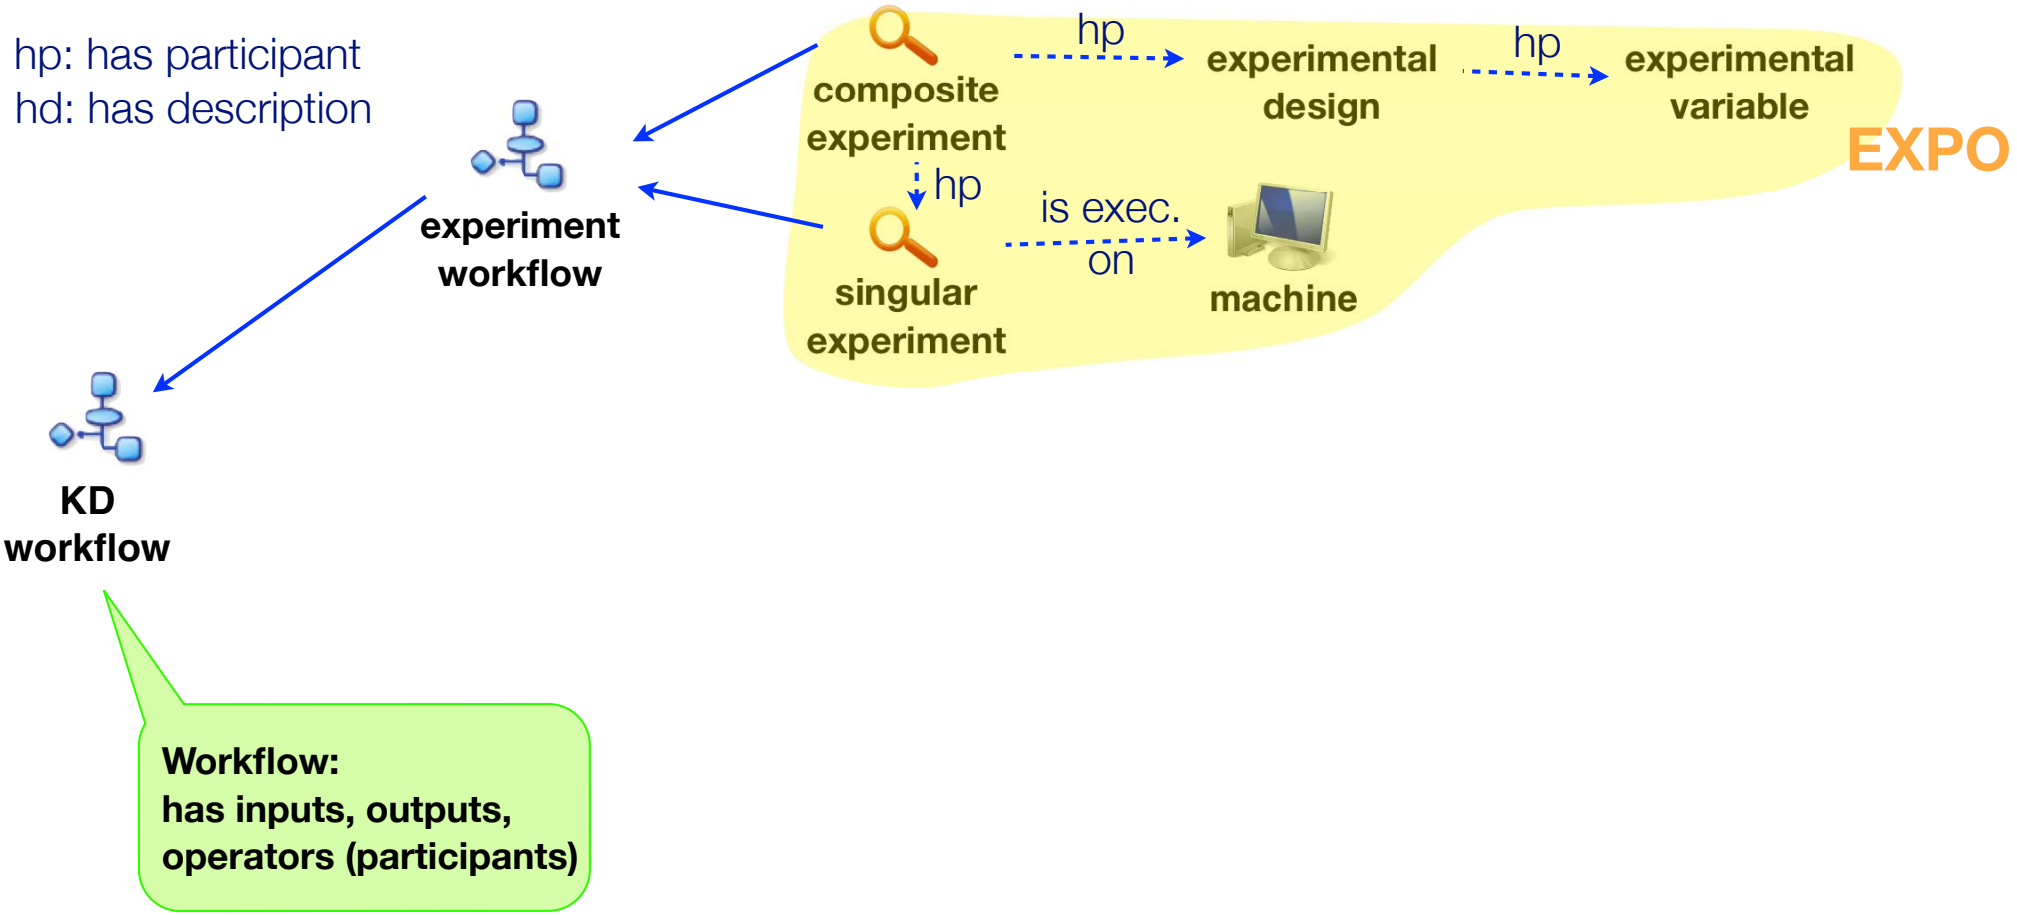

Exposé: Experiments

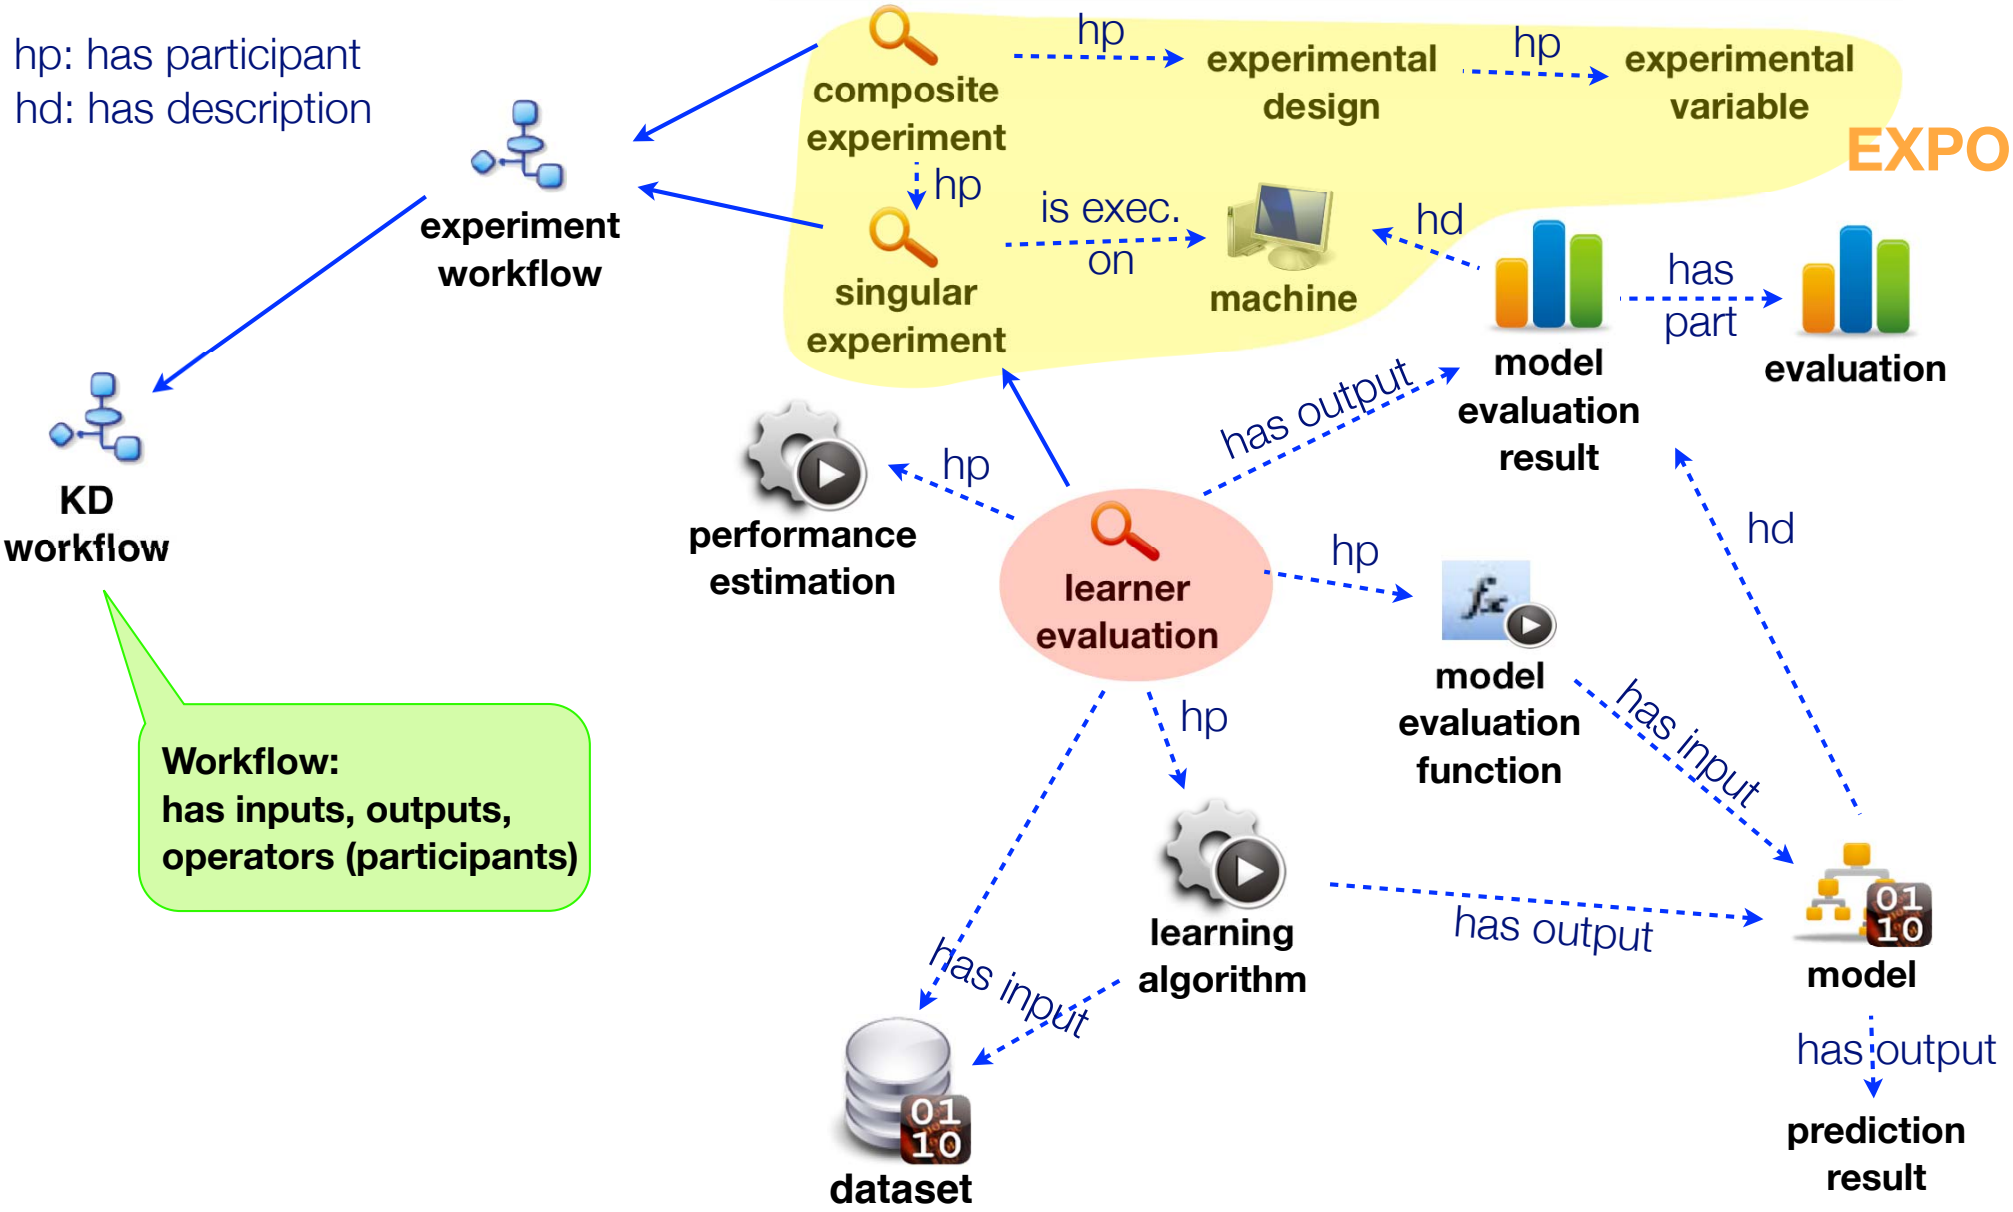

Problem 2: Algorithms

When talking about an algorithm, what is meant?

General algorithm?
Specific implementation? Which version?
When run, which parameters, components?

Exposé: Algorithms
Specification, implementation, application

hp: has participant
hd: has description
ico: is concretization of

Similar to OntoDM

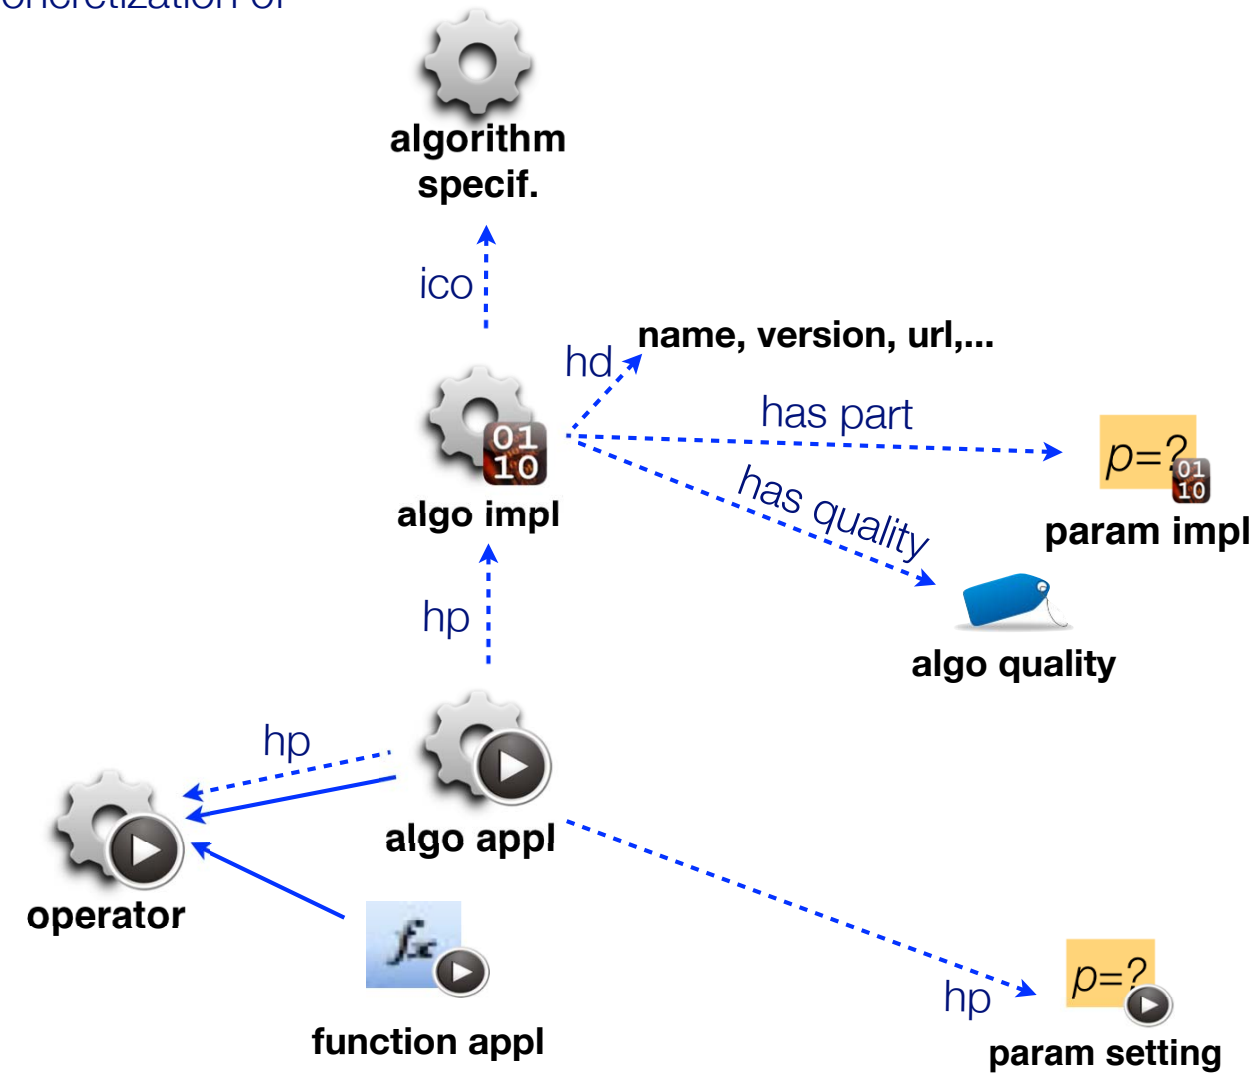

Same for functions and parameters

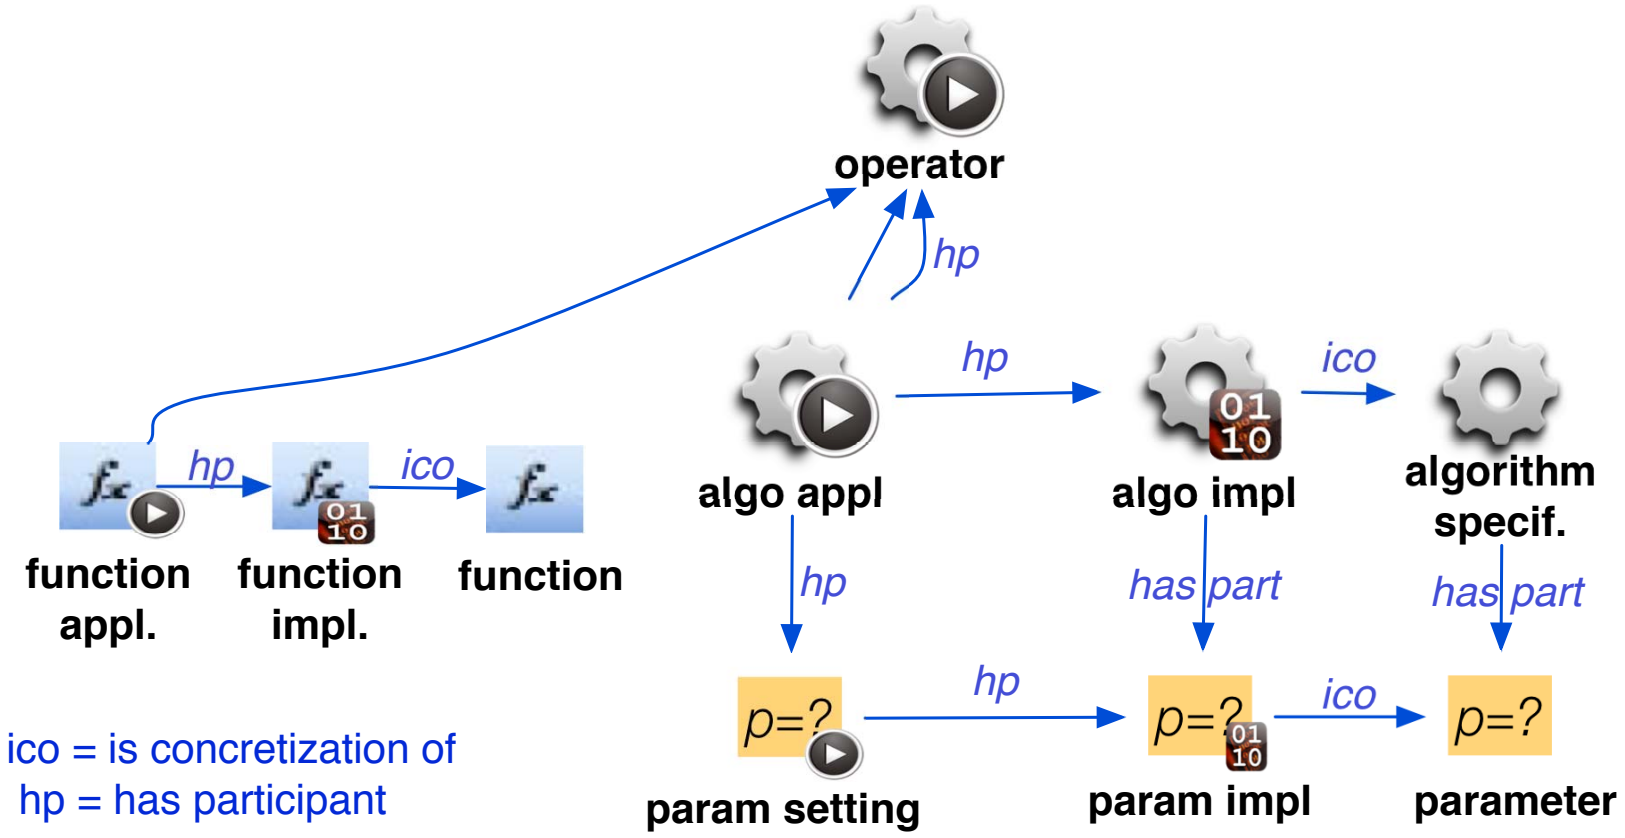

Problem 3: Algorithm composition

plug-in functions, kernels, other algorithms

such components play different roles
-> Agent-role pattern

Exposé: Algorithms

hp: has participant
hd: has description

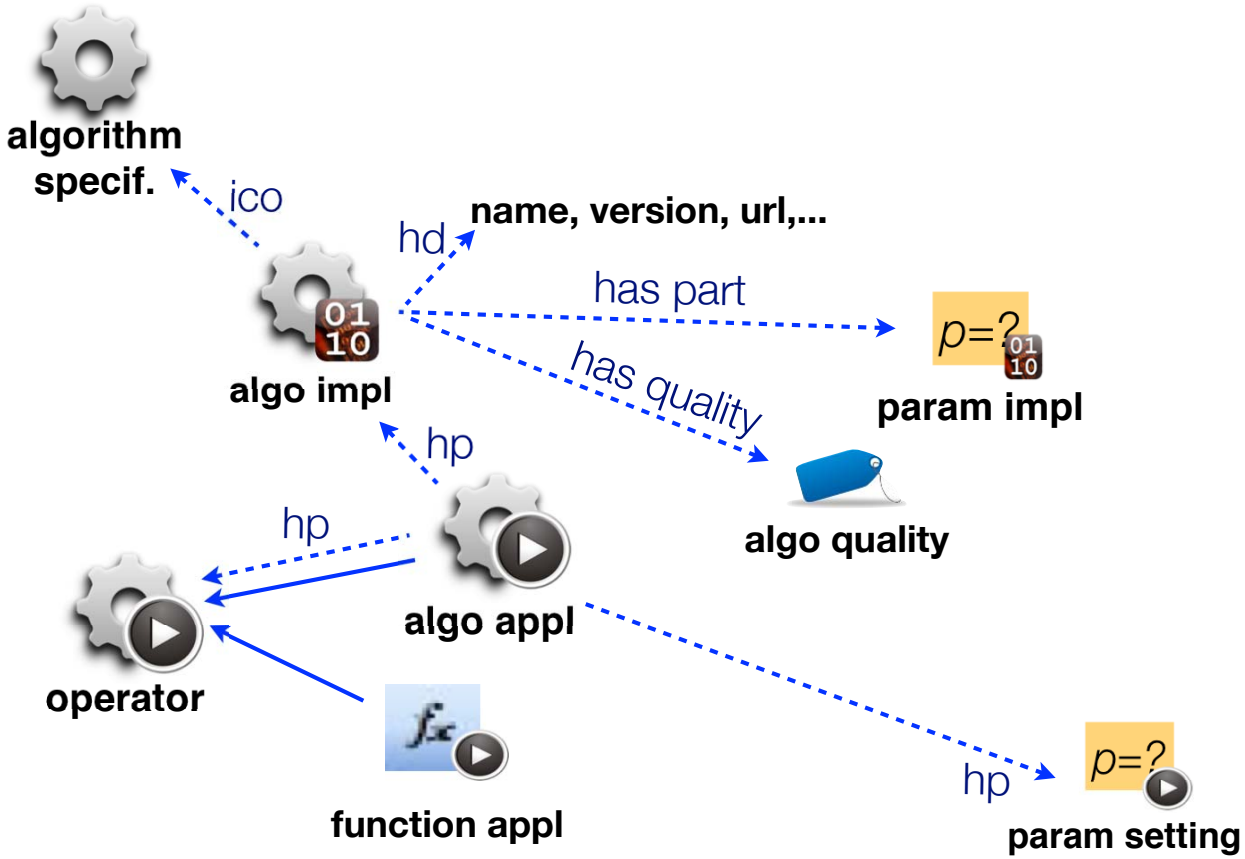

Exposé: Algorithms

hp: has participant
hd: has description

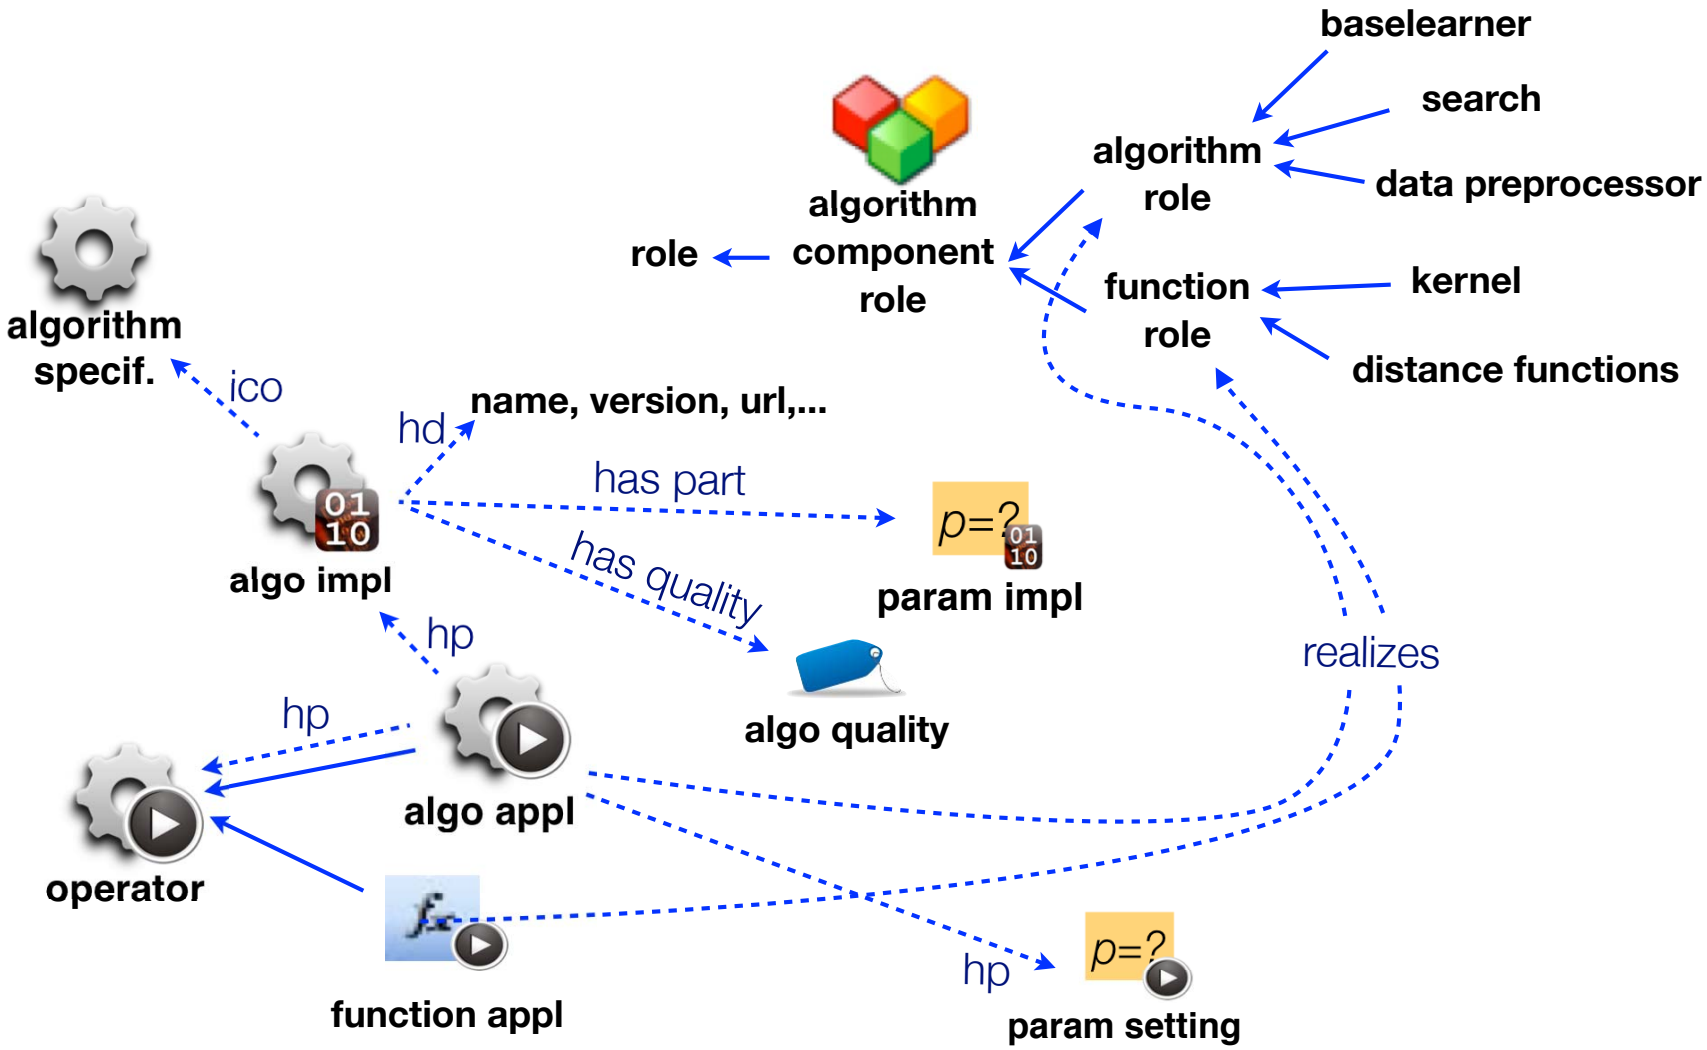

Exposé: Algorithms

hp: has participant
hd: has description

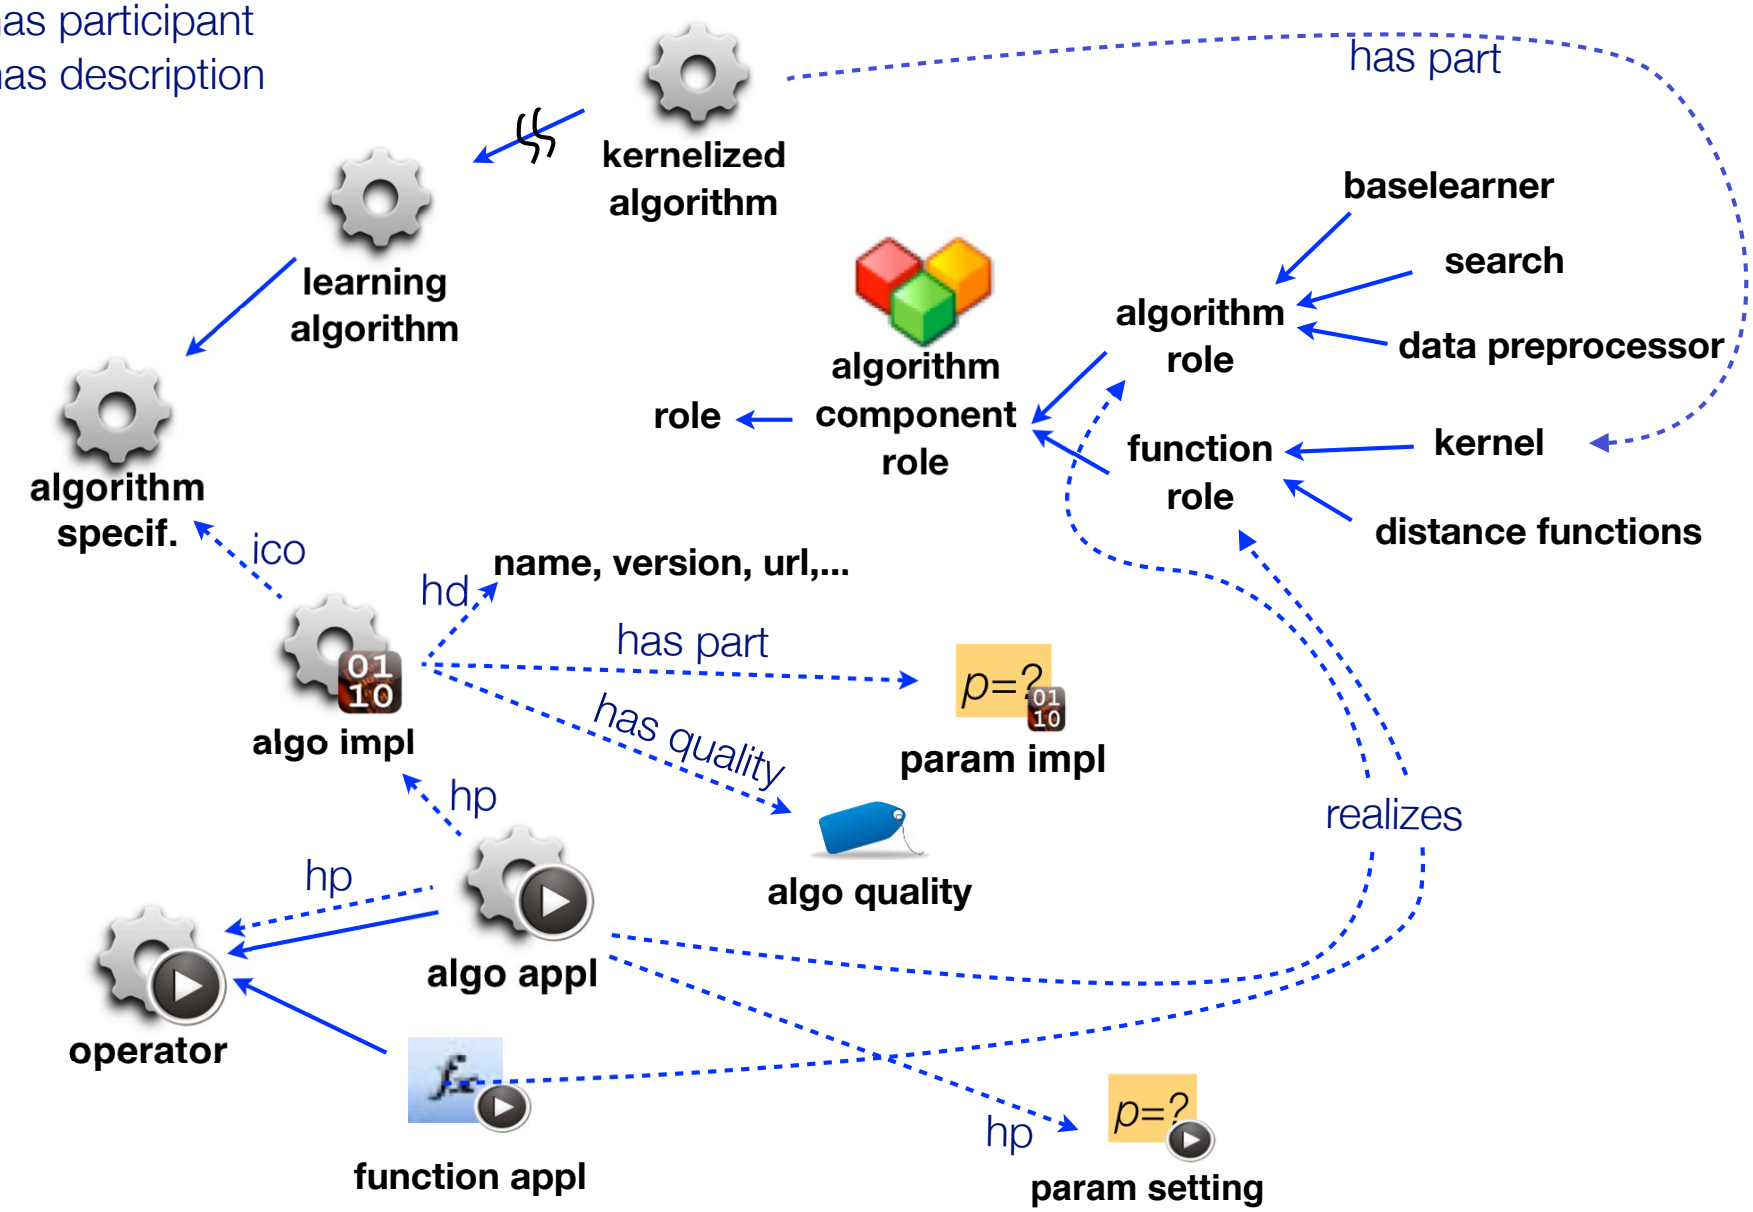

Problem 4: Workflows

Inputs, outputs, operators

Hierarchical: workflows within workflows

Reuse, parameterize common workflows, e.g. k-fold CV

Exposé: workflows

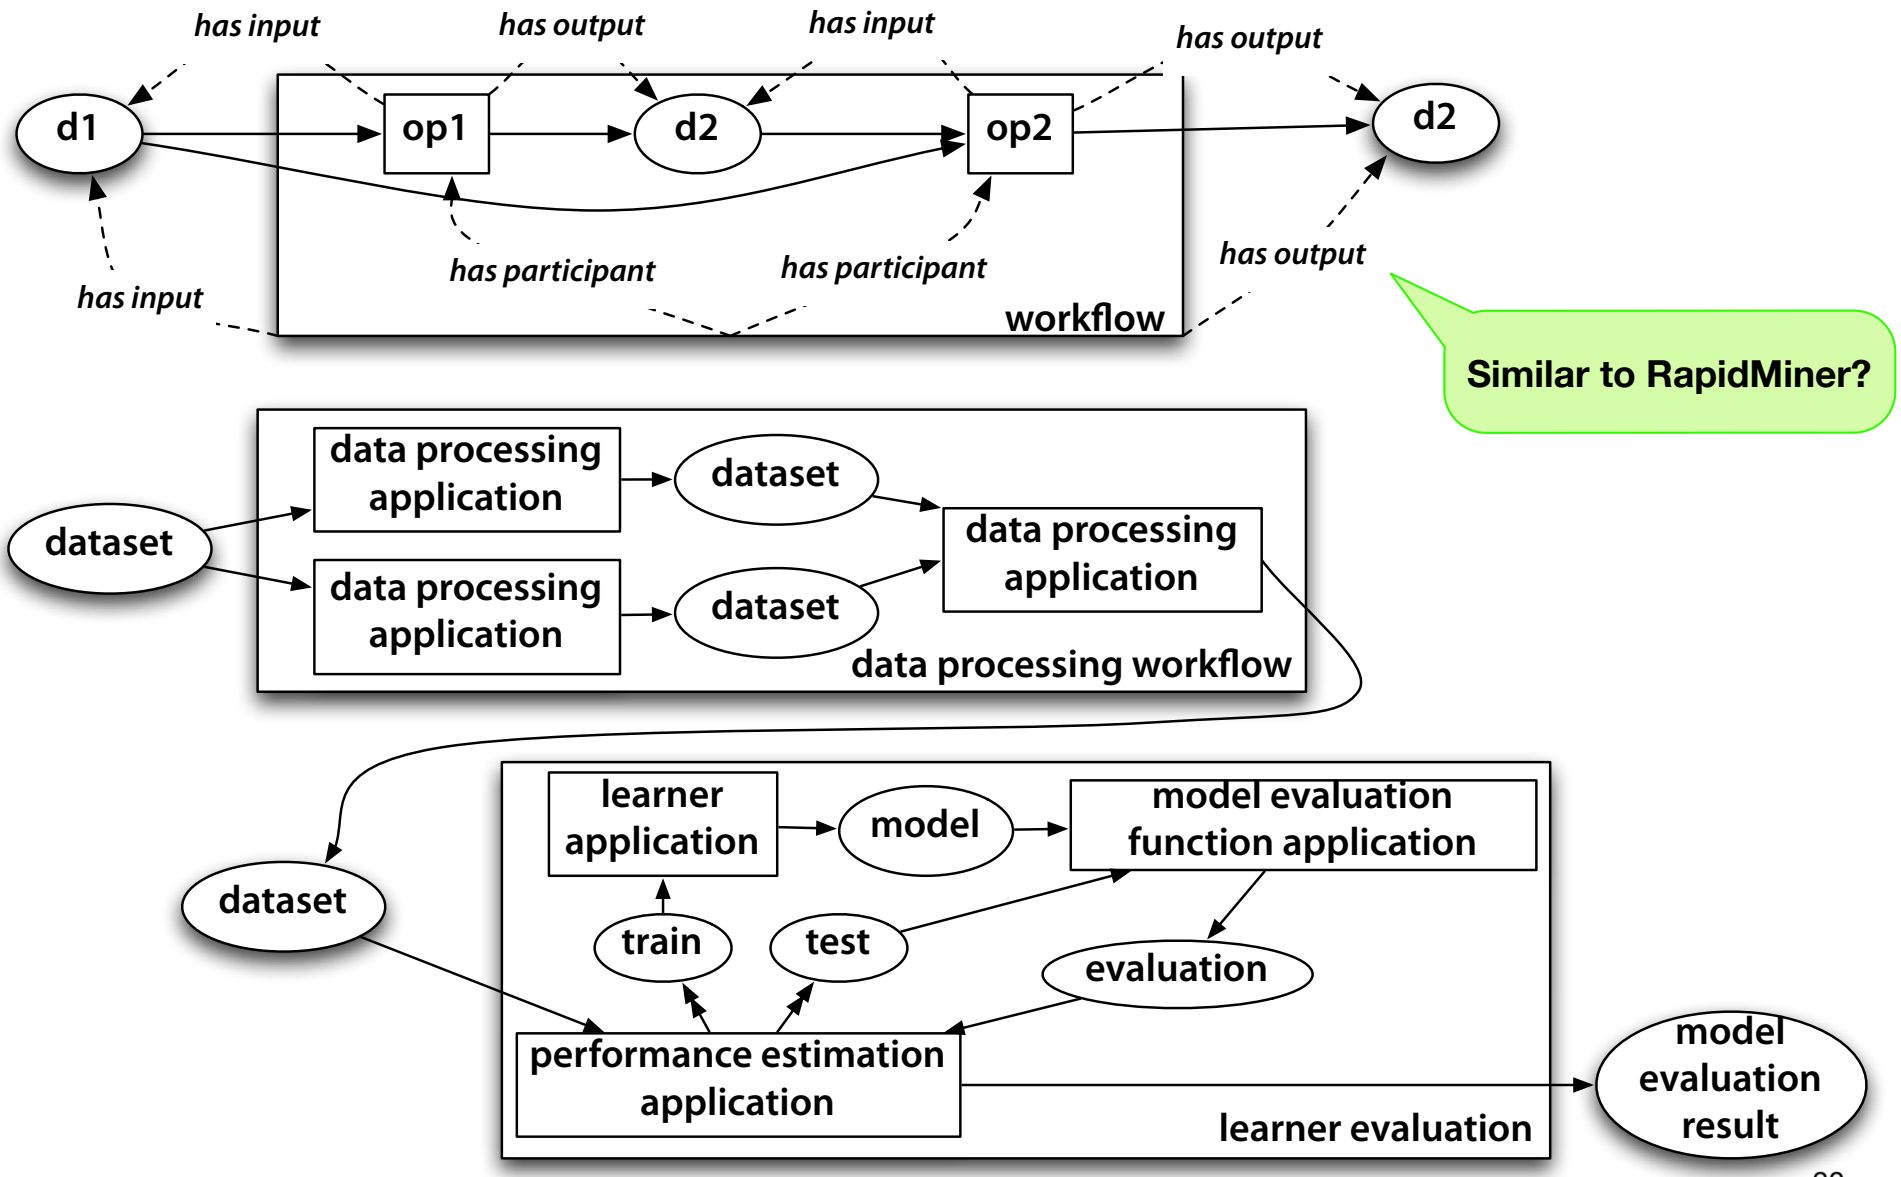

Problem 5: Reuse

How can we make maximal use of existing ontologies?

OBI: top-level
OntoDM: top-level DM concepts
DMO: operators, learning mechanisms

BFO: accepted top-level classes

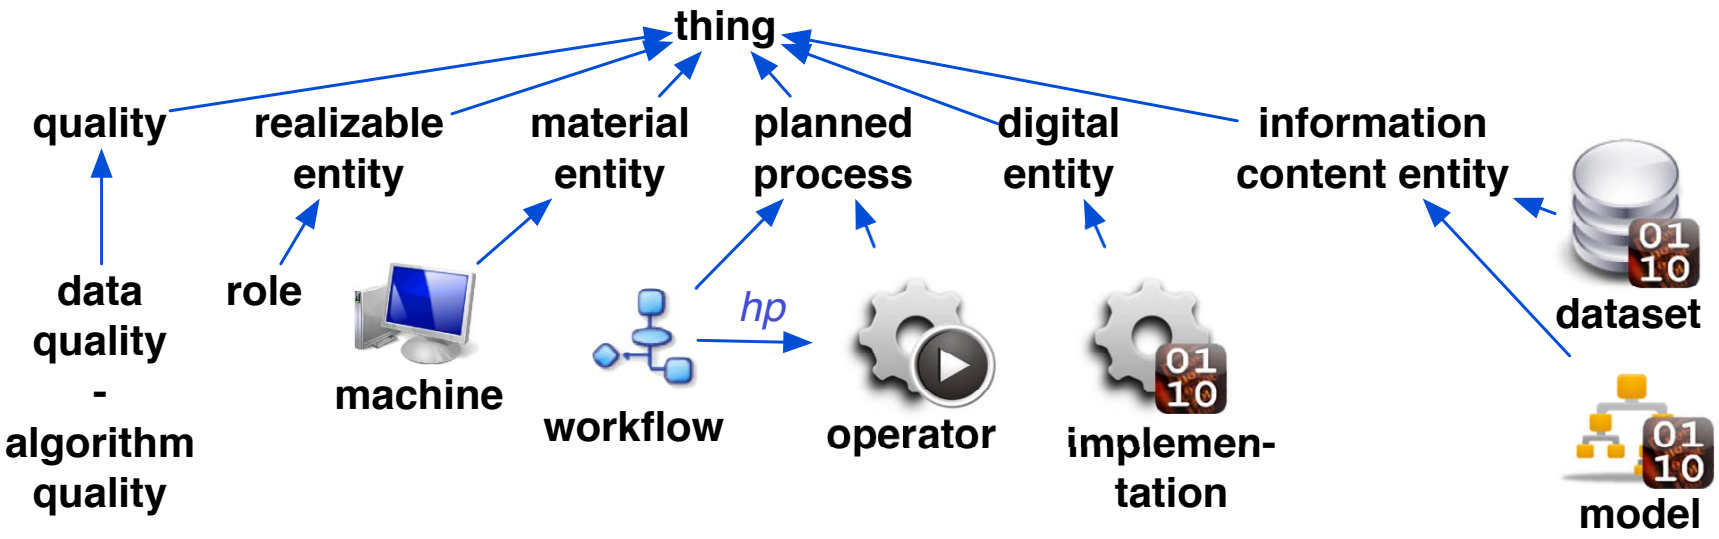

BFO: accepted top-level classes

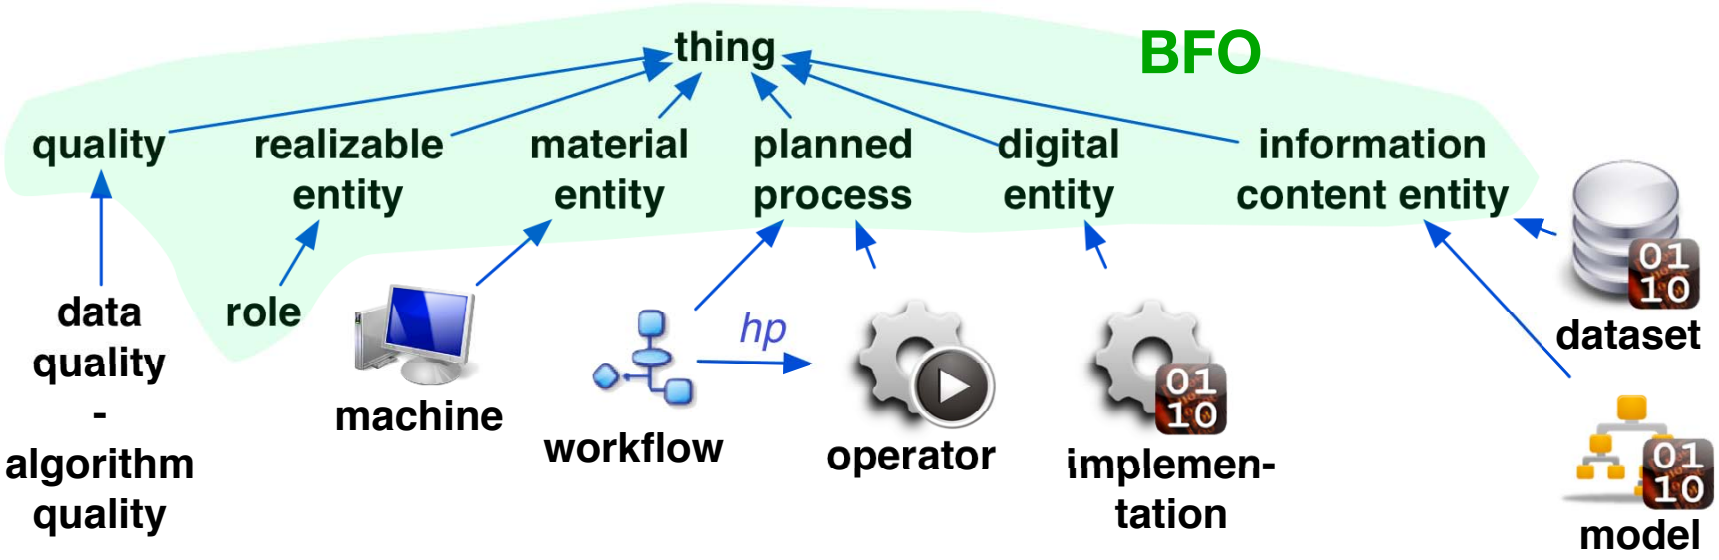

OntoDM: top-level DM concepts

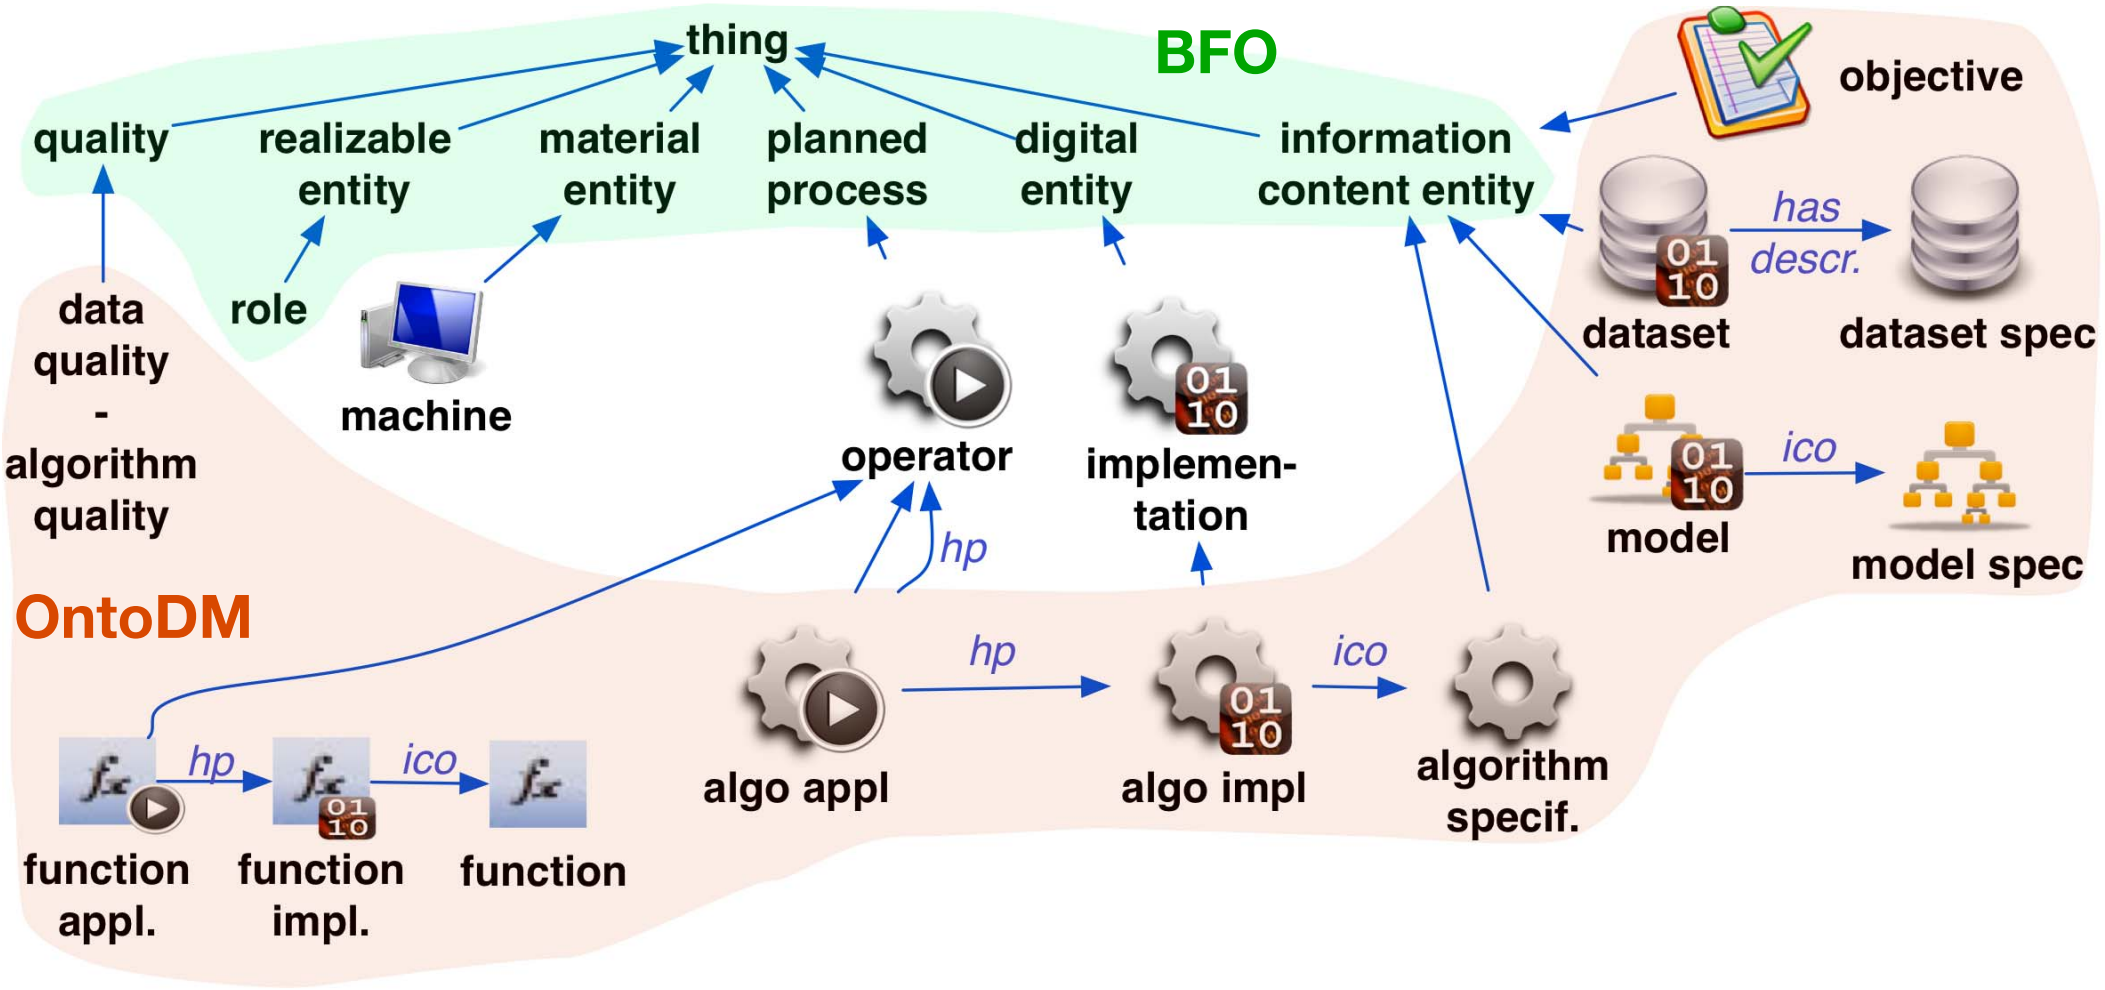

ico = is concretization of
hp = has participant

DMO: operators, learning mechanisms

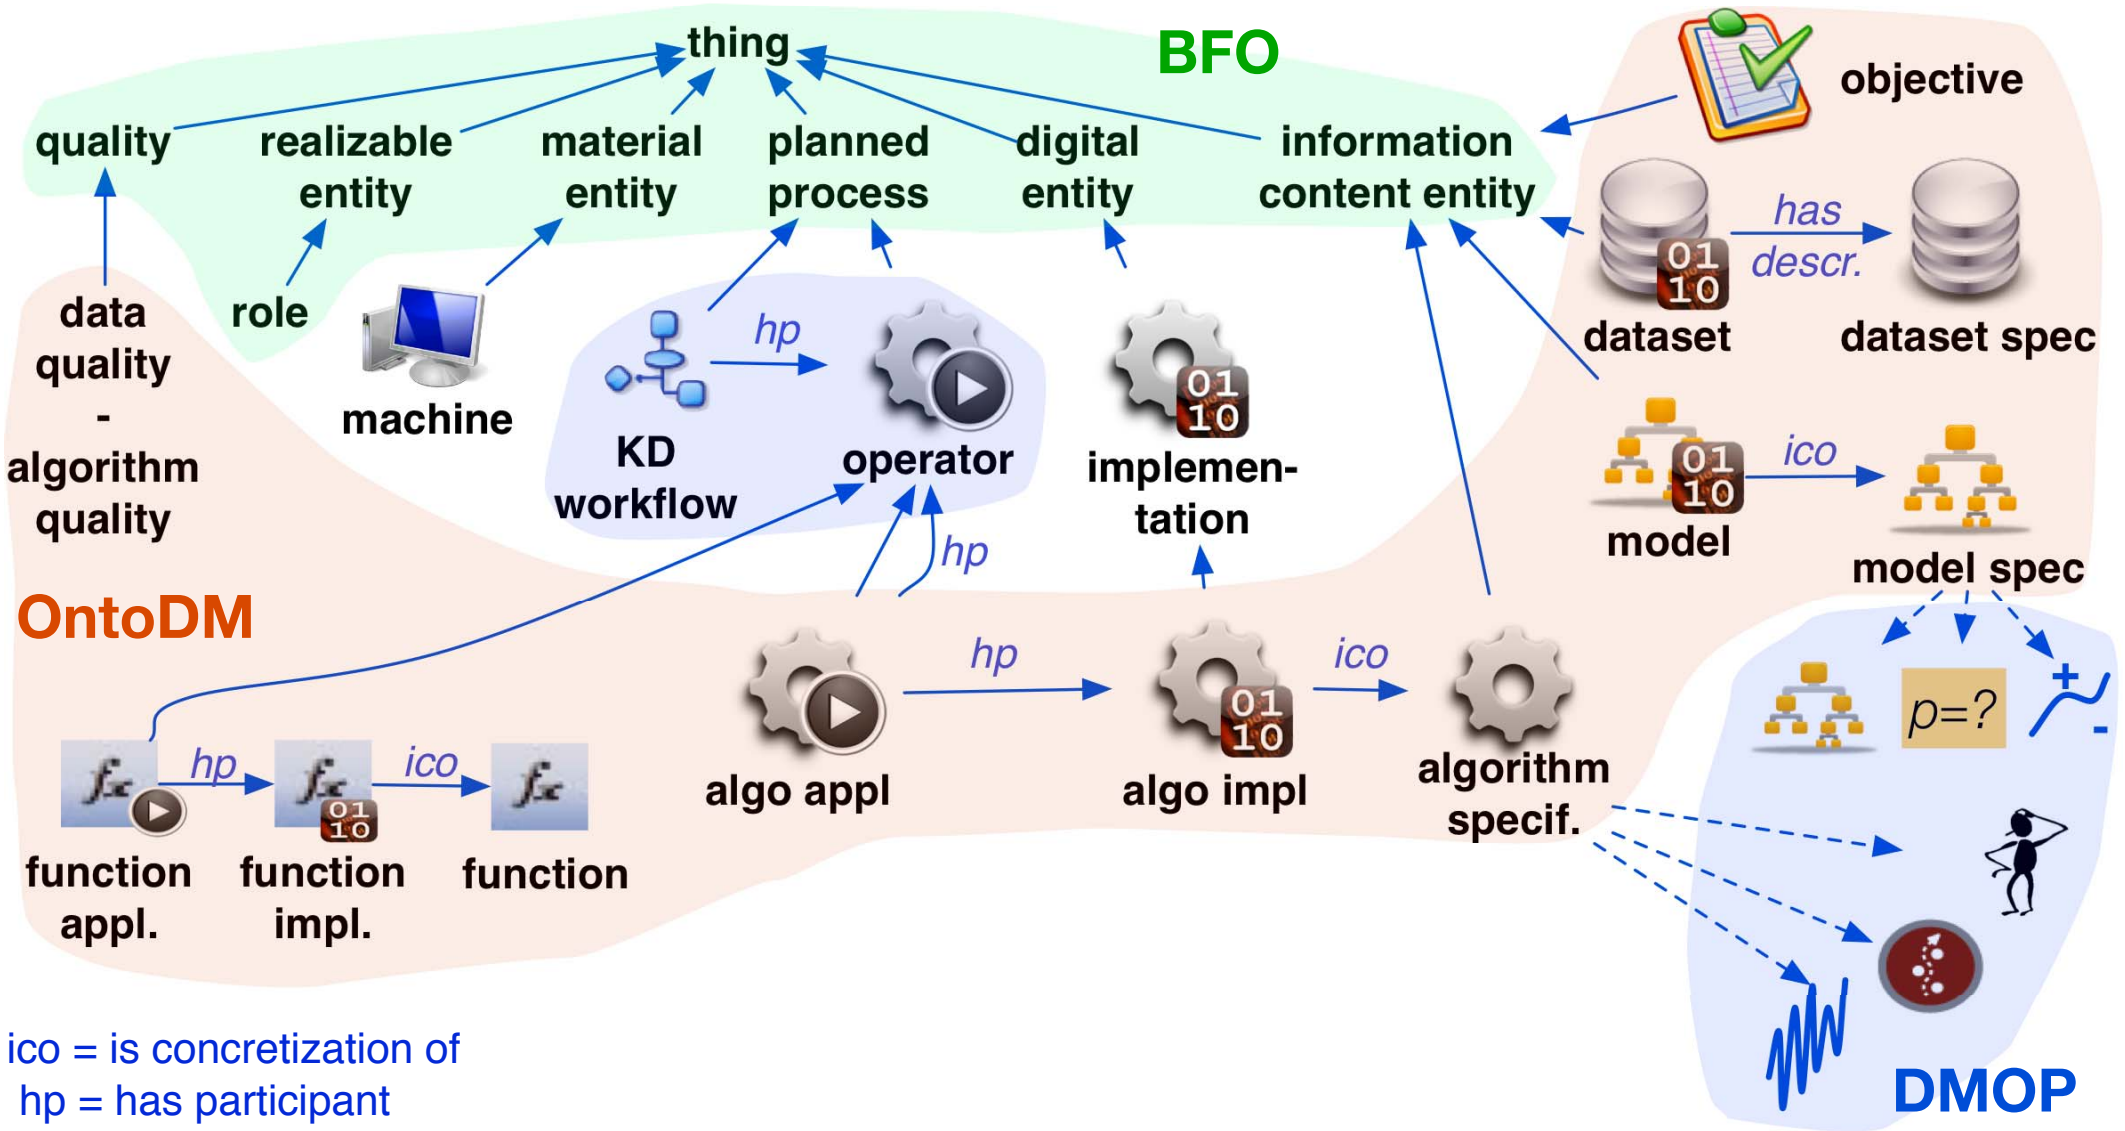

ico = is concretization of
hp = has participant

Exposé: top level classes

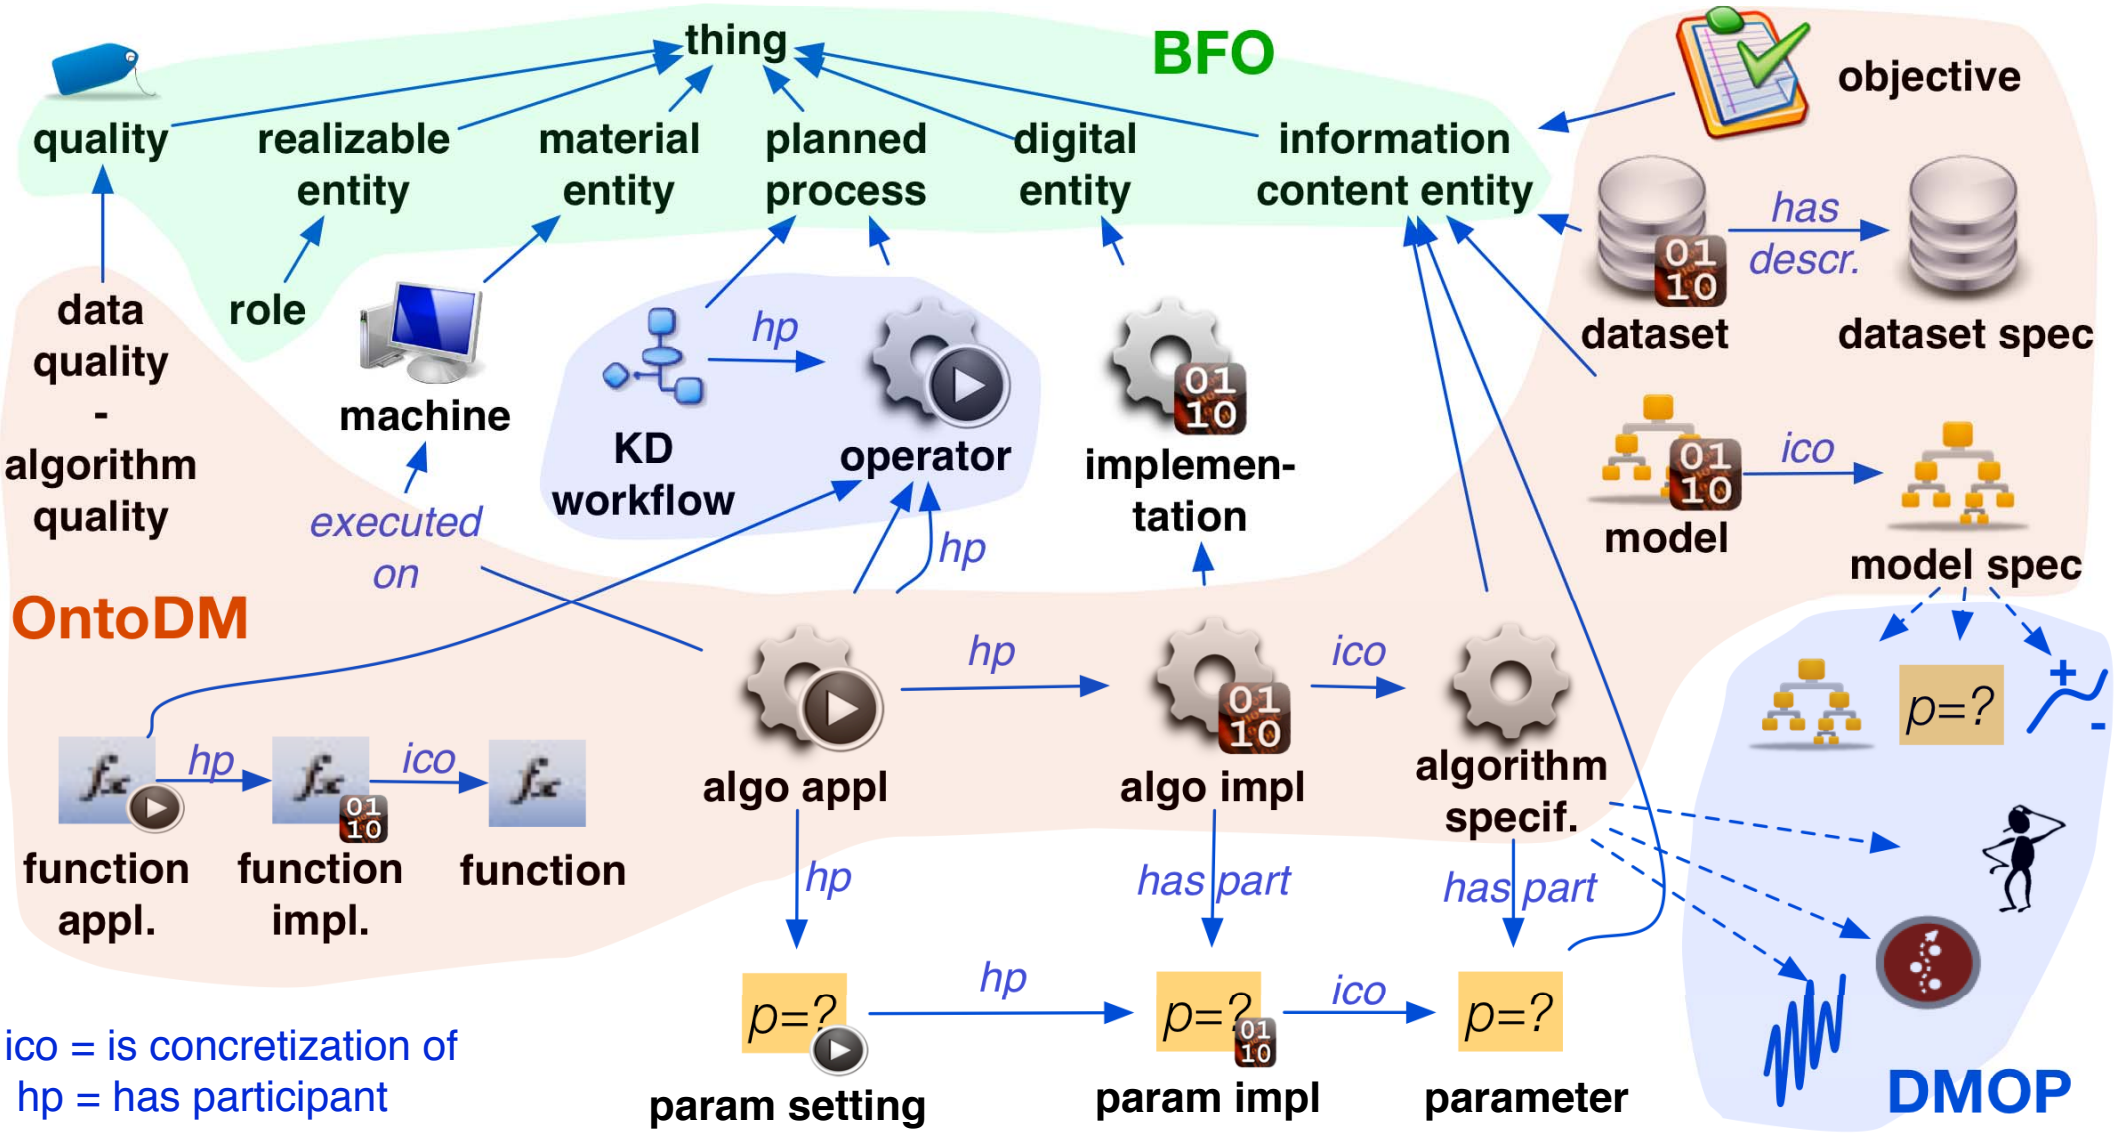

Other aspects

Datasets

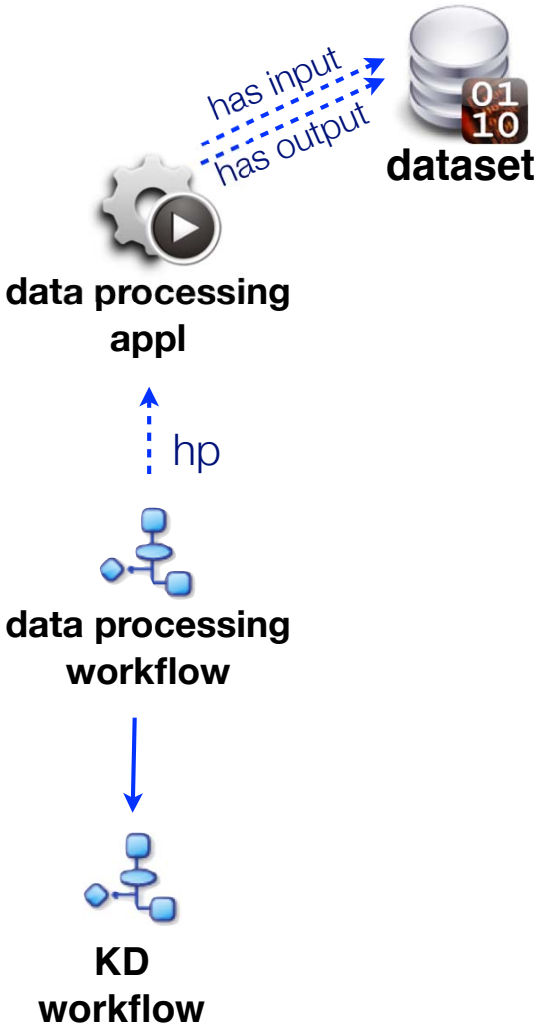
has input
has output
dataset
data processing
appl
hp
data processing
workflow
KD
workflow

Datasets

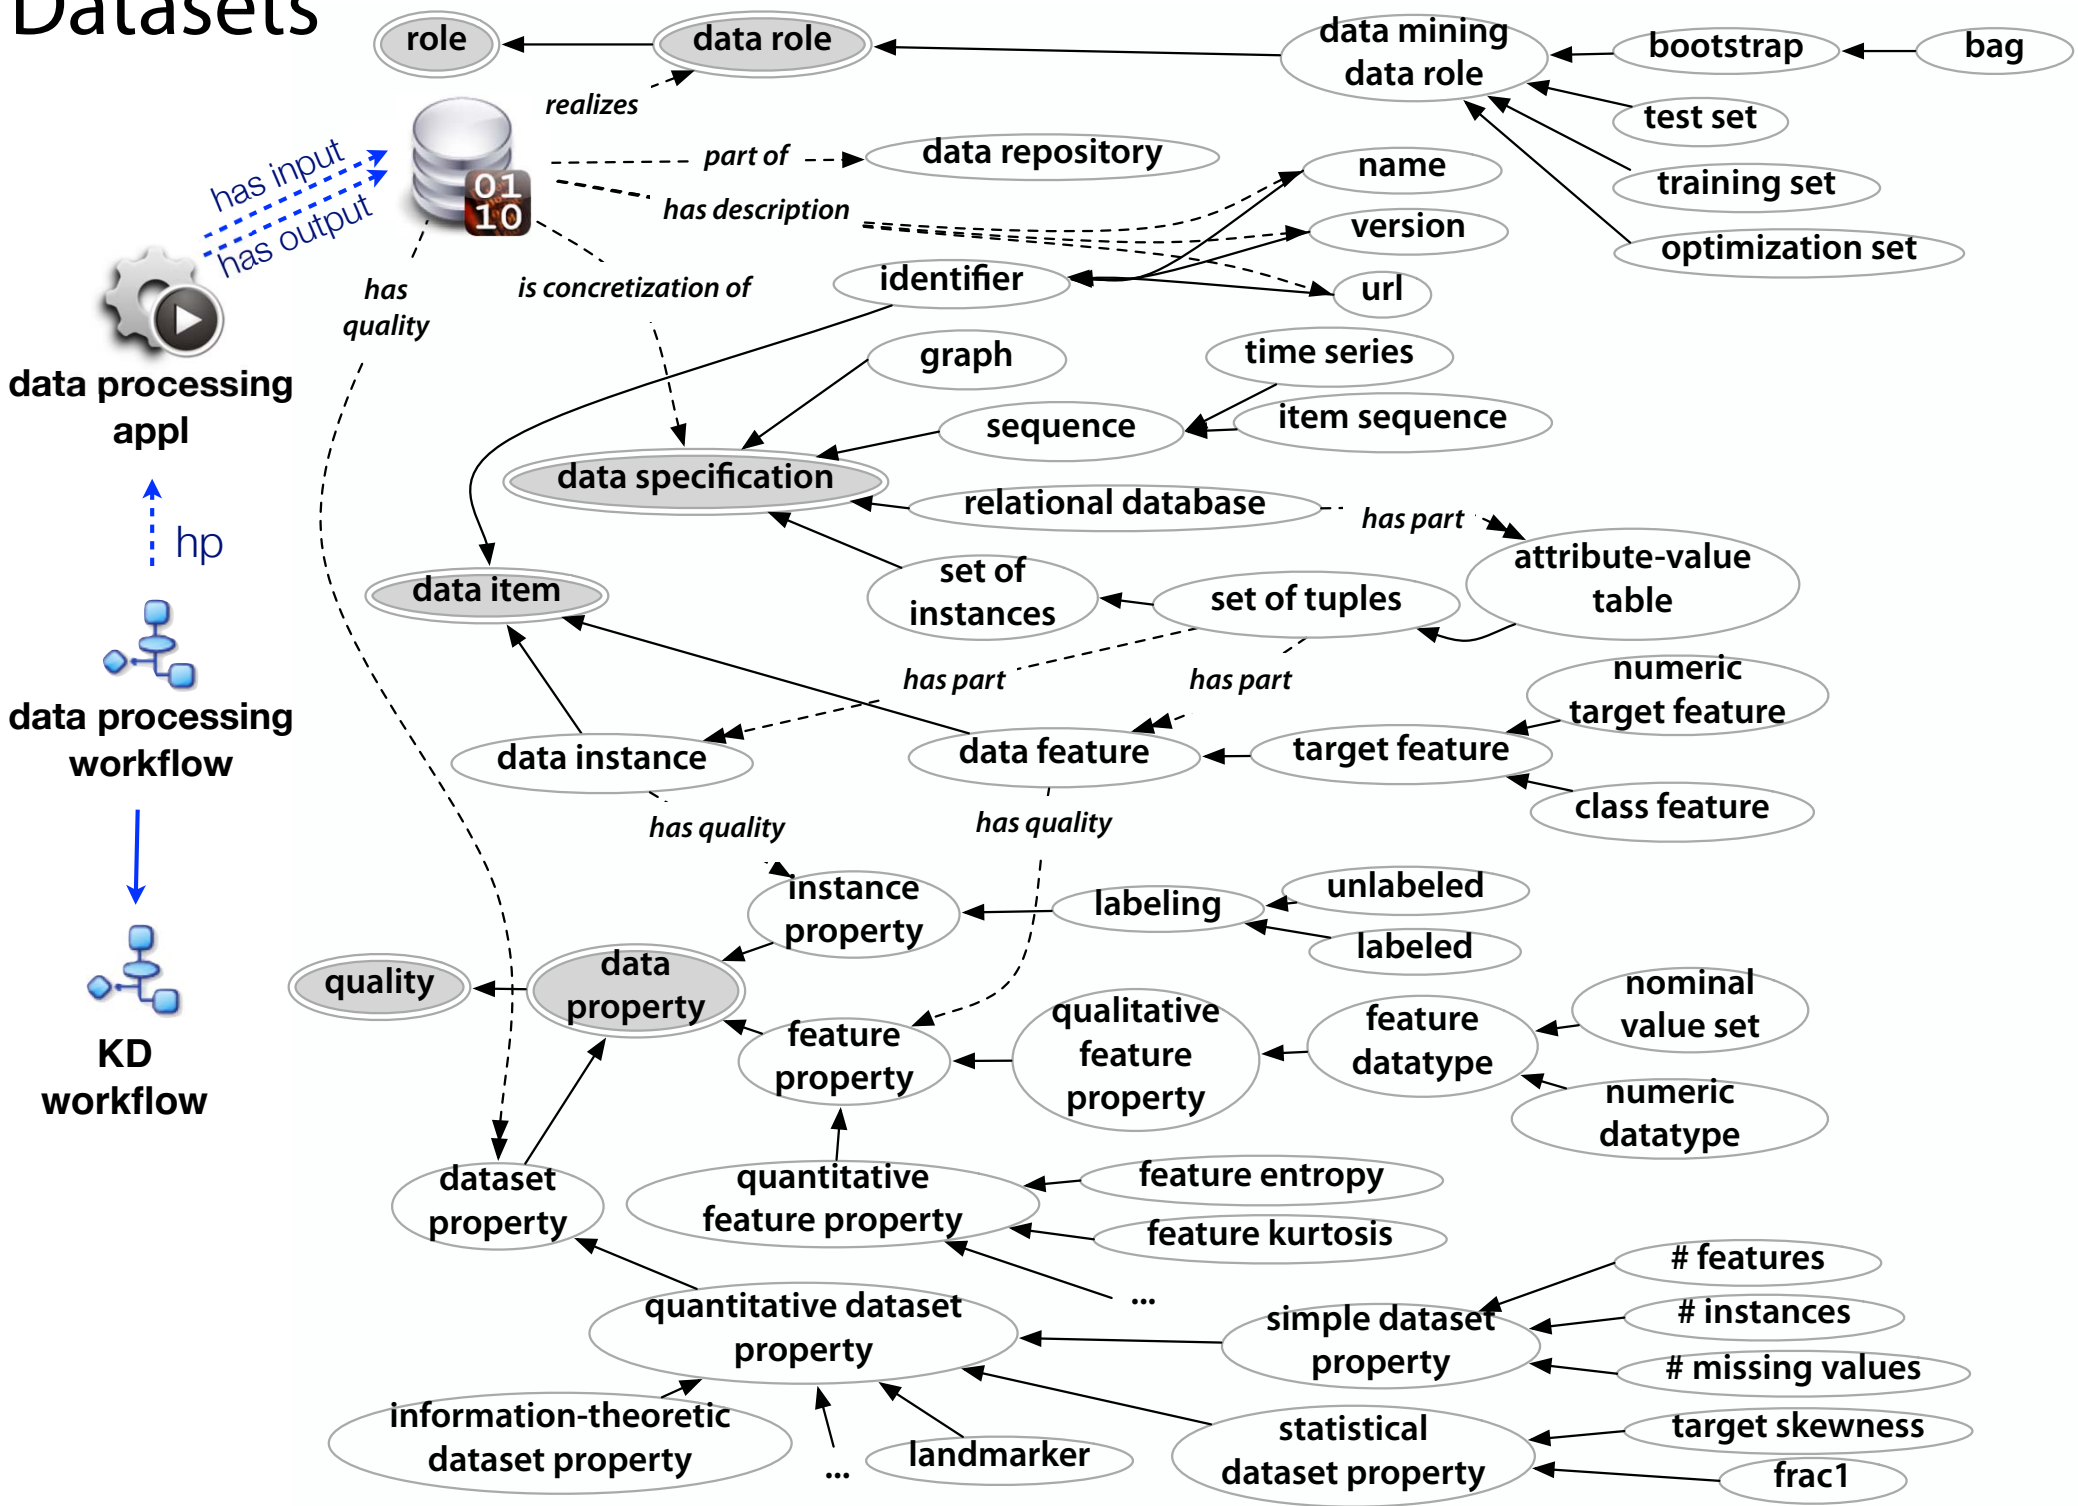
role
data role
data mining data role
bootstrap
bag
test set
training set
optimization set
realizes
part of
data repository
has description
name
version
url
identifier
has input
has output
has quality
is concretization of
data processing appl
hp
data processing workflow
KD workflow
graph
time series
sequence
item sequence
data specification
relational database
has part
attribute-value table
data item
set of instances
set of tuples
has part
has part
numeric target feature
data instance
data feature
target feature
class feature
has quality
has quality
instance property
labeling
unlabeled
labeled
quality
data property
feature property
qualitative feature property
feature datatype
nominal value set
numeric datatype
dataset property
quantitative feature property
feature entropy
feature kurtosis
...
quantitative dataset property
simple dataset property
features
instances
missing values
information-theoretic dataset property
...
landmarker
statistical dataset property
target skewness
frac1

Evaluation

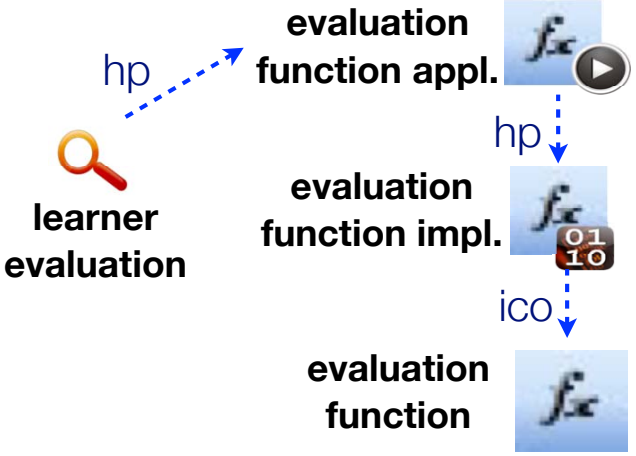
hp
evaluation
function appl.
hp
learner
evaluation
evaluation
function impl.
ico
evaluation
function

Evaluation

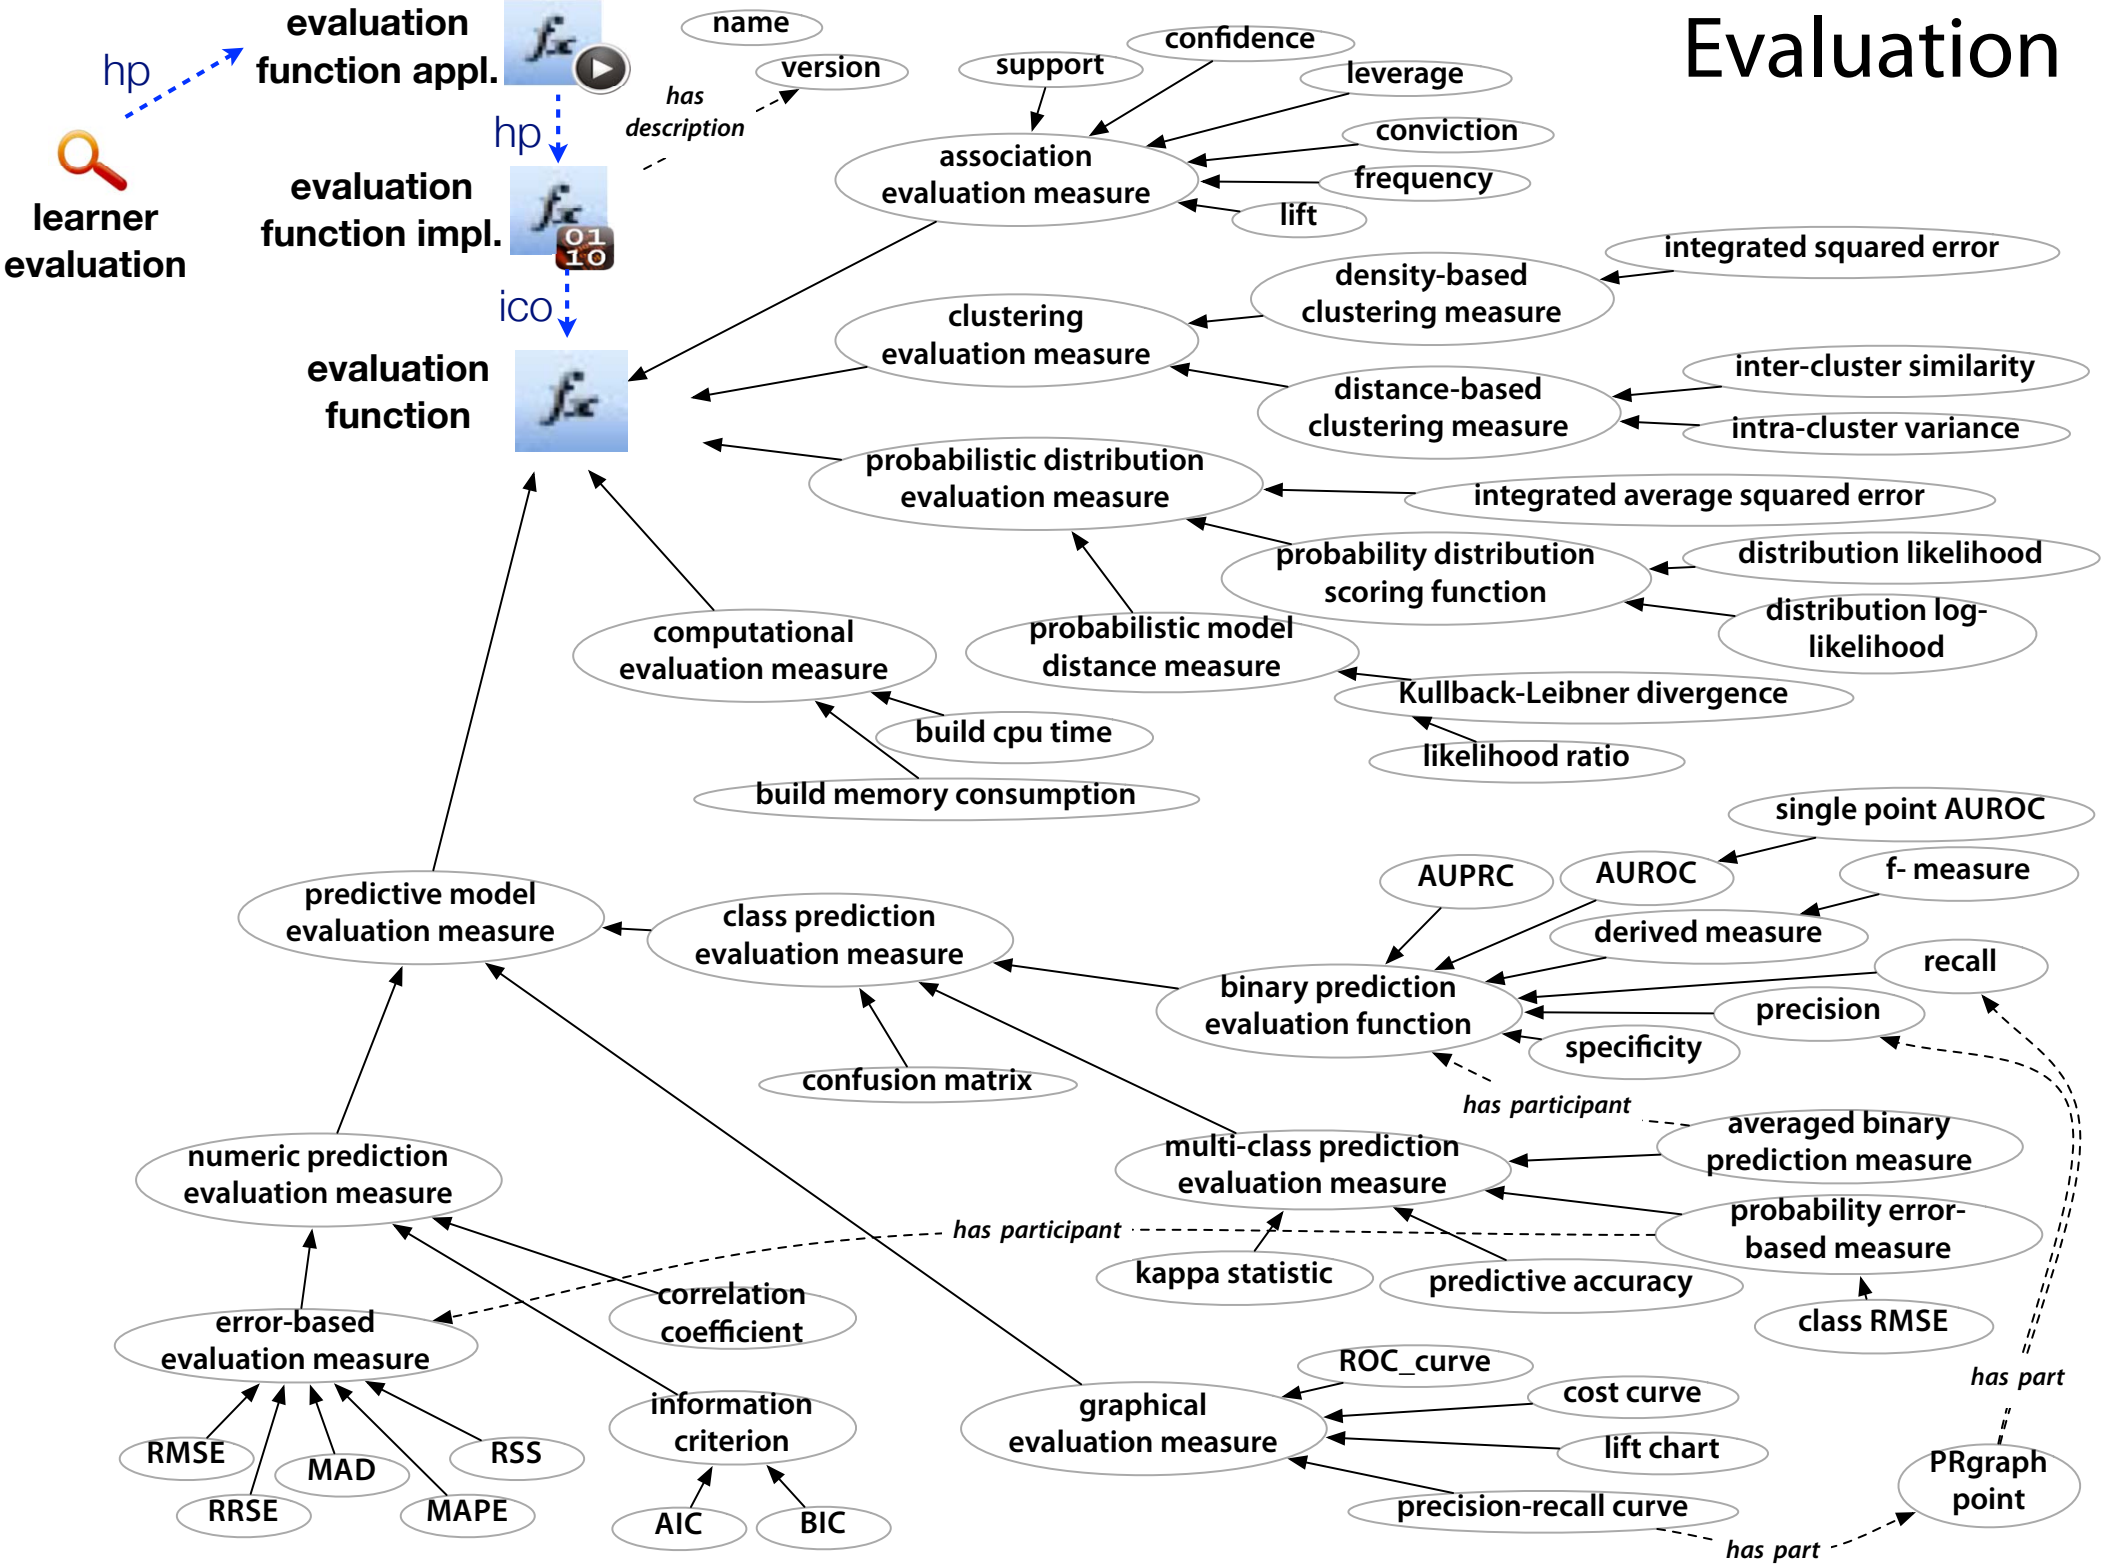
learner evaluation
hp
evaluation function appl.
hp
evaluation function impl.
ico
evaluation function
name
version
has description
support
confidence
leverage
conviction
frequency
lift
association evaluation measure
integrated squared error
density-based clustering measure
clustering evaluation measure
inter-cluster similarity
distance-based clustering measure
intra-cluster variance
probabilistic distribution evaluation measure
integrated average squared error
probability distribution scoring function
distribution likelihood
distribution log-likelihood
computational evaluation measure
probabilistic model distance measure
Kullback-Leibner divergence
likelihood ratio
build cpu time
build memory consumption
single point AUROC
AUPRC
AUROC
f- measure
derived measure
recall
precision
specificity
binary prediction evaluation function
predictive model evaluation measure
class prediction evaluation measure
confusion matrix
has participant
averaged binary prediction measure
multi-class prediction evaluation measure
probability error-based measure
numeric prediction evaluation measure
has participant
kappa statistic
predictive accuracy
class RMSE
correlation coefficient
error-based evaluation measure
ROC_curve
cost curve
information criterion
graphical evaluation measure
lift chart
has part
RMSE
MAD
RSS
RRSE
MAPE
AIC
BIC
precision-recall curve
PRgraph point
has part

Experiment context

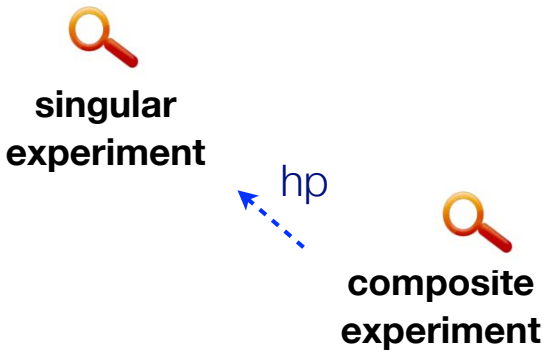

Experiment context

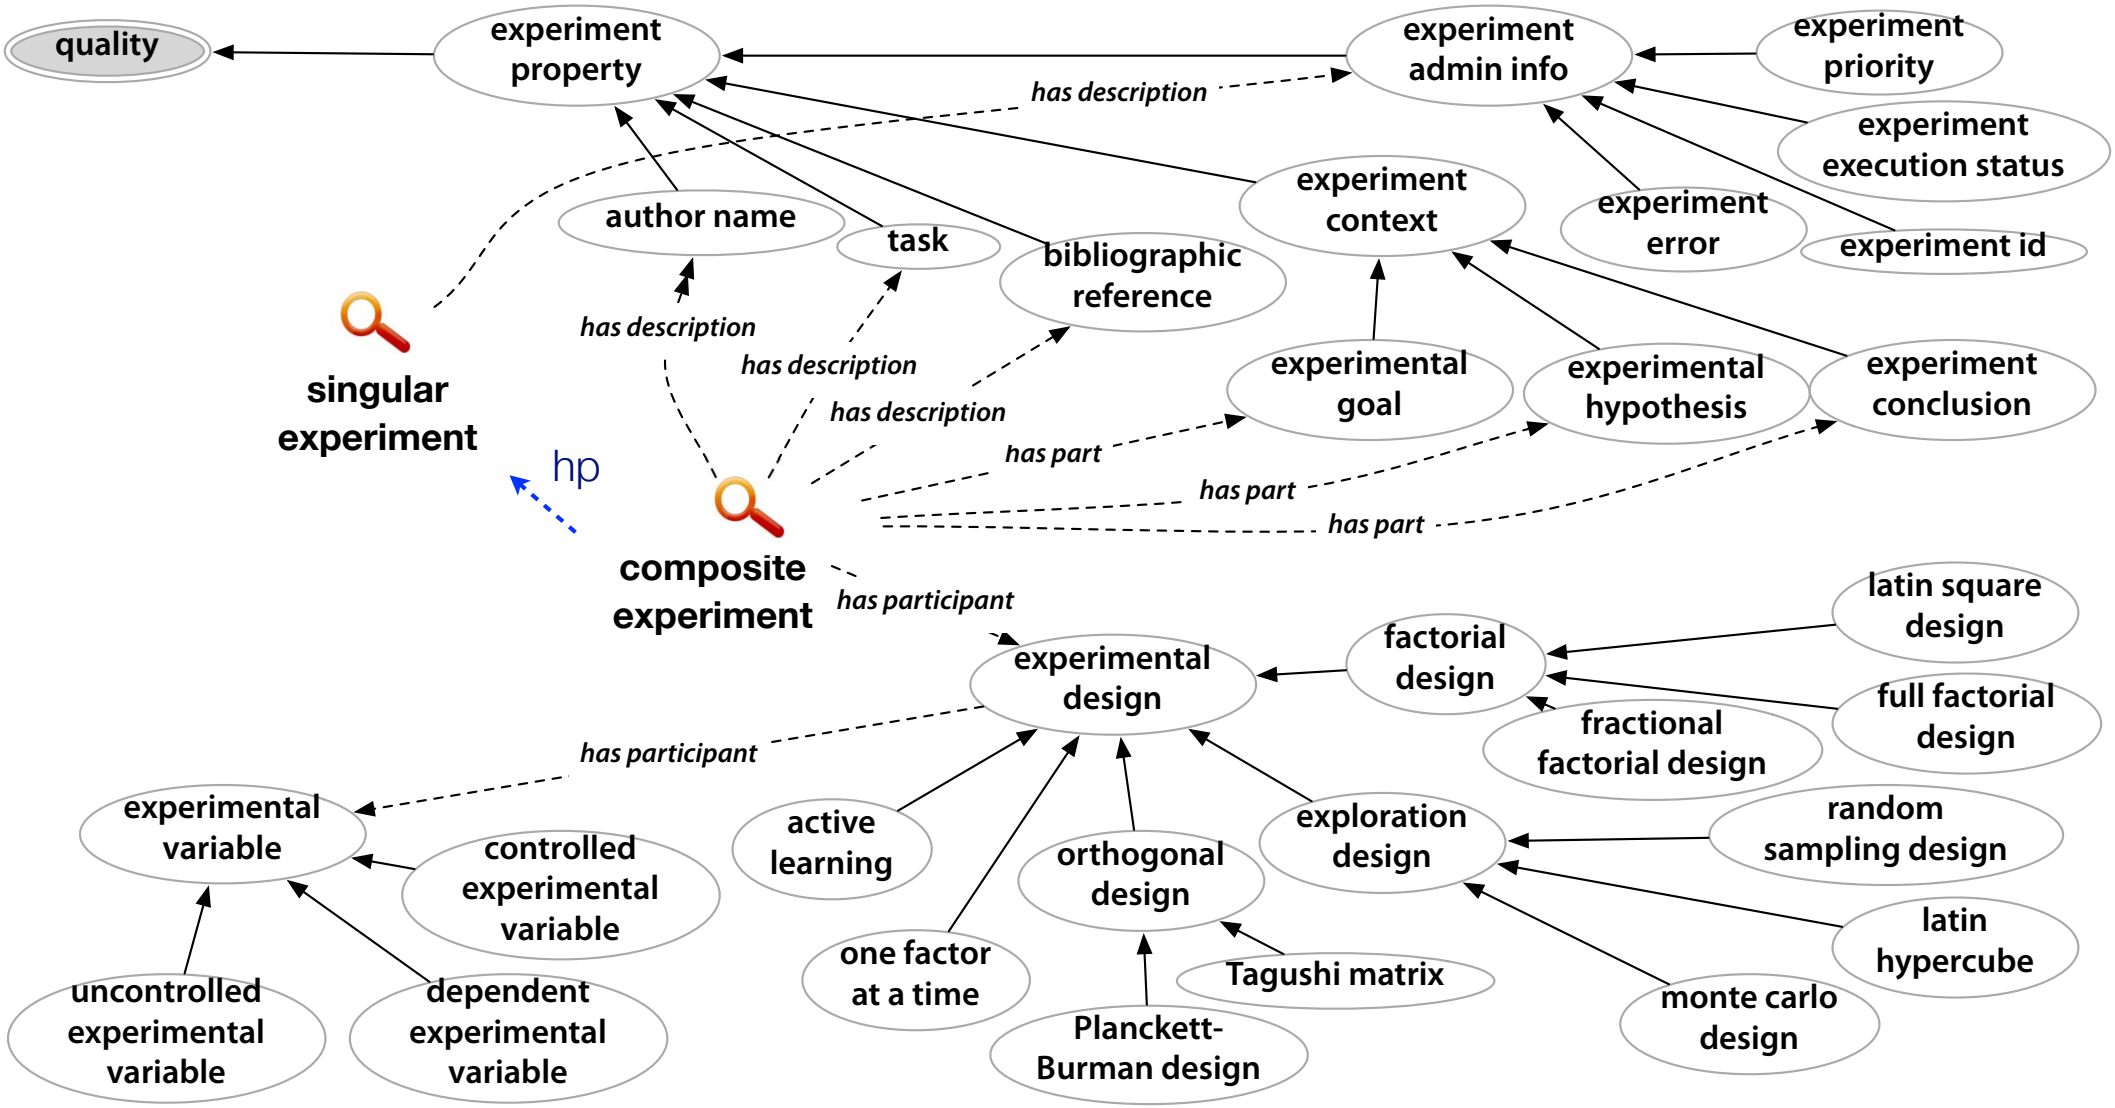

Exposé: final notes

- In total 860 classes, 32 properties (from RO + DMOP)
- Individuals: all algorithms, preprocessors, evaluation from WEKA
  - actually stored in experiment database
  - should be programmatically added (and updated)
- Written in OWL-DL, using Protégé 4.0
- Can be browsed at:
  - http://expdb.cs.kuleuven.be/expdb/expose.owl
  - http://www.e-lico.eu/OWLBrowser2/manage/

Use Cases

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

new algorithm

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

datasets

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

preprocessing
workflows

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

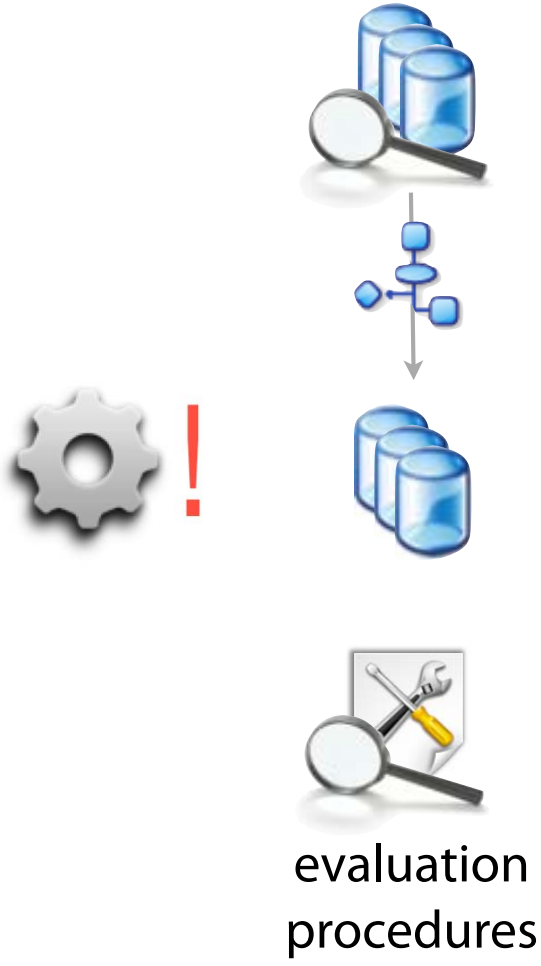

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

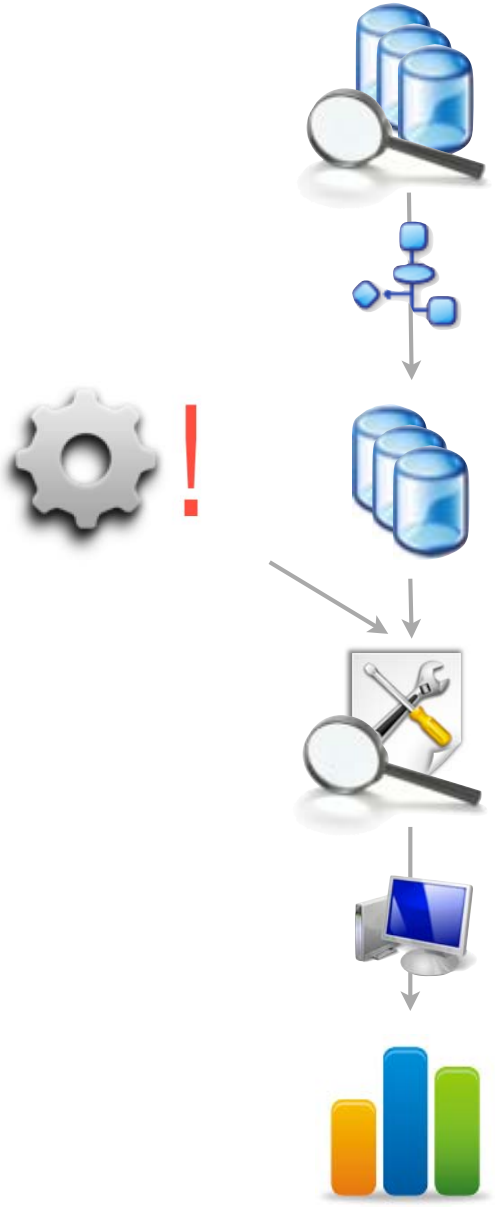

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

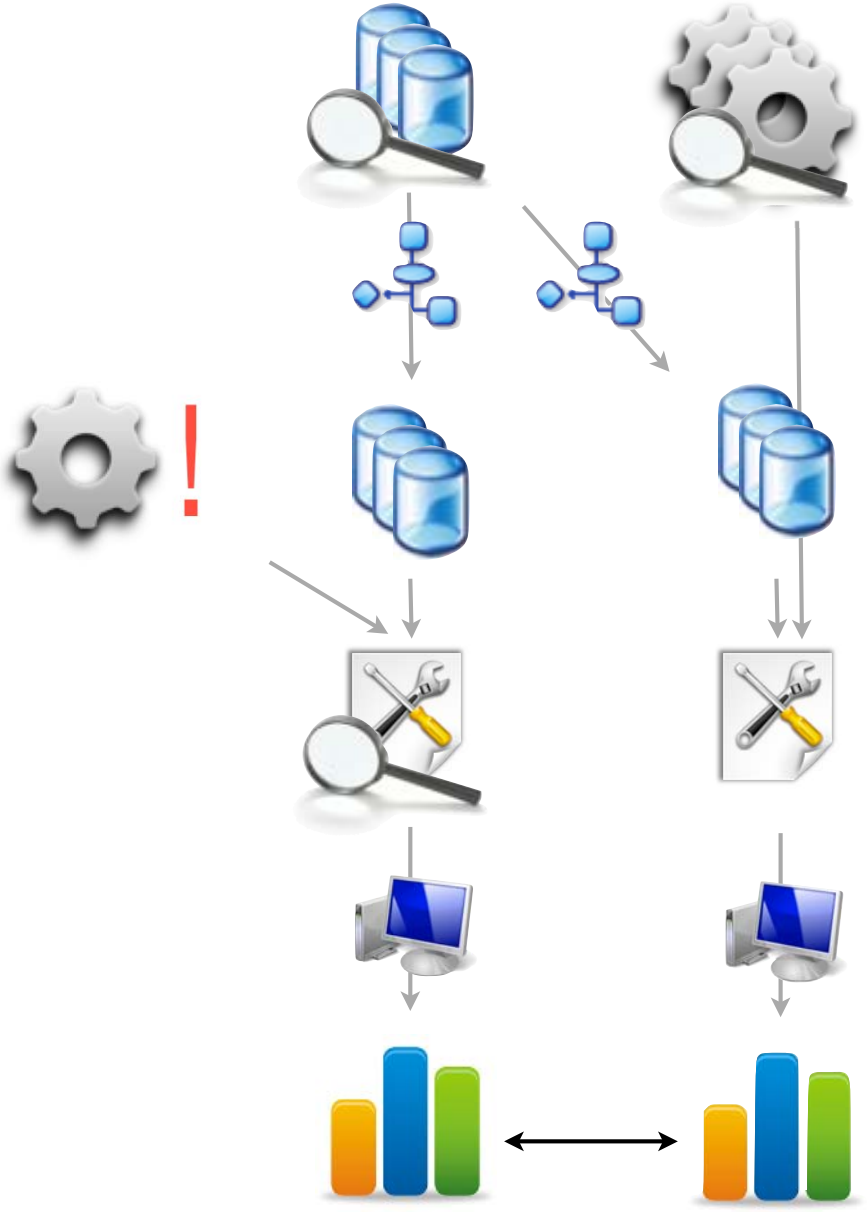

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

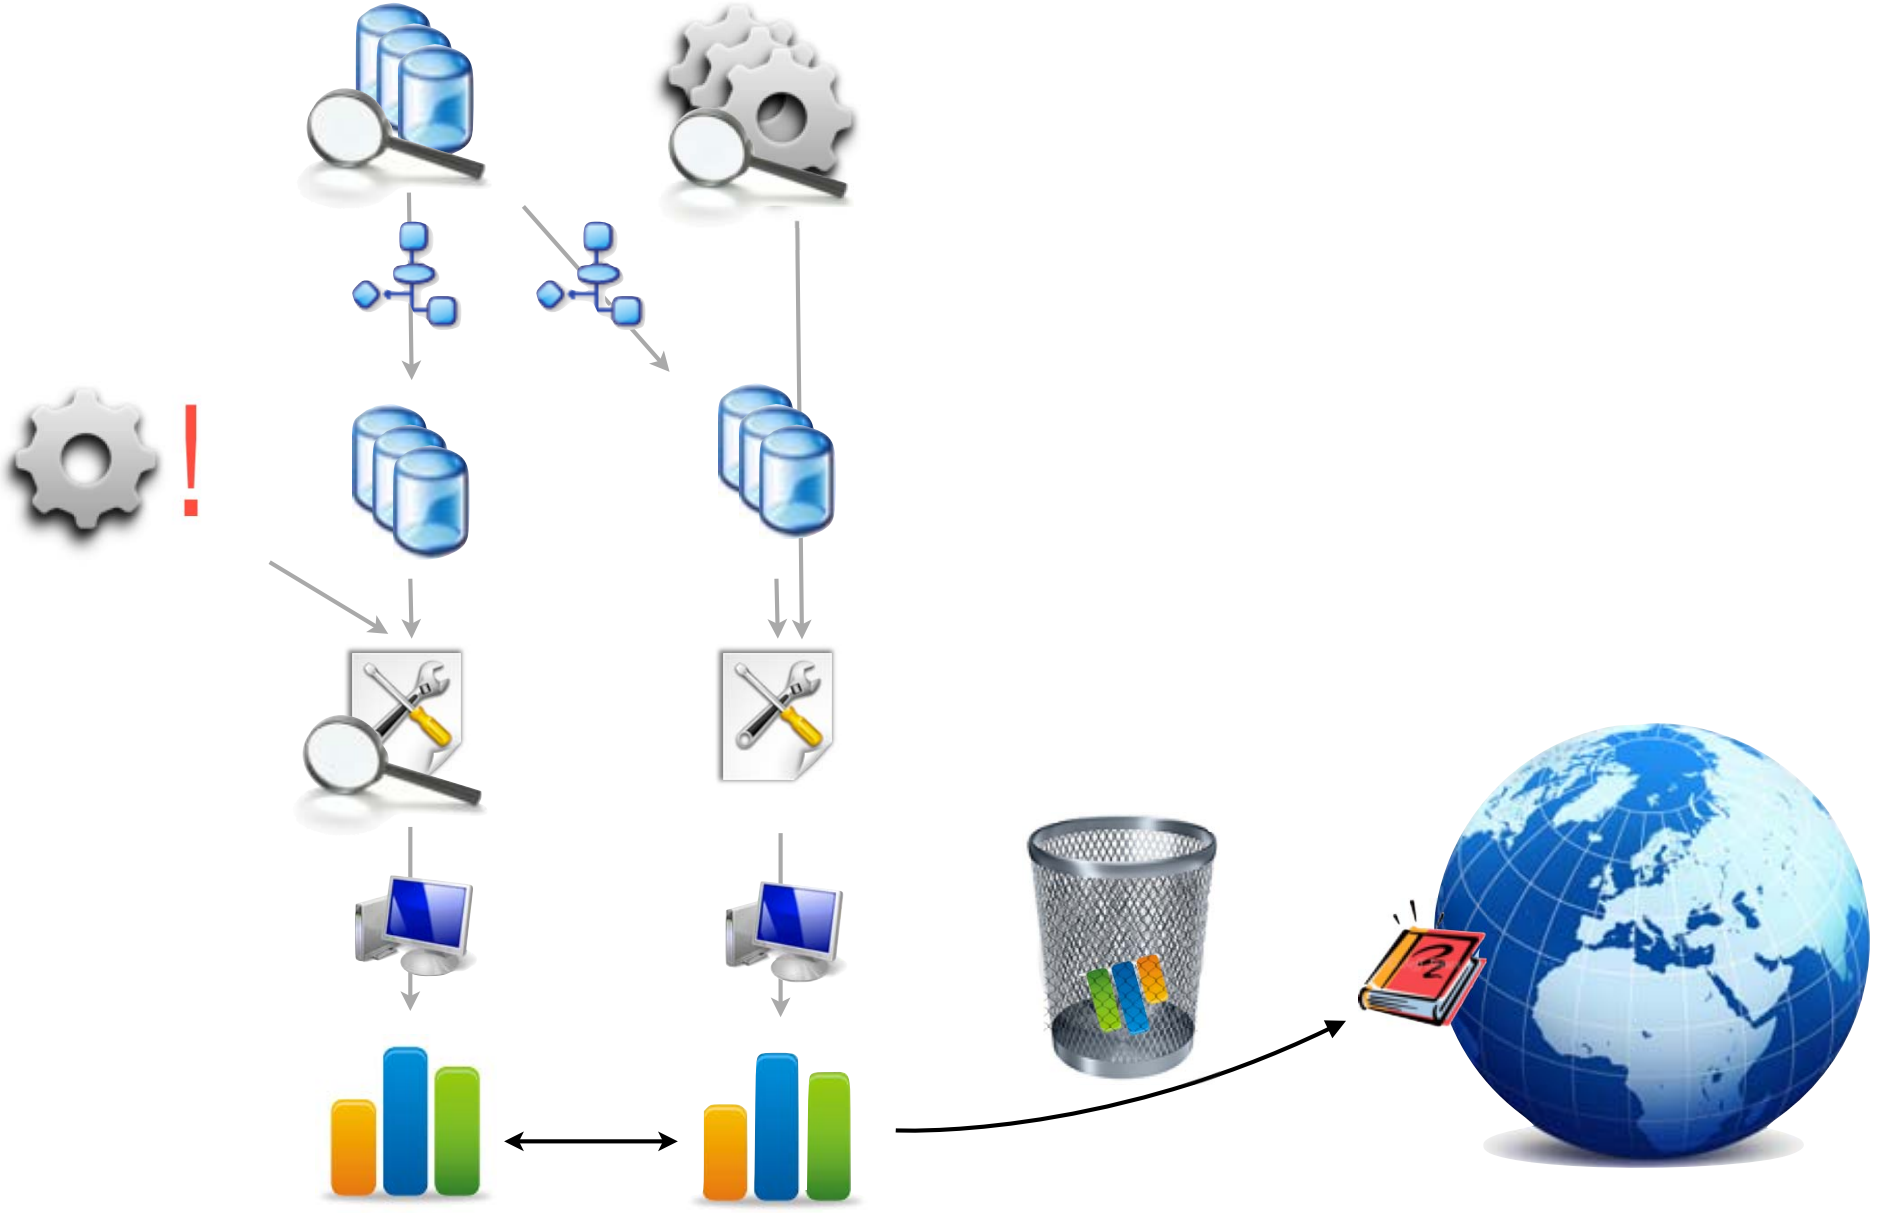

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

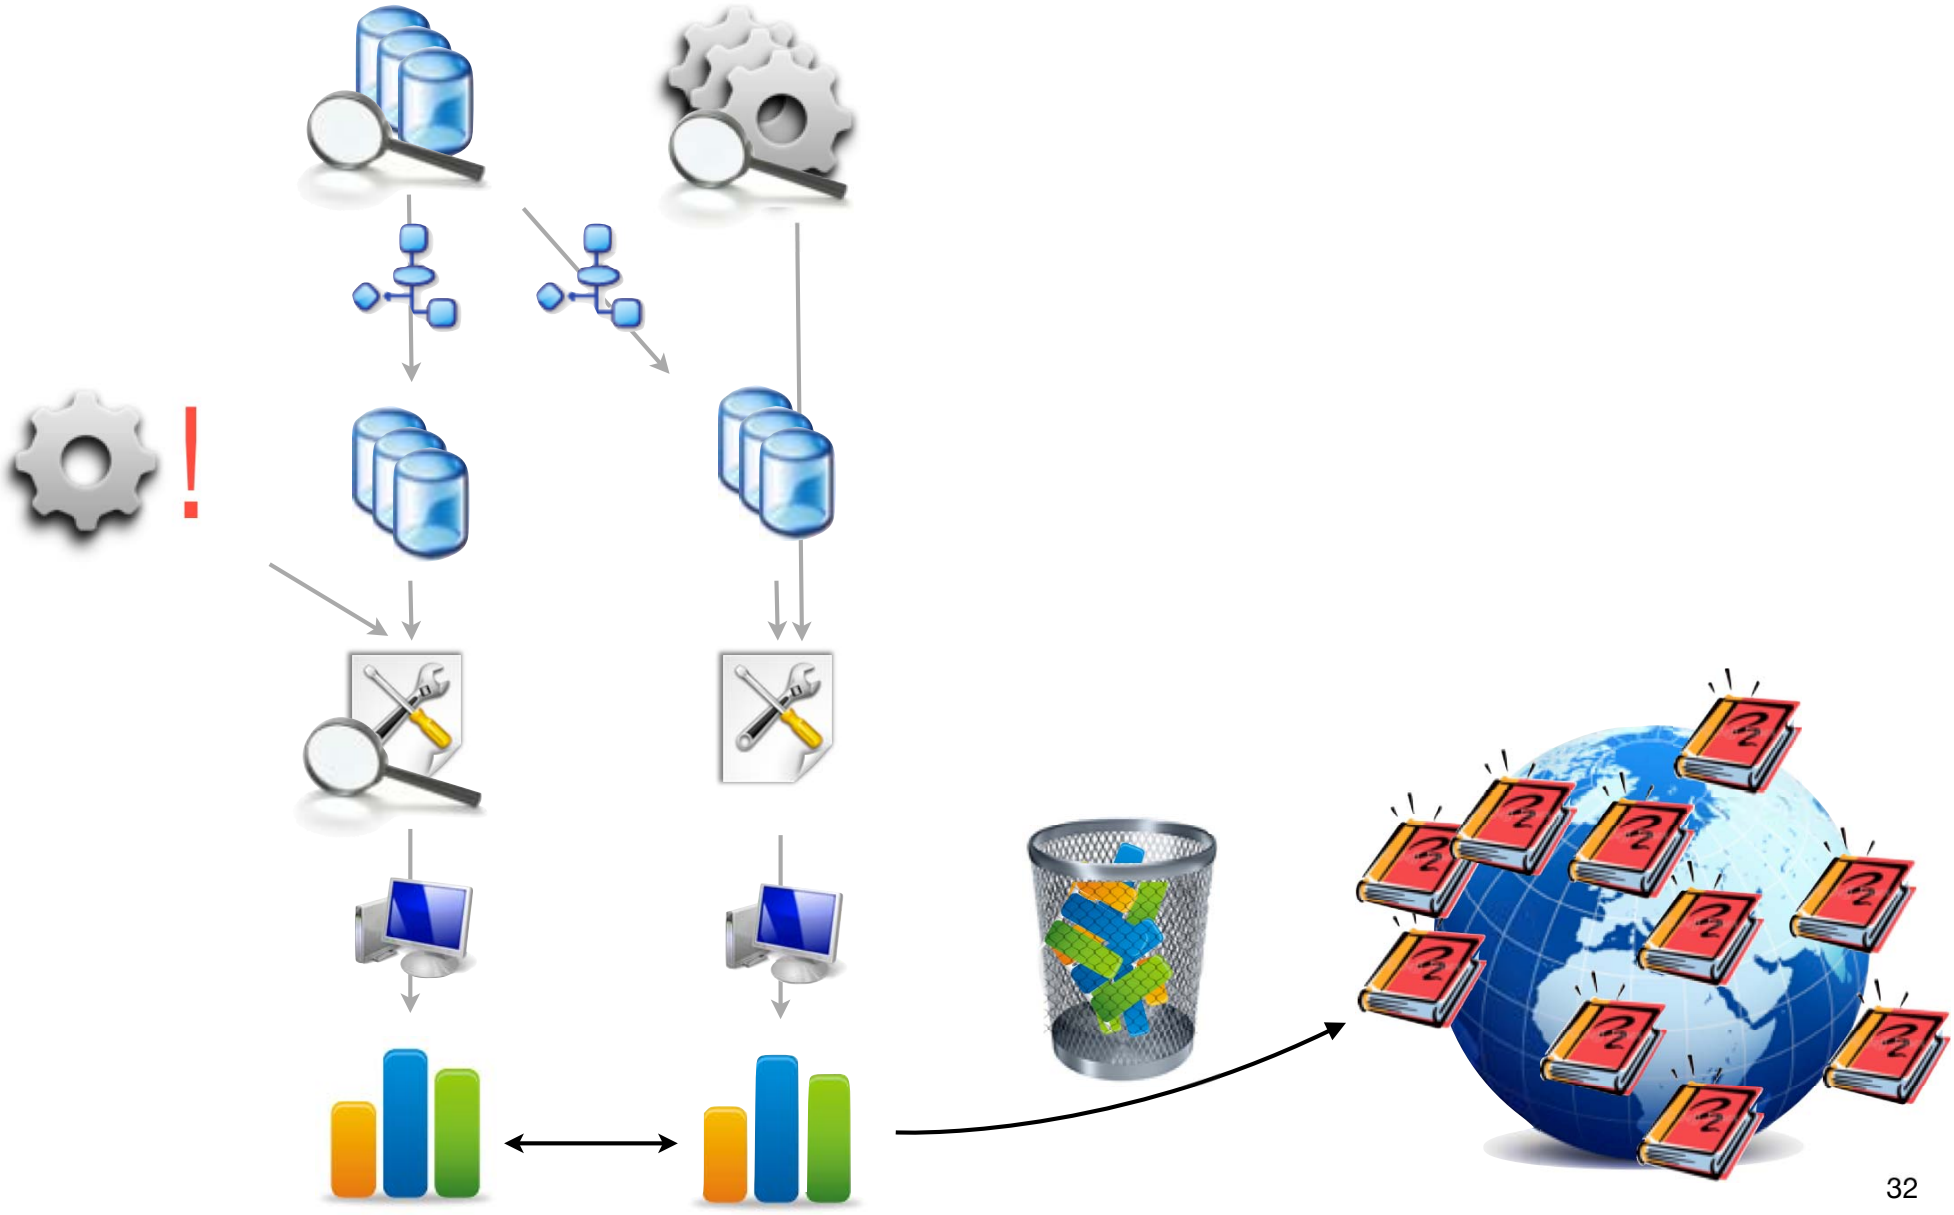

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

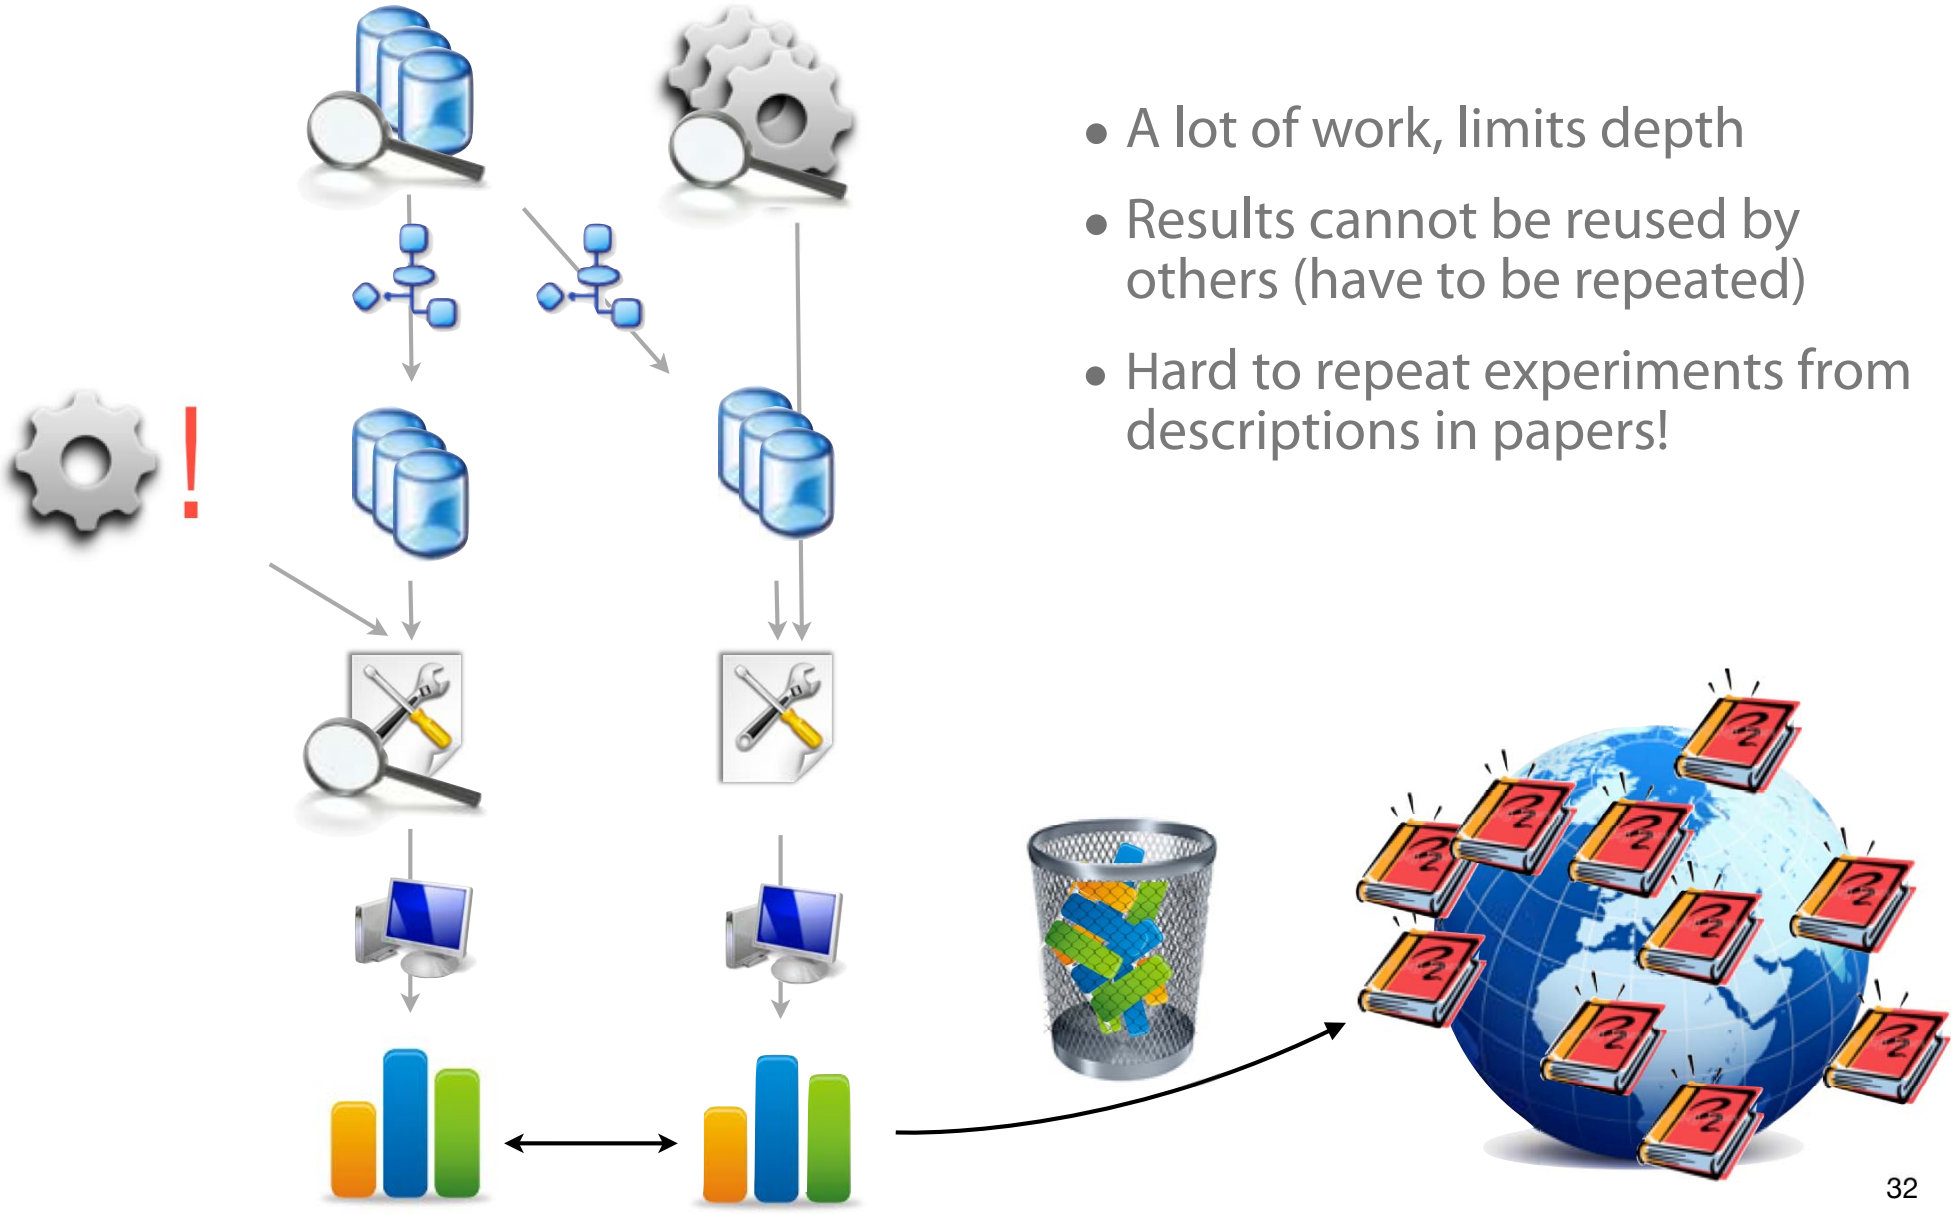

- A lot of work, limits depth
- Results cannot be reused by others (have to be repeated)
- Hard to repeat experiments from descriptions in papers!

Goal: Collaborative experimentation
Now: small-scale, not repeatable, not reusable

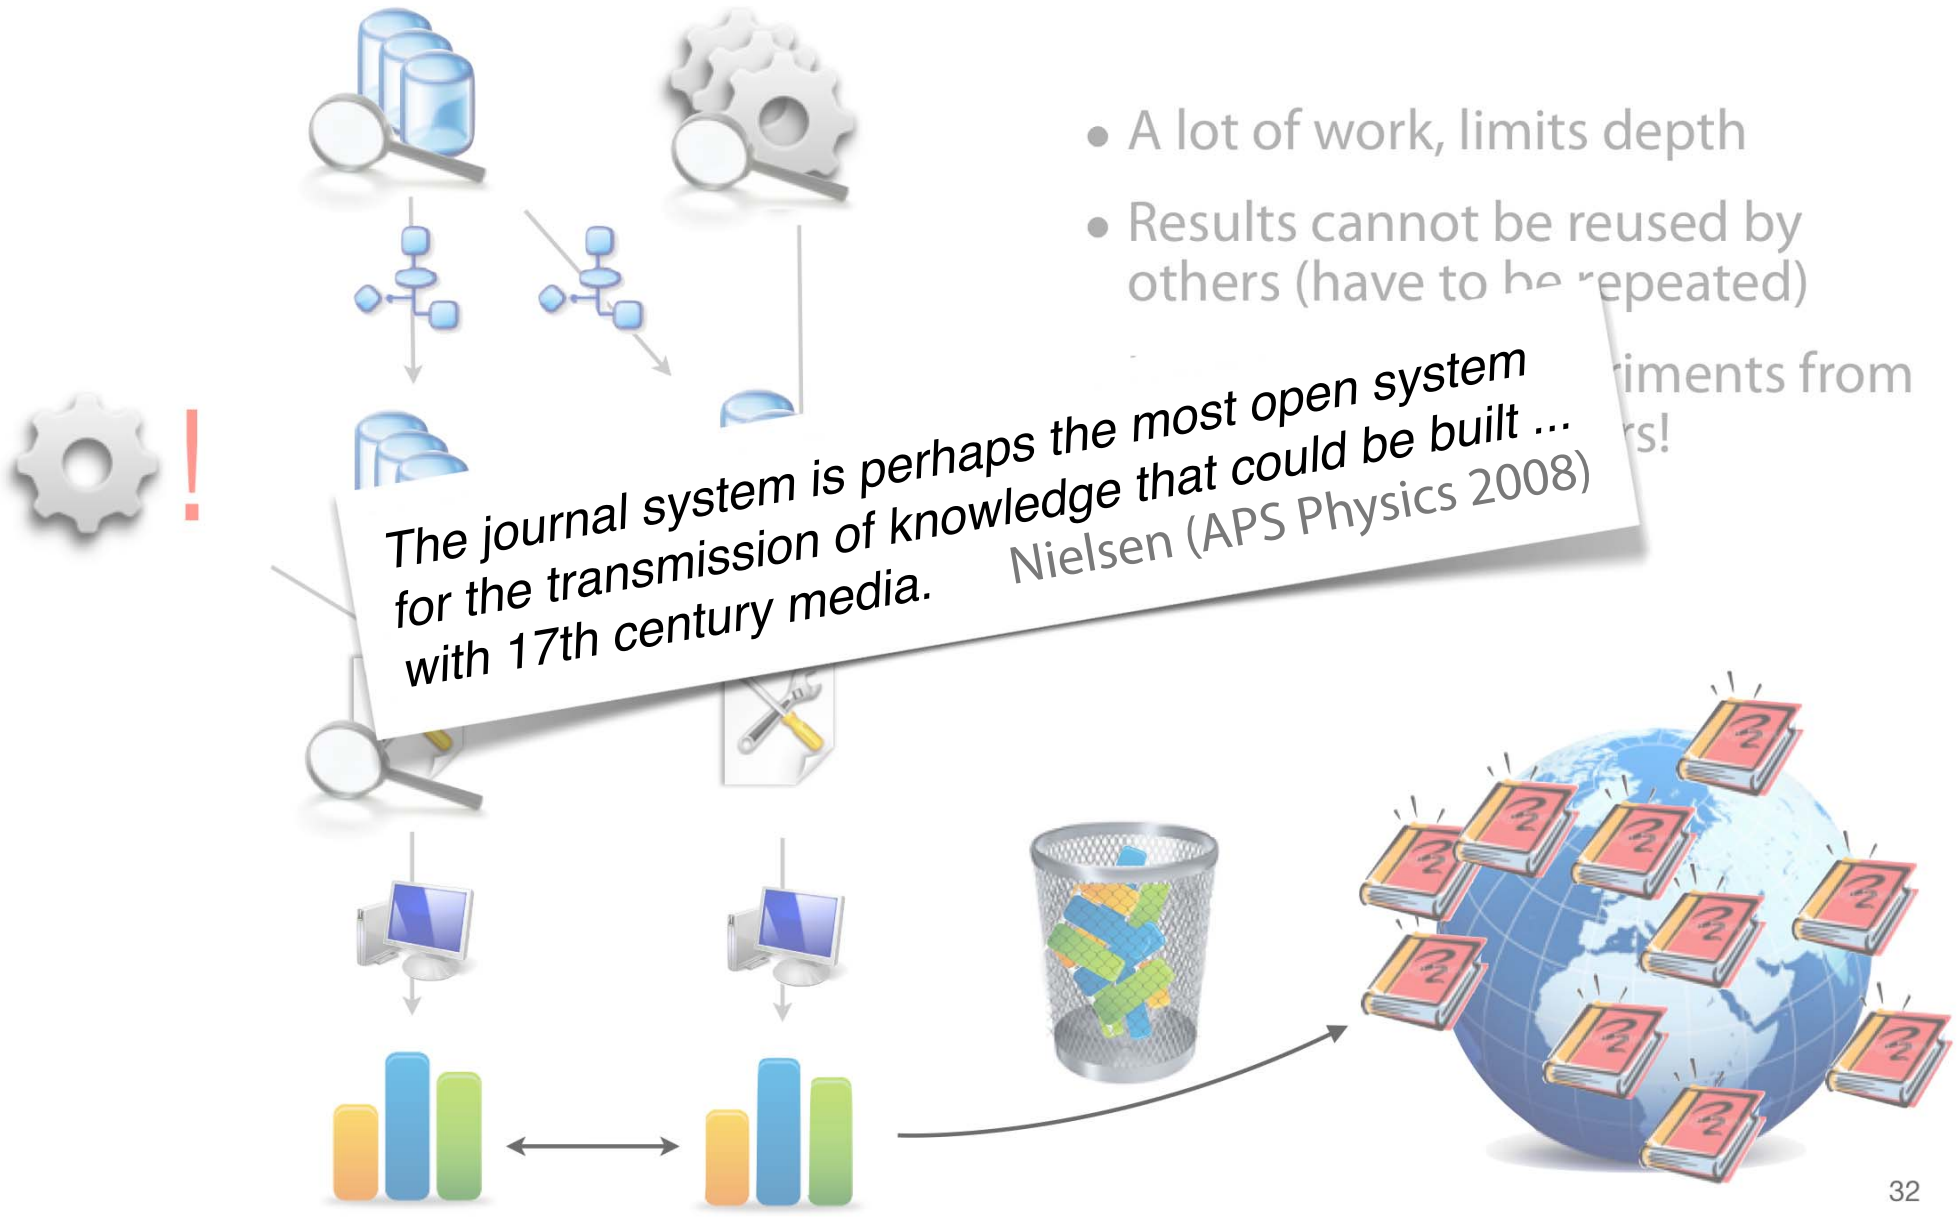

Data mining as an e-science
Ontologies: experiments shared, run automatically

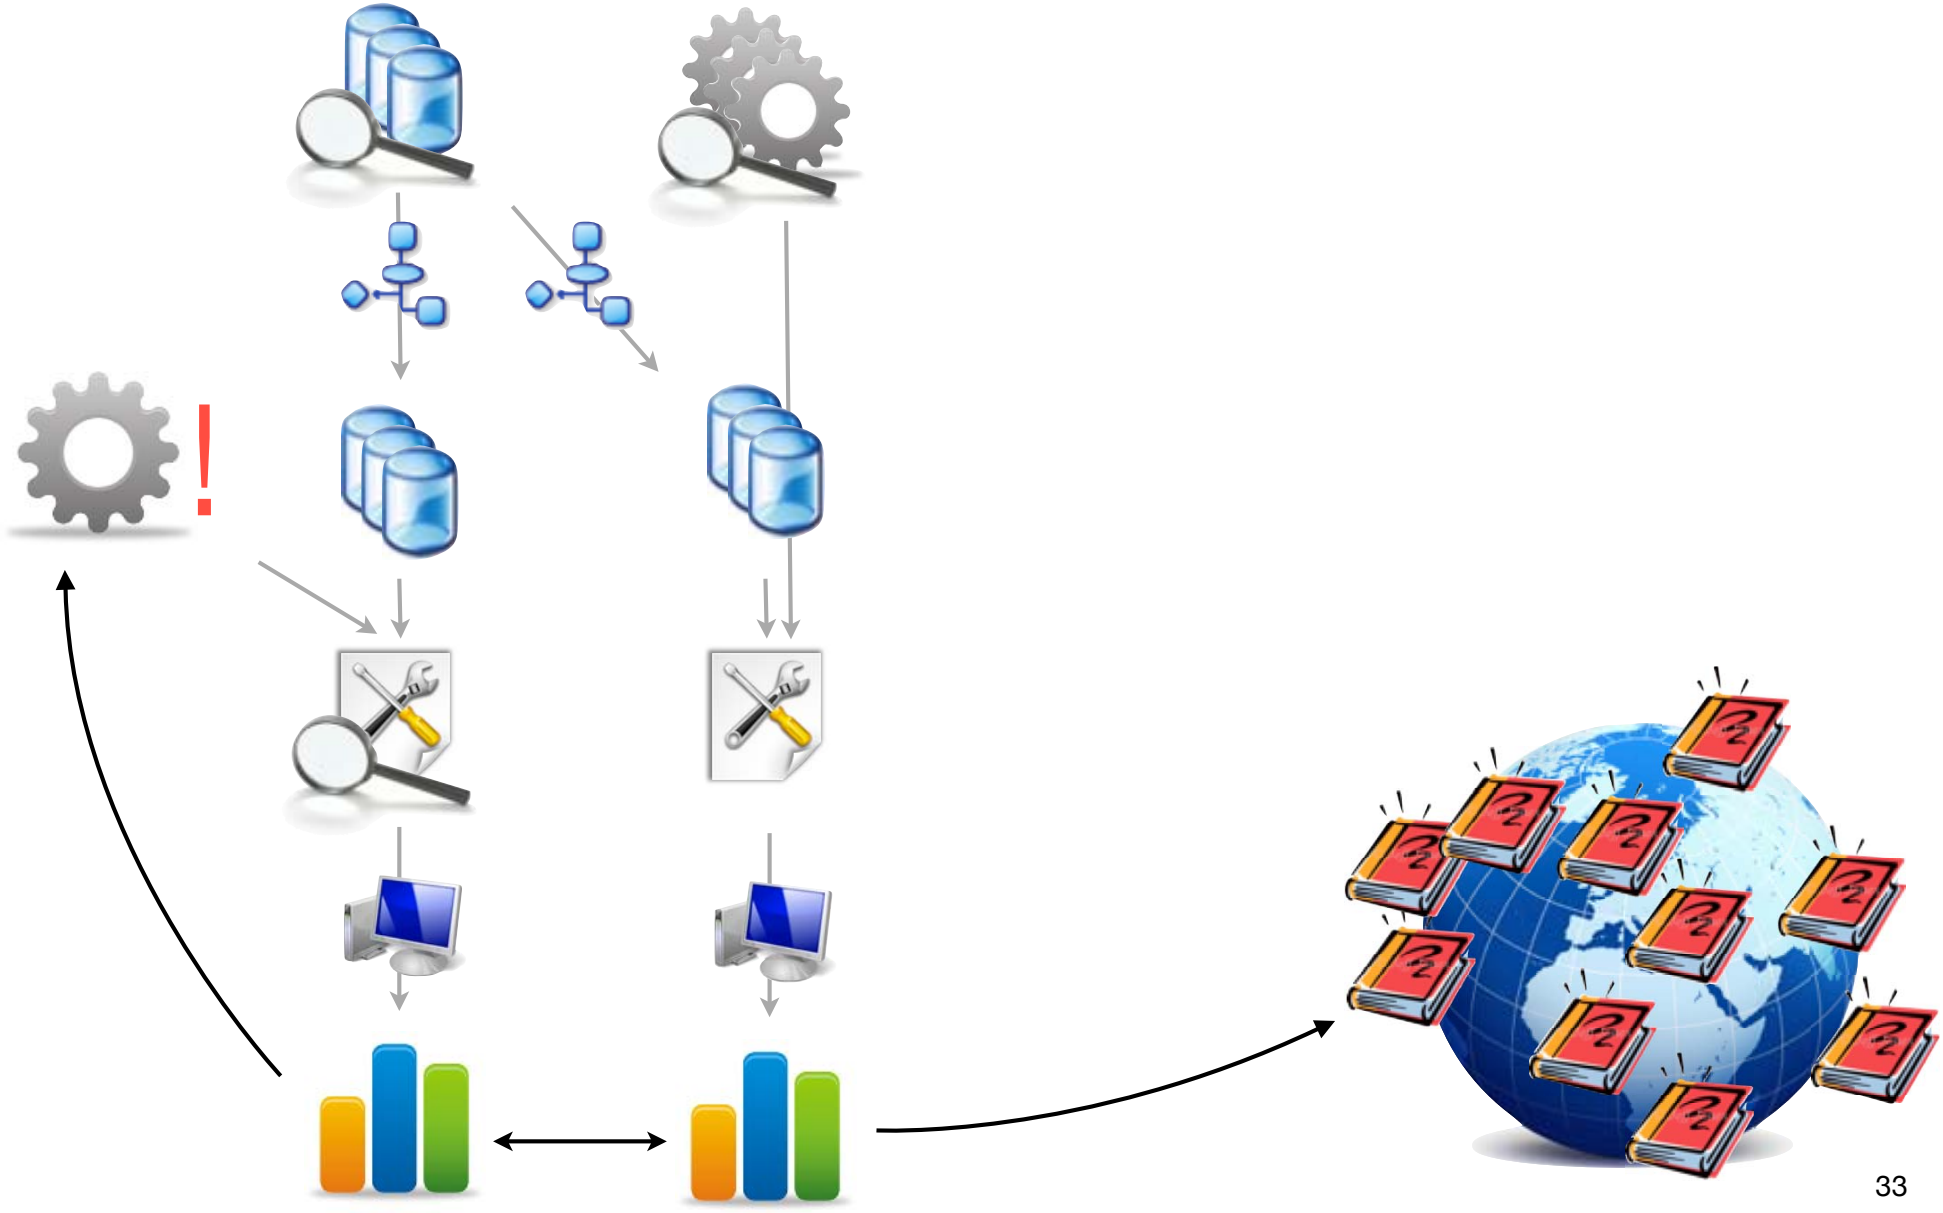

Data mining as an e-science
Ontologies: experiments shared, run automatically

- Share experiments
  - Internet = large, collaborative workspace

Data mining as an e-science
Ontologies: experiments shared, run automatically

- Share experiments
  - Internet = large, collaborative workspace
- Store them in ***experiment databases***
  - Ensure reproducibility
  - Reuse millions of prior experiments
  - Use all info on algorithms, datasets
  - Results universally accessible and useful

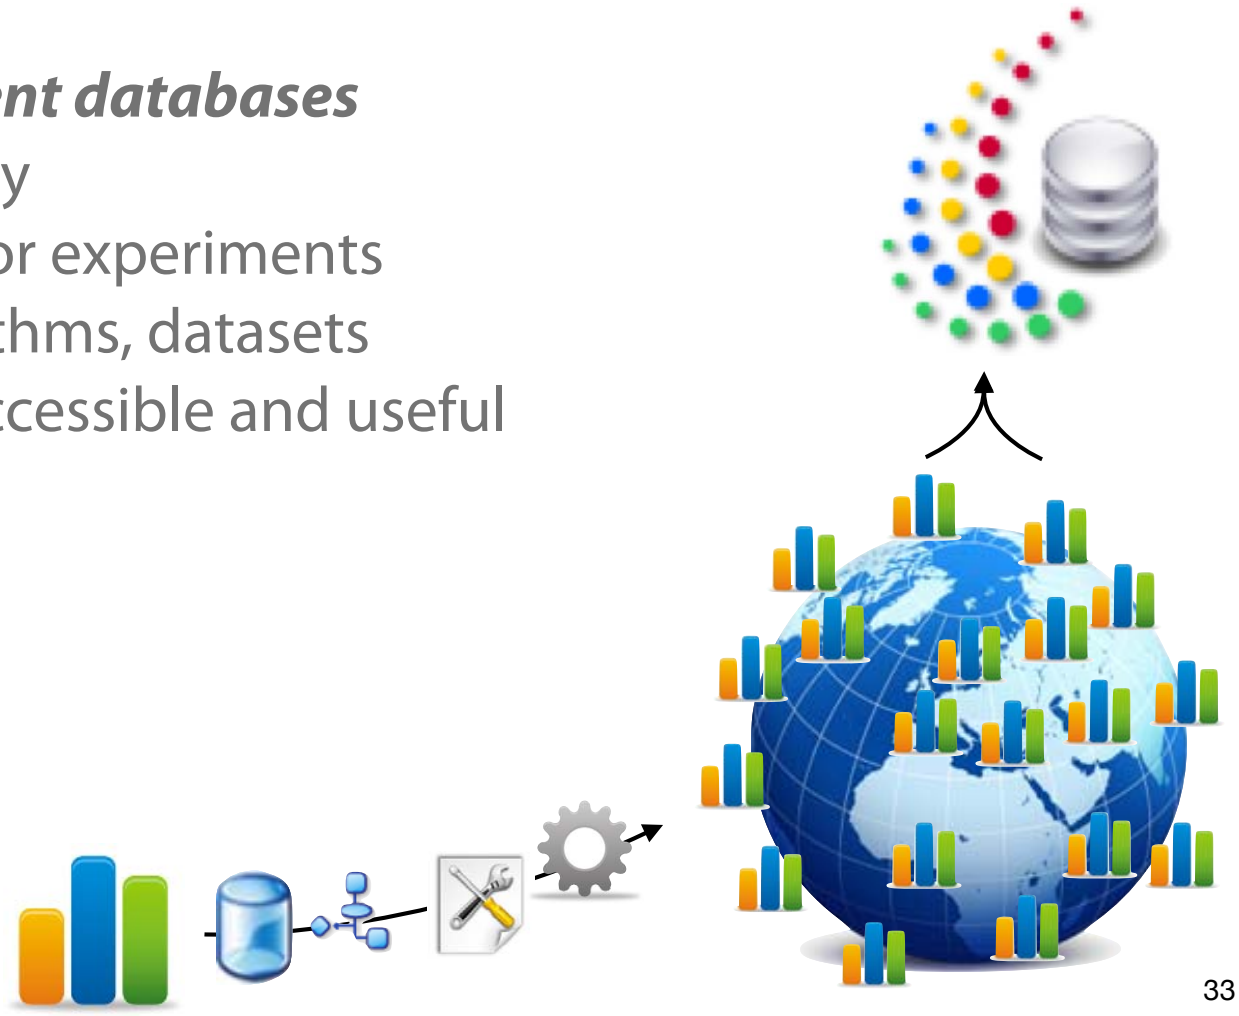

e-Sciences
Astrophysics: Virtual Observatories

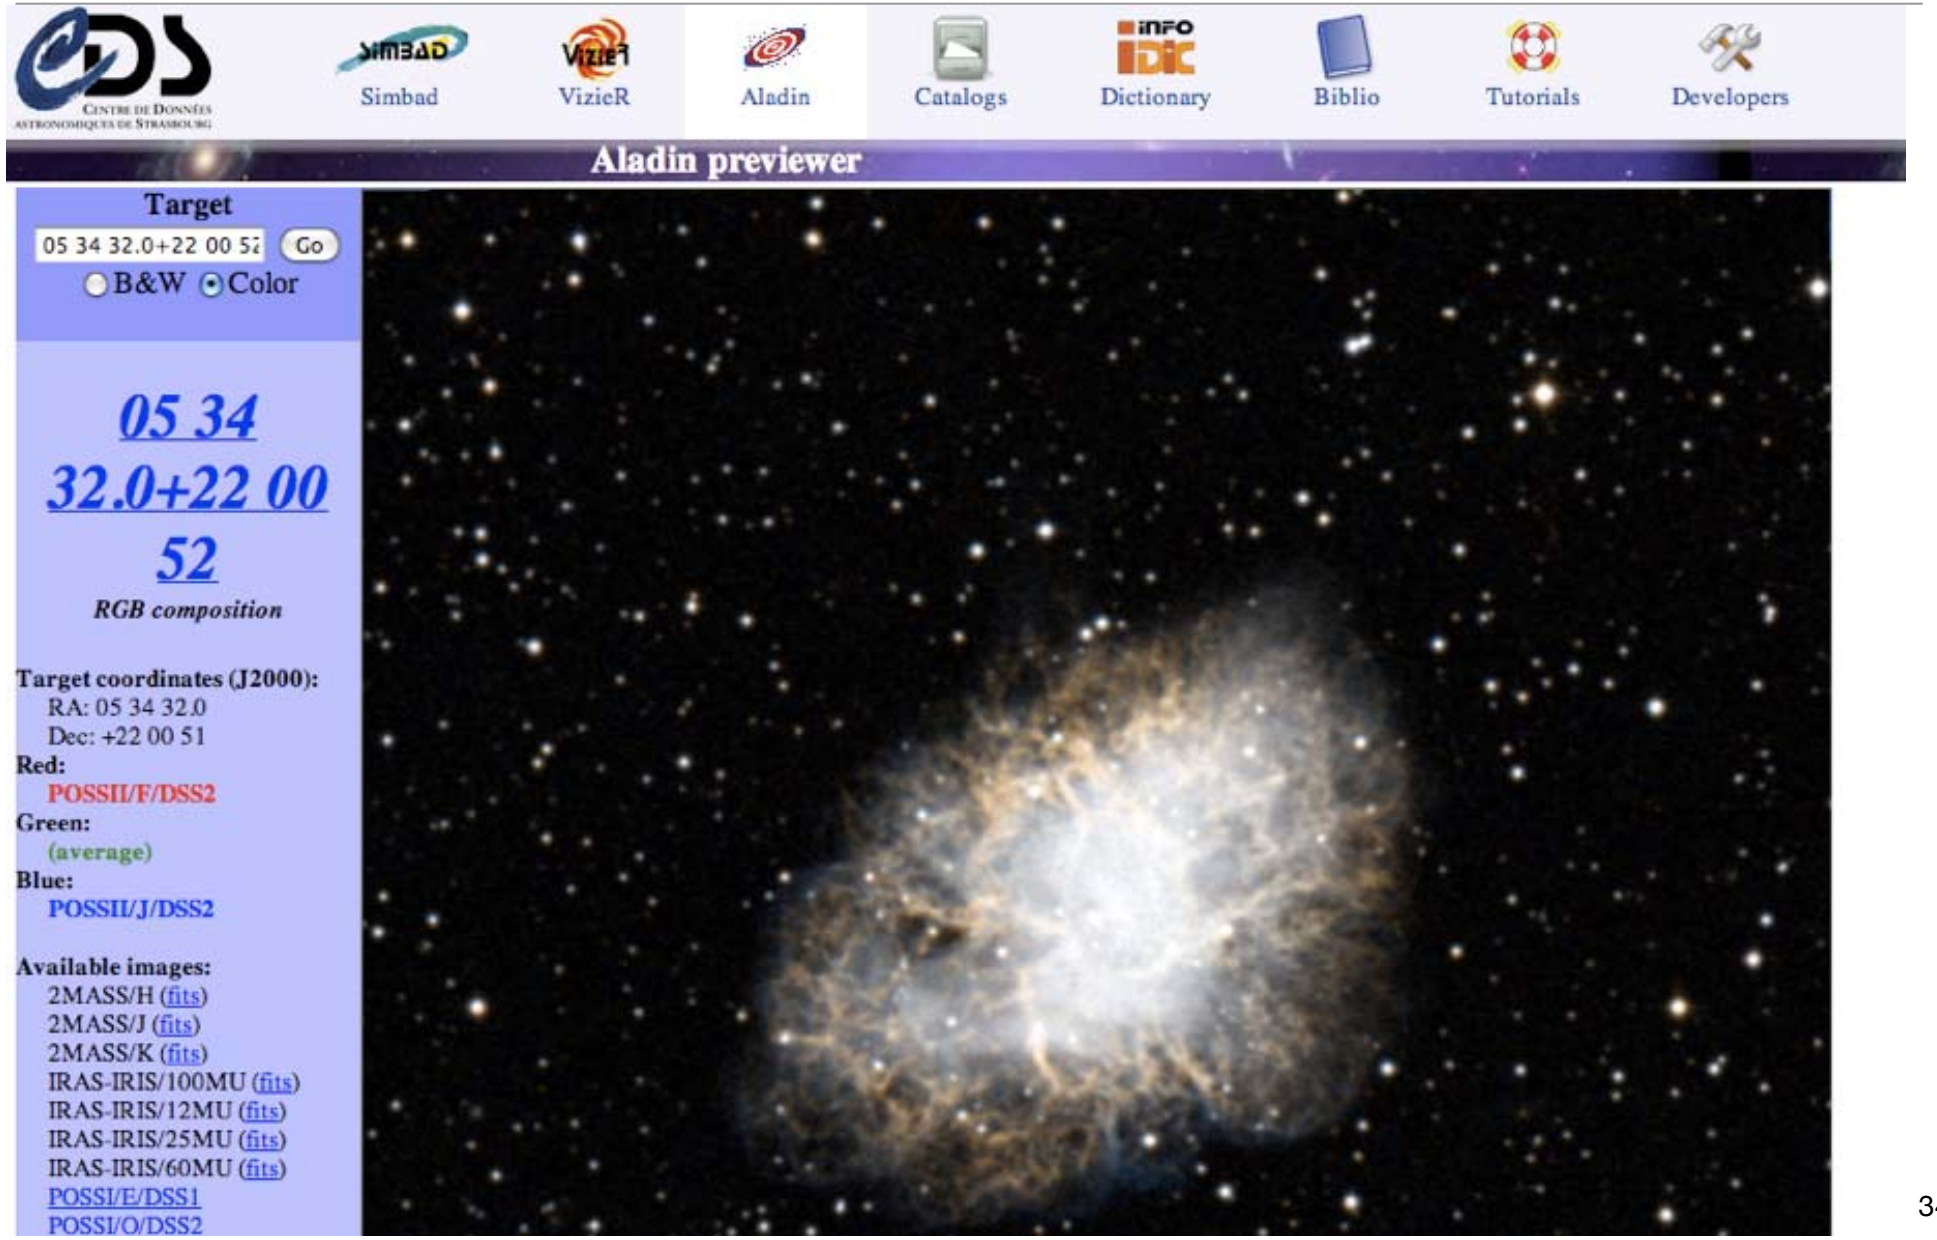

e-Sciences
Bio-informatics: Micro-array Databases

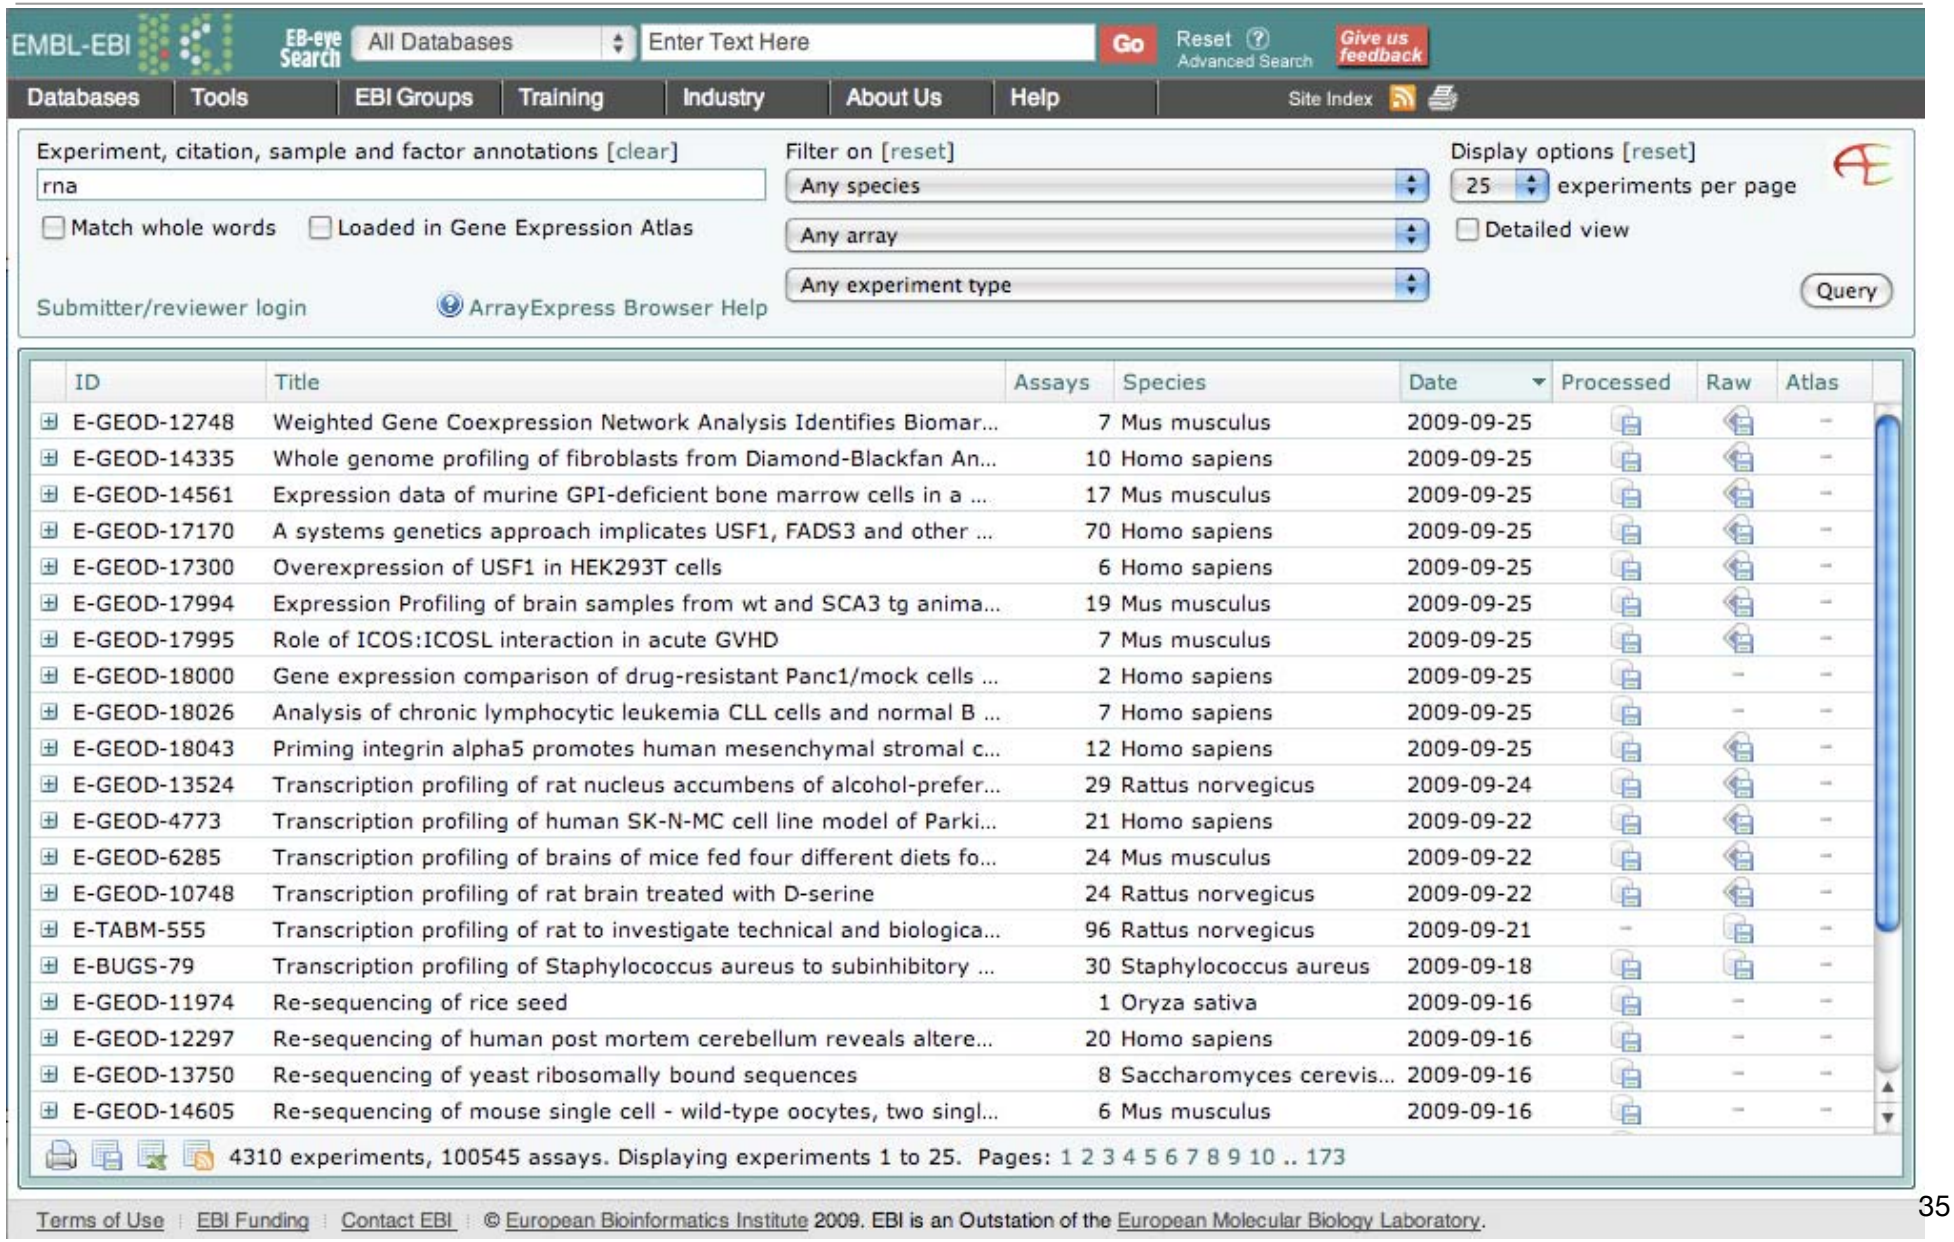

e-Sciences
Bio-informatics: Micro-array Databases

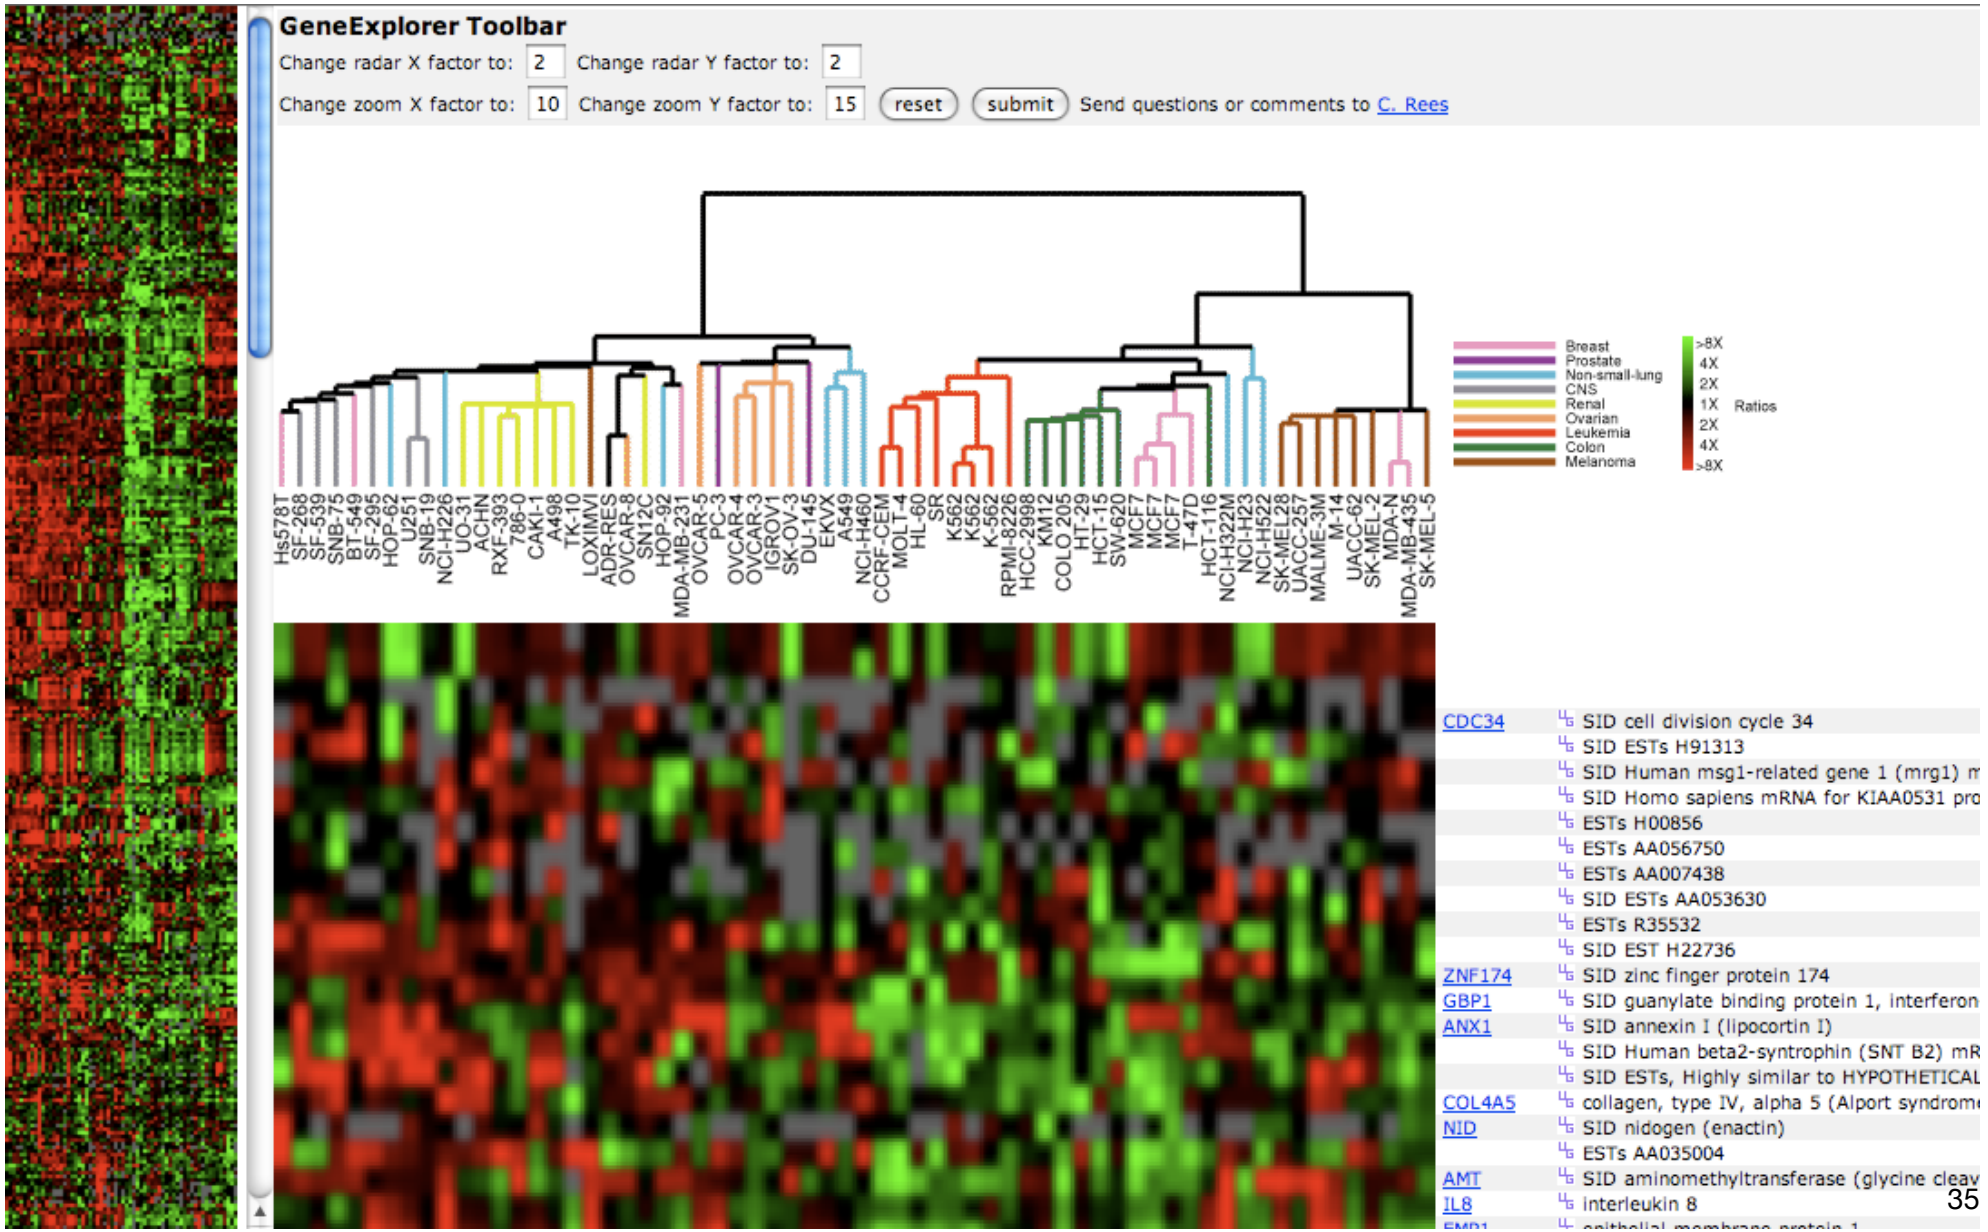

Collaborative Experimentation
Why?

Collaborative Experimentation
Why?

Reproducibility

Good science

Collaborative Experimentation
Why?

Reproducibility

Good science

Quick, easy analysis

Querying: Answer questions

Test hypotheses

Collaborative Experimentation
Why?

Reproducibility

Good science

Quick, easy analysis

Querying: Answer questions

Test hypotheses

Reuse

Save time & energy

(e.g. benchmarking)

Collaborative Experimentation
Why?

Reproducibility

Good science

Quick, easy analysis

Querying: Answer questions

Test hypotheses

Reuse

Save time & energy

(e.g. benchmarking)

Generalizability:

Plug into prior results: larger studies

Collaborative Experimentation
Why?

Reproducibility

Good science

Quick, easy analysis

Querying: Answer questions

Test hypotheses

Integration

Data mining tools

import/export

Reuse

Save time & energy

(e.g. benchmarking)

Generalizability:

Plug into prior results: larger studies

Collaborative Experimentation
Why?

Reproducibility

Good science

Quick, easy analysis

Querying: Answer questions

Test hypotheses

Integration

Data mining tools

import/export

Reuse

Save time & energy

(e.g. benchmarking)

Reference

'Map' of known approaches

Compare to state-of-the-art

Includes negative results

Generalizability:

Plug into prior results: larger studies

Use Case 1

Describe experiments in a common language
-> sharing or running experiments on grid

Use Exposé to define common language: ExpML

Use Exposé to define common language: ExpML

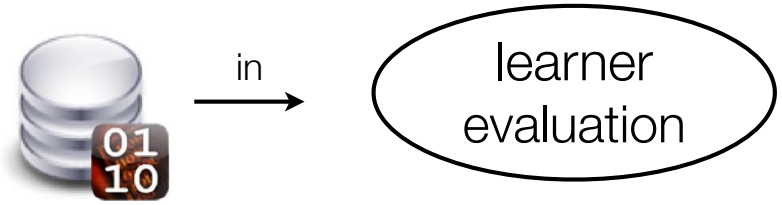

Use Exposé to define common language: ExpML

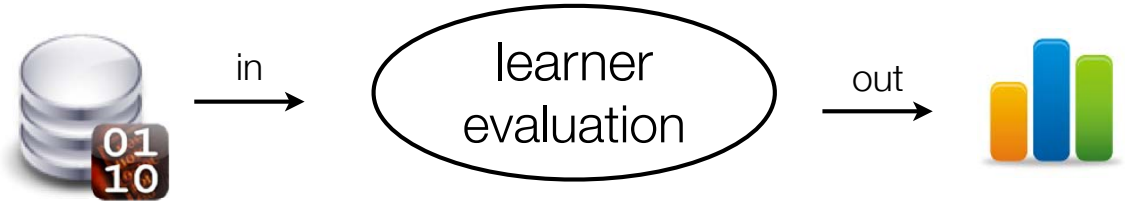

Use Exposé to define common language: ExpML

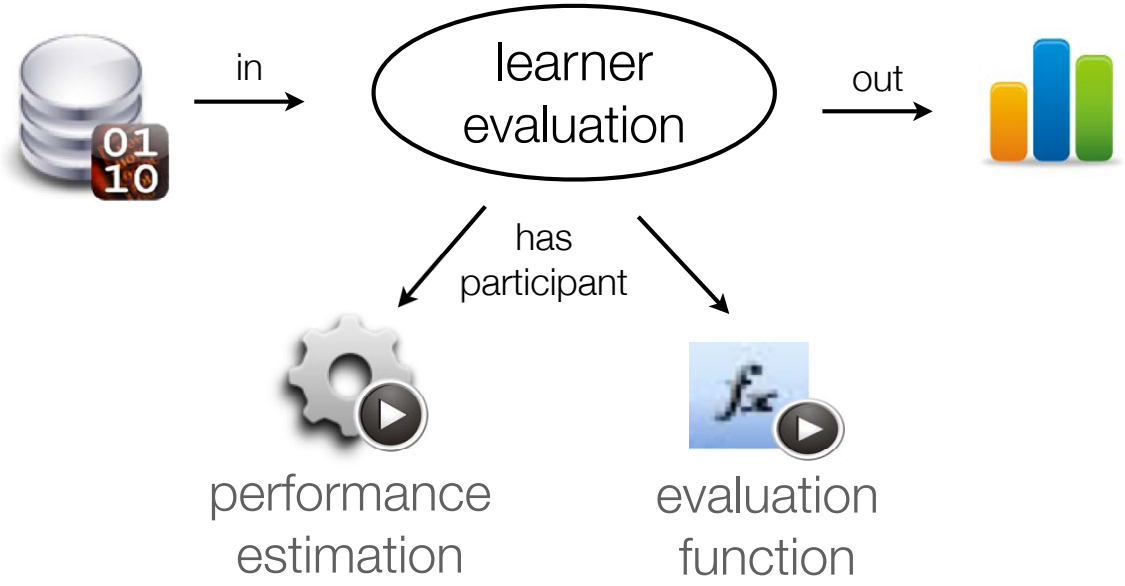

Use Exposé to define common language: ExpML

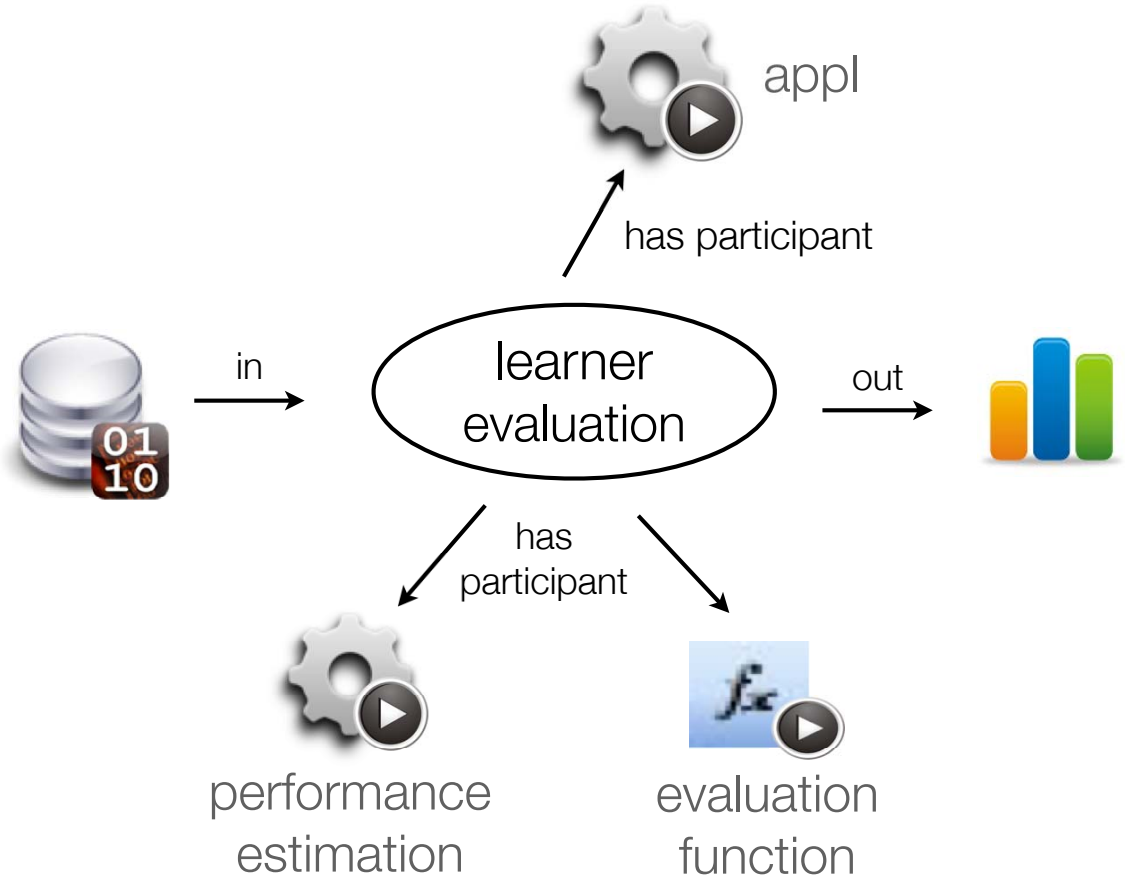

Use Exposé to define common language: ExpML

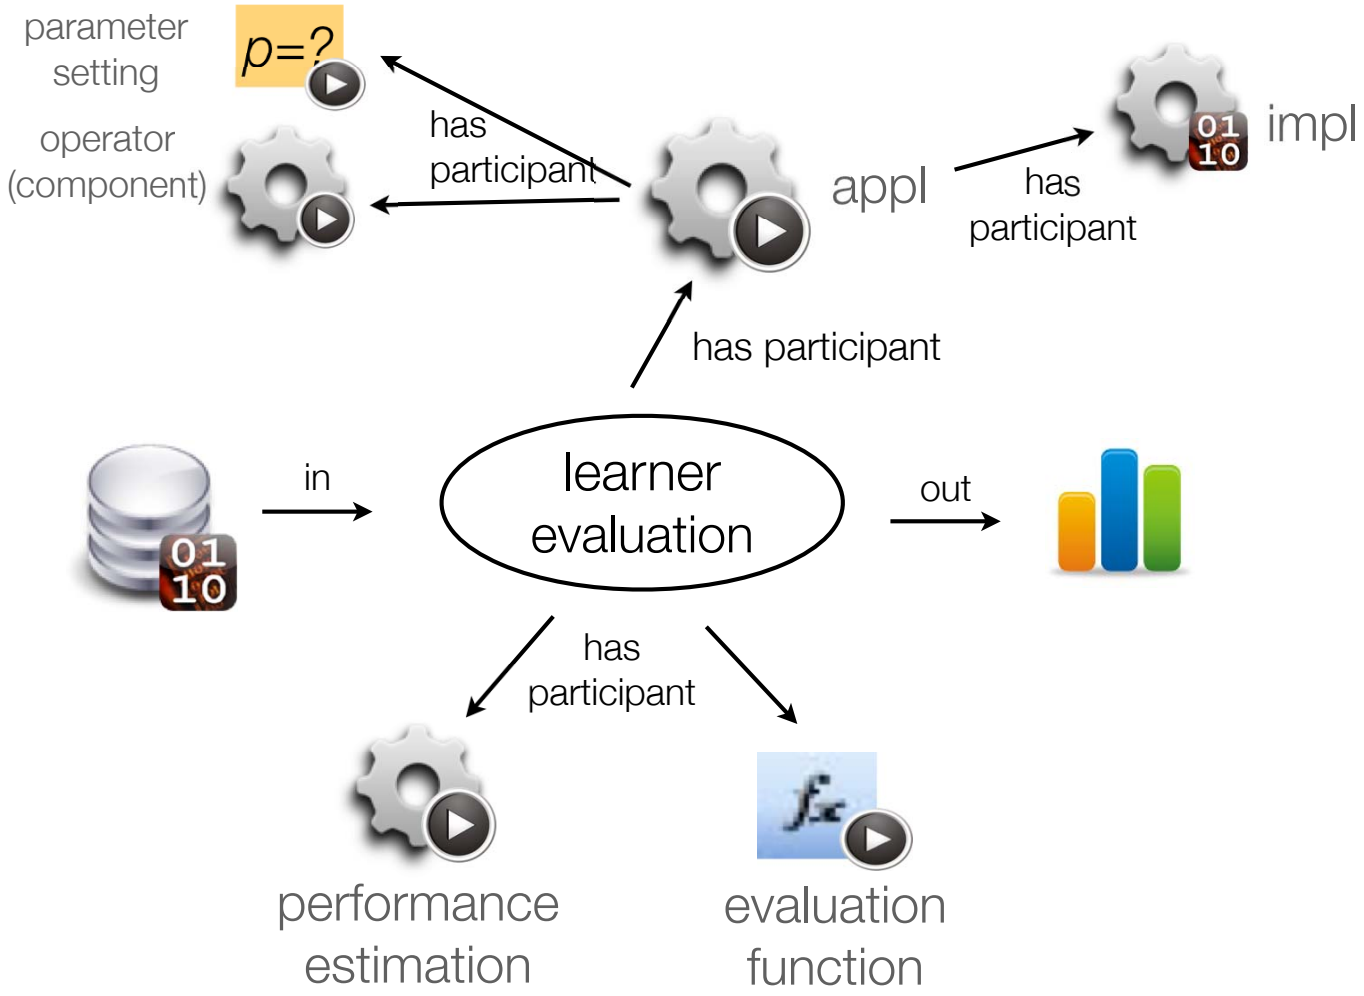

ExpML: a markup language for DM experiments

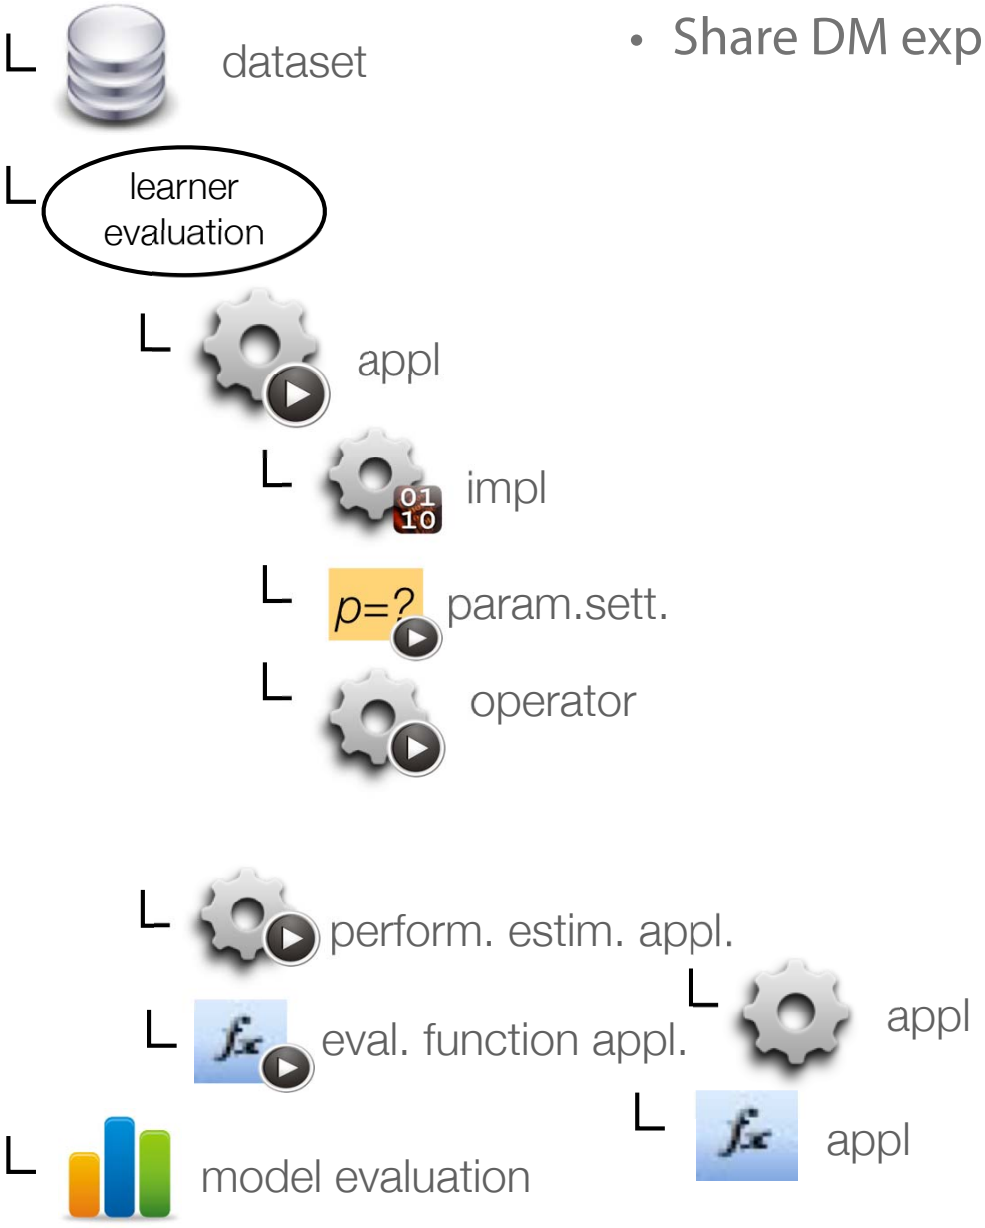

- Share DM experiments, XML-based

ExpML: a markup language for DM experiments

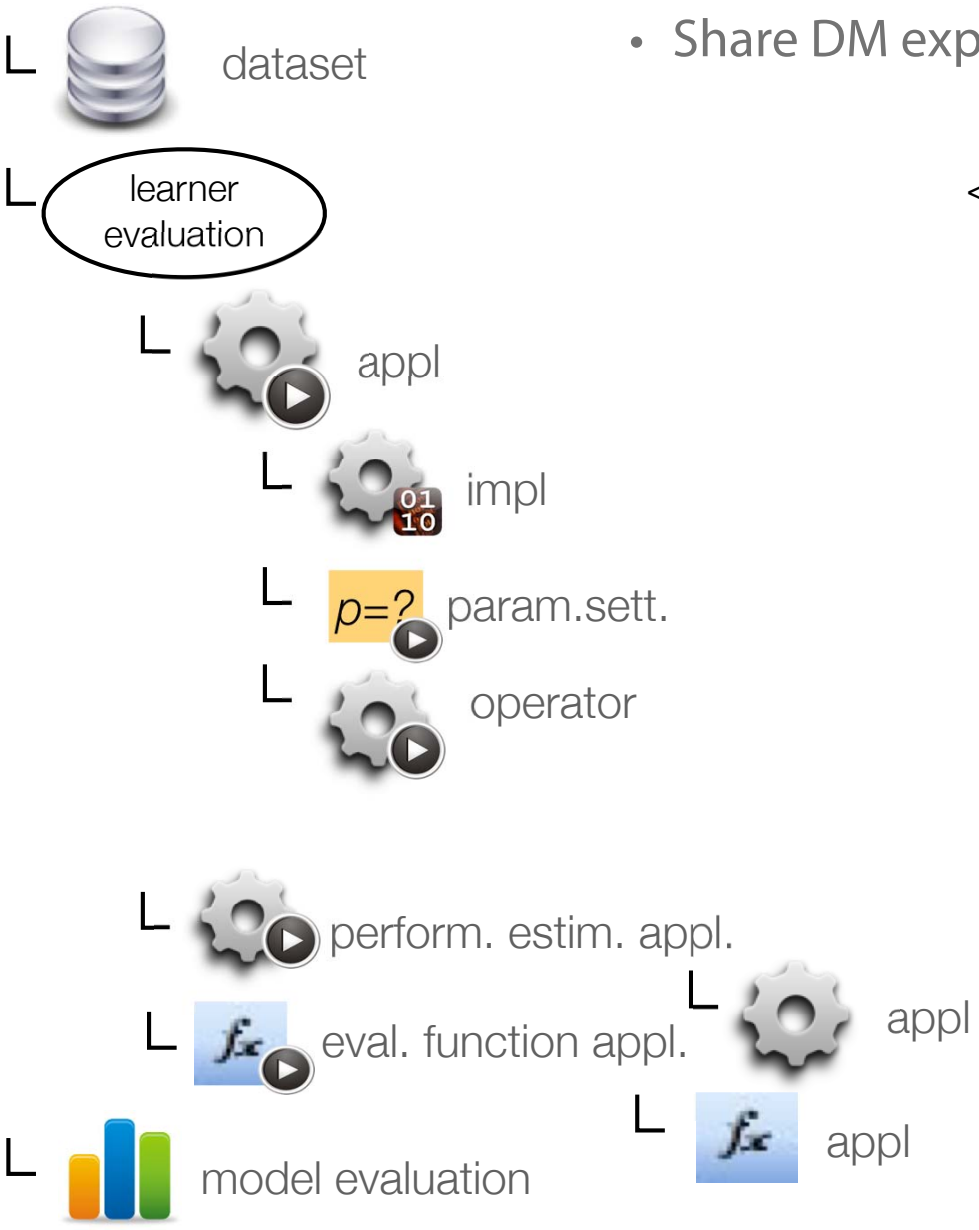

- Share DM experiments, XML-based

```
<expml>
  <dataset id='d1'>
  <learner evaluation id='e1' input_data='d1'>
    <learner_appl>
      <learner_impl name=... version=...>
      <parameter_setting name='P' value='100'/>
      <learner_appl role= 'base-learner'>
        ...
    </learner_appl>
    <performance_estimation_appl>
    ...
    <model_evaluation_function_appl>
    ...
  </learner_evaluation>
  <model_evaluation_result output_of='e1'>
    <evaluation name='accuracy' value= '0.99'>
    ...
```

ExpML: a markup language for DM experiments

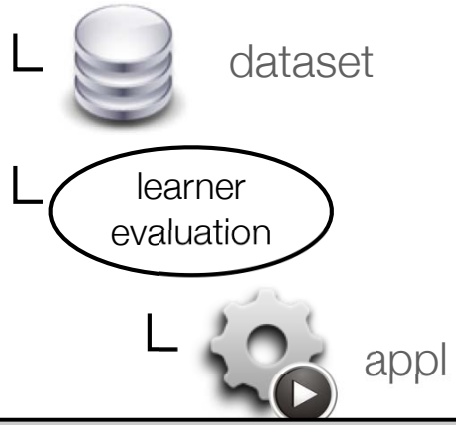

- Share DM experiments, XML-based

ontology	XML
has-part,has-participant	XML subelement
	(with role attribute)
has-description	(required) attribute
has-quality	`property' subelement
is-concretization-of	implementation_of attr.
part-of	attributes
has-specific-input	input_data attribute
has-specified-output	output_of attribute

```
<expml>
   <dataset id='d1'>
   <learner evaluation id='e1' input_data='d1'>
      <learner_appl>
         <learner_impl name=... version=...>
         <parameter_setting name='P' value='100'/>
         <learner_appl role= 'base-learner'>
            ...
   </learner_appl>
   <performance_estimation_appl>
   ...
   <model_evaluation_function_appl>
   ...
</learner_evaluation>
<model_evaluation_result output_of='e1'>
   <evaluation name='accuracy' value= '0.99'>
   ...
```

Use Case 2

Collect experiments in a database
to query all empirical results

ExpDB: a database to share experiments

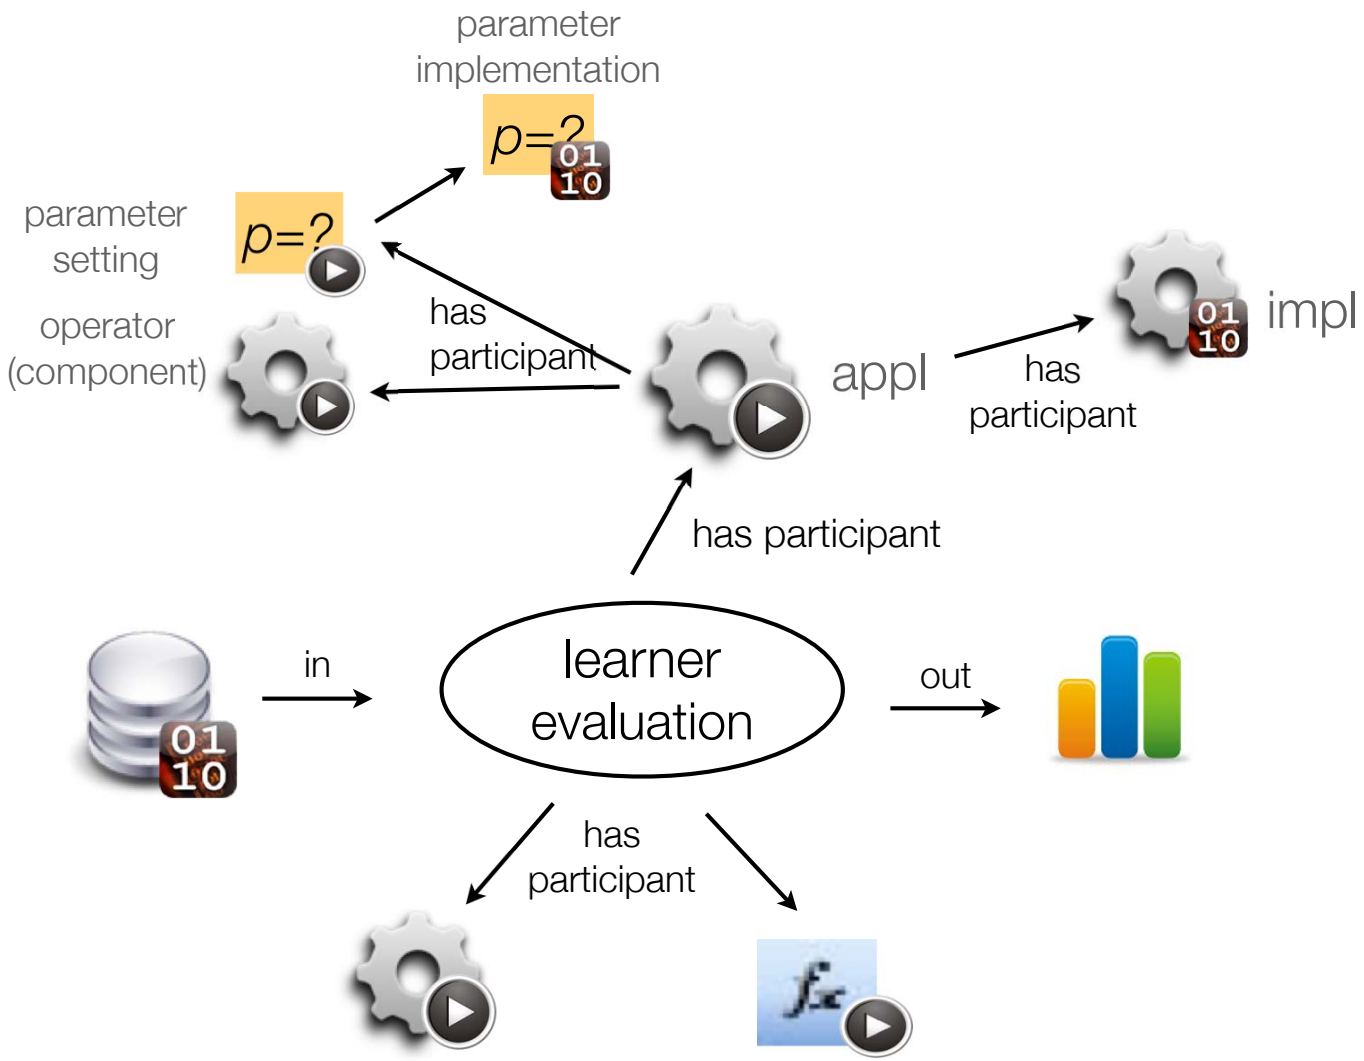

Experiment Database

>650,000 experiments, 54 algorithms,
>87 datasets, 45 evaluation measures,
2 data processors, bias-variance analysis

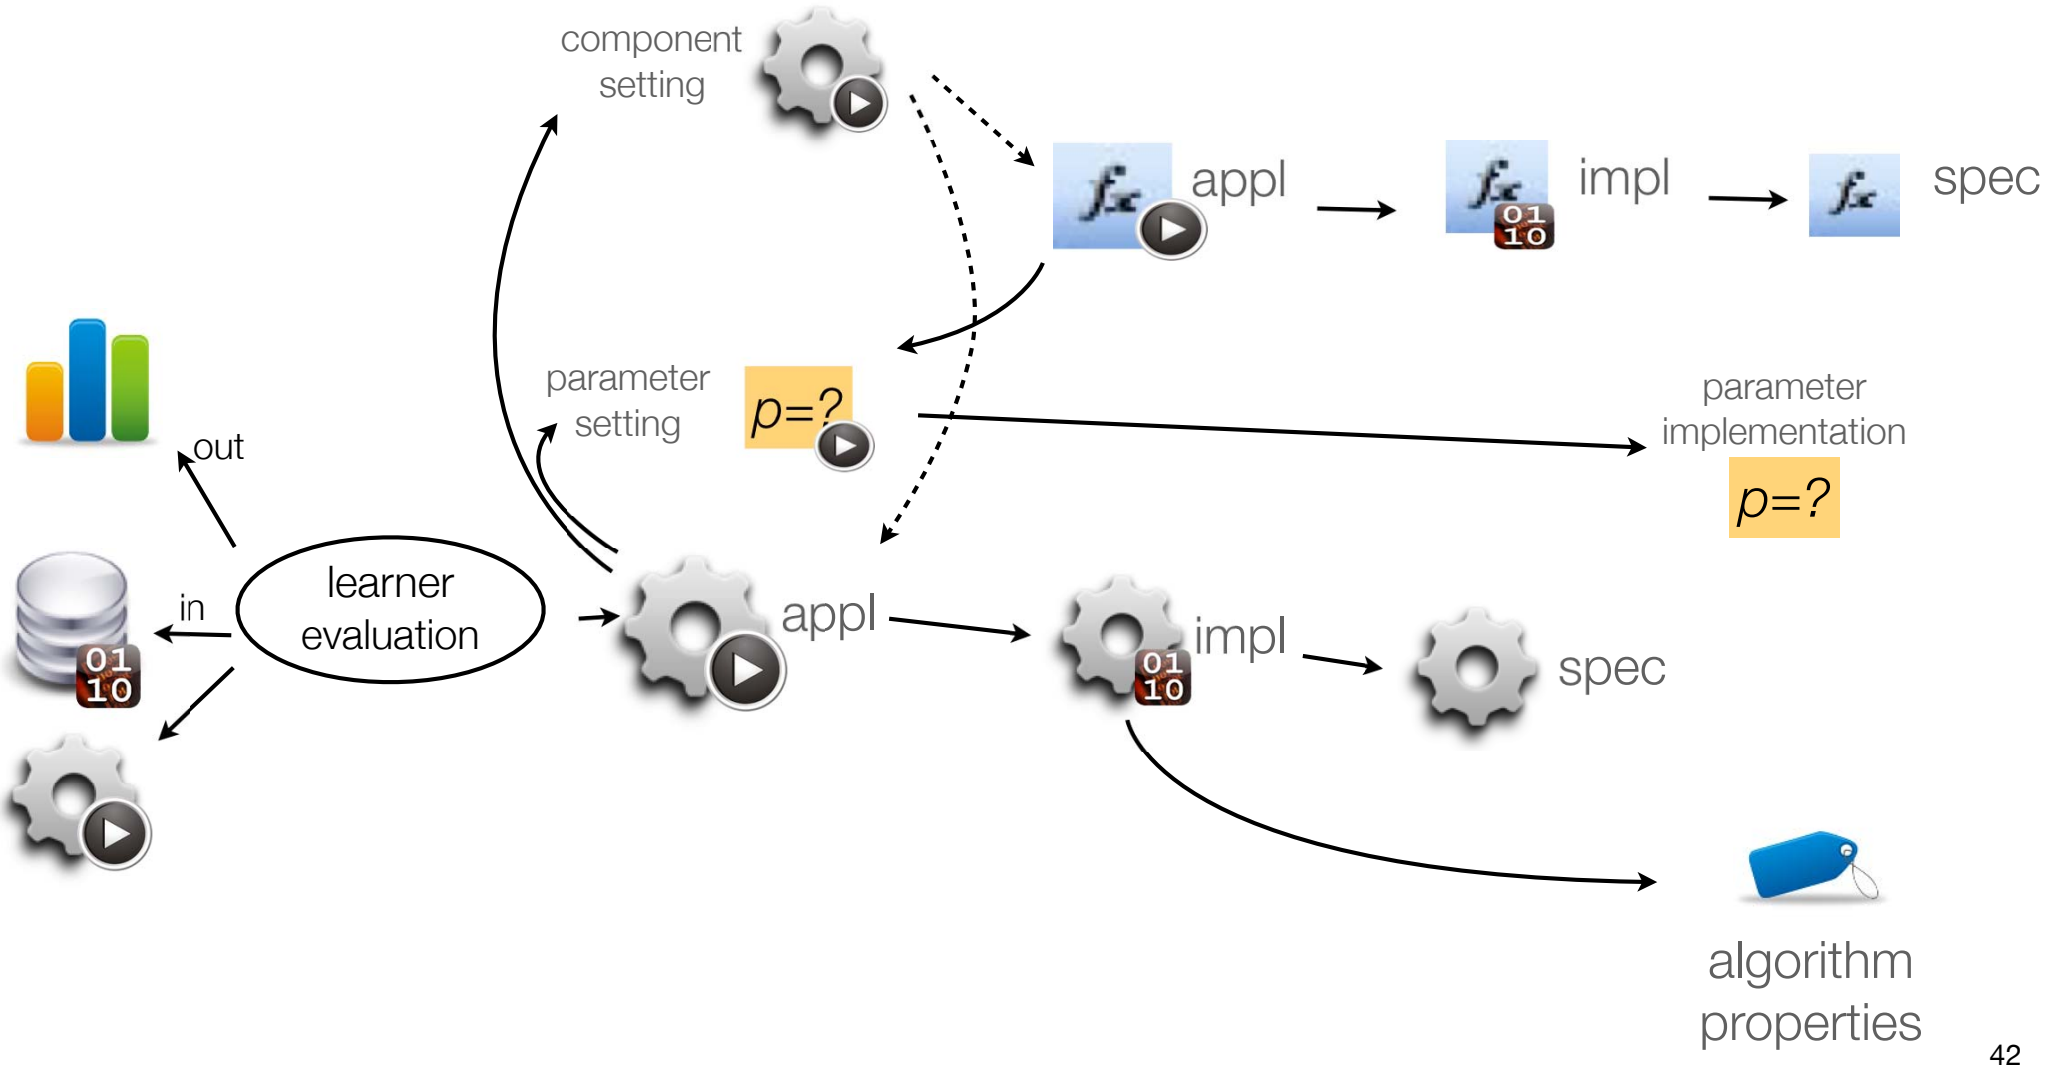

Experiment Database

>650,000 experiments, 54 algorithms,
>87 datasets, 45 evaluation measures,
2 data processors, bias-variance analysis

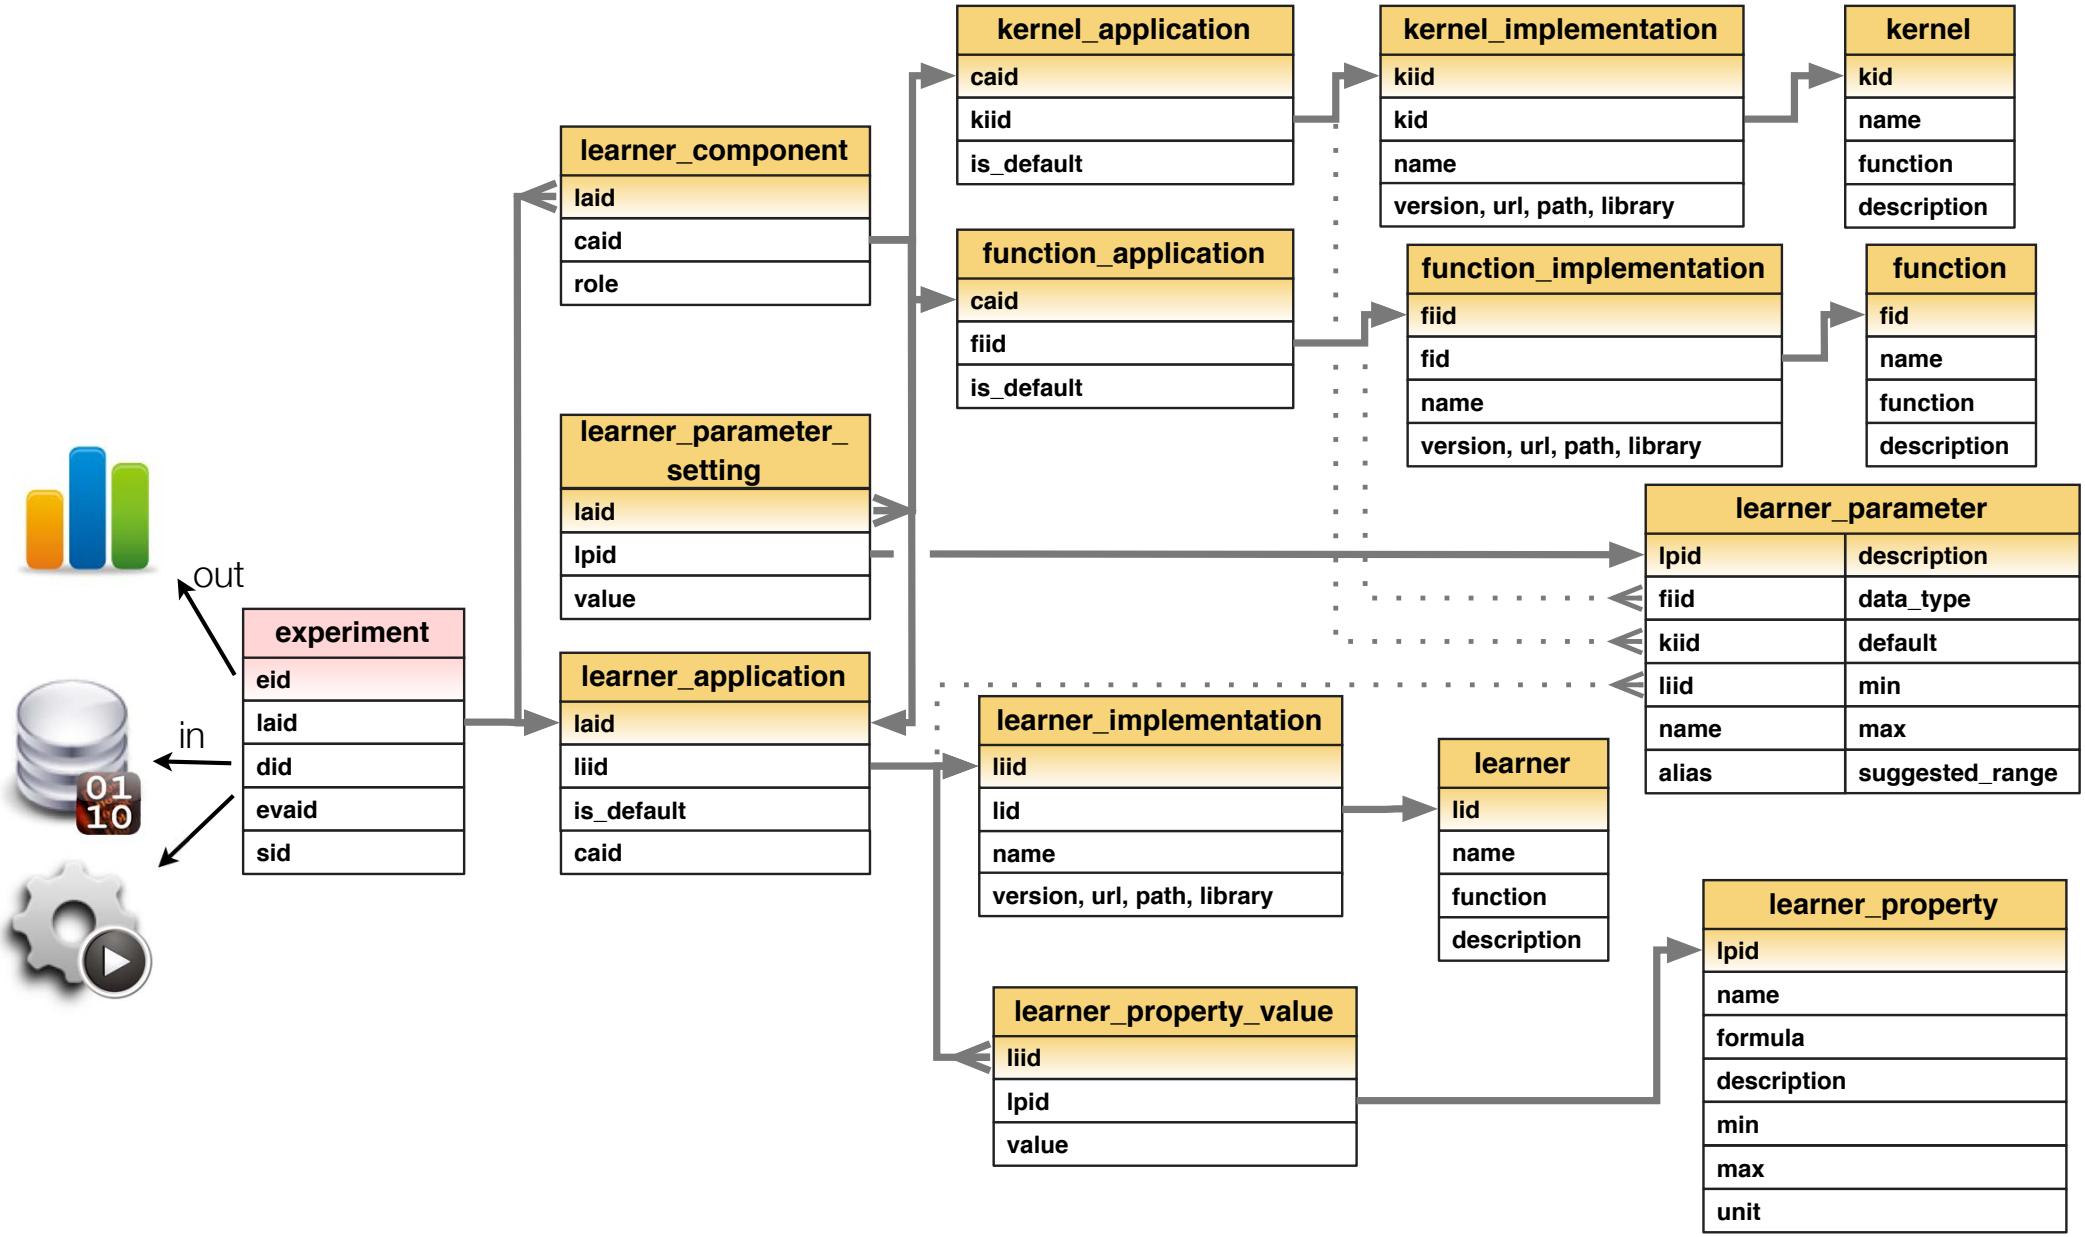

Use Case 3

Intuitive querying

Query Interface (YouTube “experiment database”)

http://expdb.cs.kuleuven.be

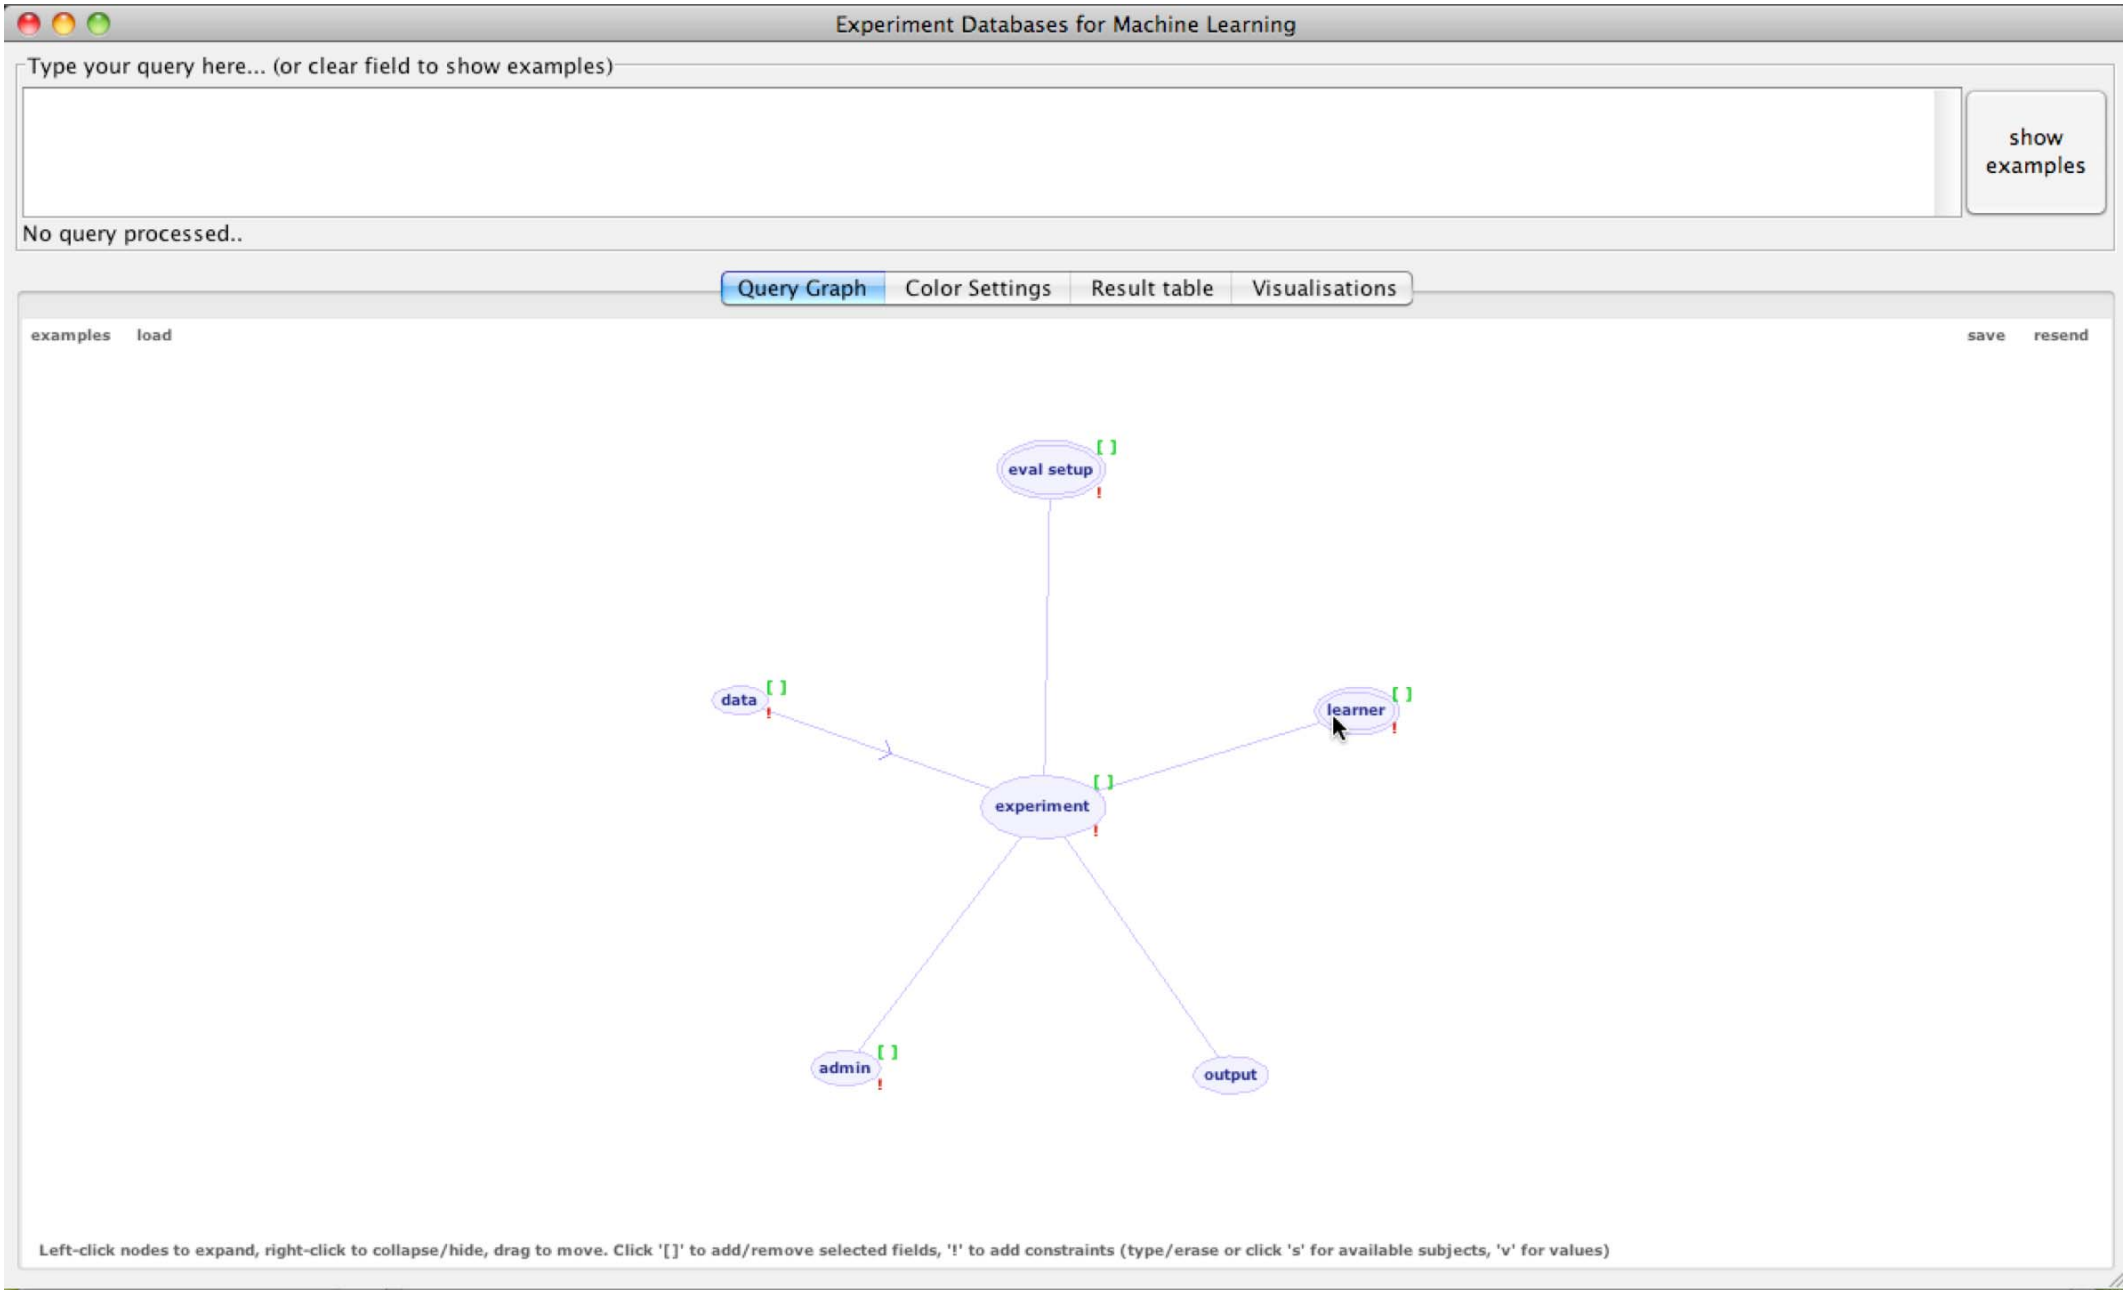

The way ahead

- 3rd generation of tools could make data mining into e-science
  - Experiments shared, reused, run worldwide
  - Repeatable, generalizable, reusable
- Cooperation on a standardized ontology for data mining?
- Automatic ontology extraction: DM paper -> ontology extension
- RDF experiment databases?
- Open problems:
  - Queriable models, auto-population (active meta-learning), quality control

Hvala

Danke

Thanks

Xie Xie

Diolch

Toda

Merci

Grazie

Spasiba

Efharisto

Obrigado

Arigato

Köszönöm

Tesekkurler

Dank U

Dhanyavaad

Gracias

http://expdb.cs.kuleuven.be